

# A COMPARATIVE STUDY OF TECHNIQUES FOR ESTIMATION AND INFERENCE OF NONLINEAR STOCHASTIC TIME SERIES

---

Dexter Barrows, B.Sc.

Masters Thesis Defence  
McMaster University  
April 2016

1. Framing
2. Stochastic SIR Model
3. Hamiltonian HMCMC
4. Iterated Filtering 2
5. Forecasting Frameworks
6. S-maps & Seasonal Outbreaks
7. Spatiotemporal Epidemics
8. Parallelism & Future Directions

Framing




SIResampling  
ParticleMCMC  
ParticleFilterDewSIS  
KalmanFilter  
HMRFHIF2  
MCMCProjection  
KalmanFilter  
ABC  
S-map  
HamiltonianMCMC  
ARIMA  
SeasonalARIMA  
EnsembleKF  
GARMA  
Regression  
GLM  
EnKF  
HamiltonianMCMC  
ARIMA



## Stochastic SIR model

$$\begin{aligned}\frac{dS}{dt} &= -\beta SI \\ \frac{dI}{dt} &= \beta SI - \gamma I \\ \frac{dR}{dt} &= \gamma I \\ &+ \end{aligned}$$

$$\beta_{t+1} = \exp \left[ \beta_t + \eta \left( \bar{\beta} - \beta_t \right) + \mathcal{N}(0, \sigma_{\text{proc}}) \right]$$



Hamiltonian MCMC

3



# MCMC

Iteratively construct  
Markov chain to  
approximate posterior

1. Choose starting parameter set
2. Generate N samples by
  - 2.1 Propose new sample
  - 2.2 Compute acceptance ratio
  - 2.3 Accept/reject sample

```
/* Select a starting point */
Input : Initialize  $\theta^{(1)}$ 
1 for i = 2 : N do
    /* Sample */
    2  $\theta^* \sim q(\cdot | \theta^{(i-1)})$ 
    3  $u \sim \mathcal{U}(0, 1)$ 
    /* Evaluate acceptance ratio */
    4  $r \leftarrow \frac{\mathcal{L}(\theta^*)p(\theta^*)}{\mathcal{L}(\theta)p(\theta)}$ 
    /* Step acceptance criterion */
    5 if  $u < \min\{1, r\}$  then
    6      $\theta^{(i)} = \theta^*$ 
    7 else
    8      $\theta^{(i)} = \theta^{(i-1)}$ 

/* Samples from approximated posterior distribution */
Output: Chain of samples  $(\theta^{(1)}, \theta^{(2)}, \dots, \theta^{(N)})$ 
```

# Hamiltonian Dynamics

Energy

Kinetic

$$K(r) = \frac{1}{2} r^T M^{-1} r$$

Potential

$$U(\theta) = -\log(\mathcal{L}(\theta)p(\theta))$$

Hamiltonian

$$H(\theta, r) = U(\theta) + K(r)$$

Dynamics simulation

$$\begin{aligned} \frac{d\theta}{dt} &= M^{-1} r \\ \frac{dr}{dt} &= -\nabla U(\theta) \end{aligned}$$

# HMCMC algorithm

1. Choose starting parameter set
2. Generate N samples by
  - 2.1 Resampling moments
  - 2.2 Simulate Hamiltonian dynamics using Leapfrog integration
  - 2.3 Compute acceptance ratio
  - 2.4 Accept/reject sample


```
/* Select a starting point */
Input : Initialize  $\theta^{(1)}$ 
1 for i = 2 : N do
    /* Resample moments */
2     for i = 1 : n do
3          $r(i) \leftarrow \mathcal{N}(0, I)$ 

    /* Leapfrog initialization */
4      $\theta_0 \leftarrow \theta^{(i-1)}$ 
5      $r_0 \leftarrow r - \nabla U(\theta_0) \cdot \varepsilon/2$ 
    /* Leapfrog intermediate steps */
6     for j = 1 : L - 1 do
7          $\theta_j \leftarrow \theta_{j-1} + M^{-1} r_{j-1} \cdot \varepsilon$ 
8          $r_j \leftarrow r_{j-1} - \nabla U(\theta_j) \cdot \varepsilon$ 

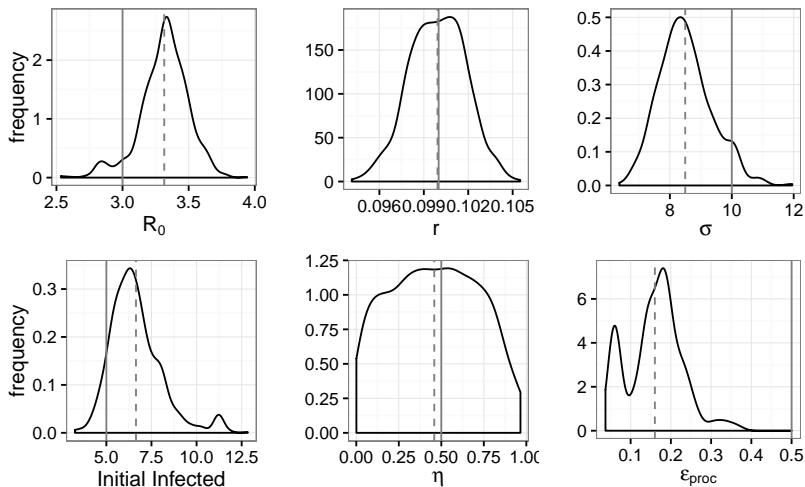
    /* Leapfrog last steps */
9      $\theta^* \leftarrow \theta_{L-1} + M^{-1} r_{L-1} \cdot \varepsilon$ 
10     $r^* \leftarrow \nabla U(\theta_L) \cdot \varepsilon/2 - r_{L-1}$ 
    /* Evaluate acceptance ratio */
11     $r = \exp [H(\theta^{(i-1)}, r) - H(\theta^*, r^*)]$ 
    /* Sample */
12     $u \sim \mathcal{U}(0, 1)$ 
    /* Step acceptance criterion */
13    if  $u < \min \{1, r\}$  then
14         $\theta^{(i)} = \theta^*$ 
15    else
16         $\theta^{(i)} = \theta^{(i-1)}$ 

/* Samples from approximated posterior distribution */
Output: Chain of samples
 $(\theta^{(1)}, \theta^{(2)}, \dots, \theta^{(N)})$ 
```

## HMCMC state reconstruction sample



## HMCMC parameter estimation kernels sample



Iterated Filtering 2



# Basic Particle Filter

Iterative prediction-update cycle  
prunes particle cohort of poor  
parameter estimates

1. Initialize particles with parameter sets
2. For each data point
  - 2.1 Evolve particle states
  - 2.2 Weight via likelihood
  - 2.3 Resample proportional to weights

```
/* Select a starting point */
Input  : Observations  $D = y_1, y_2, \dots, y_T$ , initial particle
         distribution  $P_0$  of size  $J$ 

/* Setup */
1 Initialize particle cohort by sampling
  ( $p^{(1)}, p^{(2)}, \dots, p^{(J)}$ ) from  $P_0$ 
2 for  $t = 1 : T$  do
    /* Evolve */
    3   for  $j = 1 : J$  do
    4      $X_t^{(j)} \leftarrow f_1(X_{t-1}^{(j)}, \theta^{(j)})$ 

    /* Weight */
    5   for  $j = 1 : J$  do
    6      $w^{(j)} \leftarrow P(y_t | X_t^{(j)}, \theta^{(j)}) = f_2(X_t^{(j)}, \theta^{(j)})$ 

    /* Normalize */
    7   for  $j = 1 : J$  do
    8      $w^{(j)} \leftarrow w^{(j)} / \sum_i w^{(i)}$ 

    /* Resample */
    9    $p^{(1:J)} \leftarrow \text{sample}(p^{(1:J)}, \text{prob} = w, \text{replace} = \text{true})$ 

/* Samples from approximated posterior distribution */
Output: Cohort of posterior samples
        ( $\theta^{(1)}, \theta^{(2)}, \dots, \theta^{(J)}$ )
```

# Iterated Filtering 2


Treat parameter estimates as stochastic processes

Multiple passes through data (a la Data cloning)







```
1  /* Setup */
2  Initialize particle cohort by sampling
    $(p^{(1)}, p^{(2)}, \dots, p^{(J)})$  from  $P_0$ 
3  /* Particle seeding distribution */
4   $\Theta \leftarrow P_0$ 
5  for m = 1 : M do
6      /* Pass perturbation */
7      for j = 1:J do
8           $p^{(j)} \sim h(\Theta^{(j)}, \sigma_m)$ 
9          for t = 1 : T do
10             for j = 1:J do
11                 /* Iteration perturbation */
12                  $p^{(j)} \sim h(p^{(j)}, \sigma_m)$ 
13                 /* Evolve */
14                  $X_t^{(j)} \leftarrow f_1(X_{t-1}^{(j)}, \theta^{(j)})$ 
15                 /* Weight */
16                  $w^{(j)} \leftarrow P(y_t | X_t^{(j)}, \theta^{(j)}) =$ 
17                      $f_2(X_t^{(j)}, \theta^{(j)})$ 
18             /* Normalize */
19             for j = 1:J do
20                  $w^{(j)} \leftarrow w^{(j)} / \sum_l w^{(l)}$ 
21             /* Resample */
22              $p^{(1:J)} \leftarrow \text{sample}(p^{(1:J)}, \text{prob} =$ 
23                  $w, \text{replace} = \text{true})$ 
24             /* Collect particles for next pass */
25             for j = 1 : J do
26                  $\Theta^{(j)} \leftarrow p^{(j)}$ 
27             /* Samples from approximated posterior distribution */
28             Output: Cohort of posterior samples
29                  $(\theta^{(1)}, \theta^{(2)}, \dots, \theta^{(J)})$ 
```









## IF2 state reconstruction sample



## IF2 parameter estimation kernels sample




# Mean parameter estimation distributions



IF2    HMCMC

## Running times



# Forecasting Frameworks



# IF2 + Parametric bootstrapping + Forward simulation

- Provides additional samples from posterior distribution
- Forward simulation using states point estimate, posterior samples

Input : Forward simulator  $S(\theta)$ , data set  $D$

```
/* Initial fit */
1  $\theta^* \leftarrow \text{IF2}(D)$ 


/* Generate artificial data sets */
2 for  $i = 1 : M$  do
3    $D_i \leftarrow S(\theta^*)$ 

/* Fit to new data sets */
4 for  $i = 1 : M$  do
5    $\theta_i \leftarrow \text{IF2}(D_i)$ 

/* Simulate forward to produce forecasts */
6 for  $i = 1 : M$  do
7    $F_i \leftarrow S(D, \theta_i)$ 

Output: Forecast trajectories
 $F_1, F_2, \dots, F_M$ 
```

## IF2 parametric bootstrapping forecast



# HMCMC + State reconstructions + Forward simulation

- Reconstruct states using latent process noise samples
- Simulate forward

```
/* Sample from the posterior */
Input : Posterior samples  $\theta_1, \theta_2, \dots, \theta_M$ 


1 /* Reconstruct state estimates */
2 for  $i = 1 : M$  do
3    $D_i \leftarrow S(\theta_i)$ 

/* Simulate forward to produce forecasts */
4 for  $i = 1 : M$  do
5    $F_i \leftarrow S(D_i, \theta_i)$ 

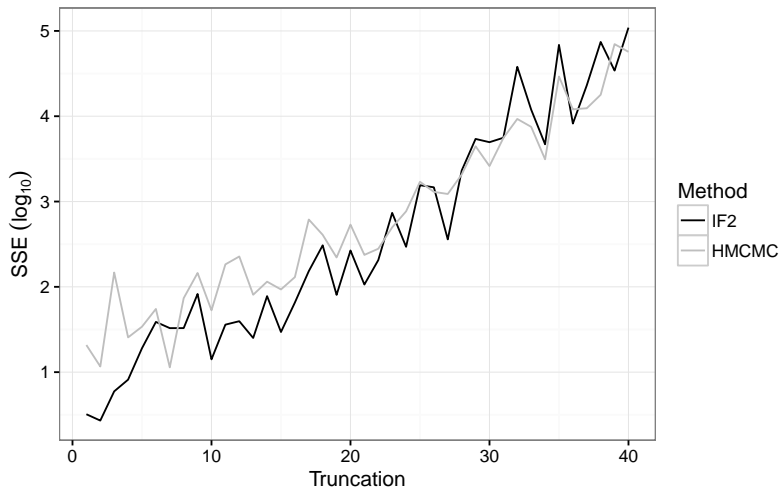
Output: Forecast trajectories  $F_1, F_2, \dots, F_M$ 
```



## HMCMC parametric bootstrapping forecast



## Forecast accuracy comparison



S-maps &  
Seasonal Outbreaks



## Stochastic SIRS model


$$\frac{dS}{dt} = -\beta SI + \alpha R$$

$$\frac{dI}{dt} = \beta SI - \gamma I$$

$$\frac{dR}{dt} = \gamma I - \alpha R$$

+

$$\beta_{t+1} = \exp \left[ \beta_t + \eta \left( \bar{\beta} - \beta_t \right) + \mathcal{N}(0, \sigma_{\text{proc}}) \right]$$



## S-map

- Construct global mapping of time-lagged vectors (the library) to future states
- Weightings are used to penalize poorly matching library vectors

```

/* Select a starting point */
Input : Time series  $x_1, x_2, \dots, x_T$ ,
        embedding dimension  $E$ , distance
        penalization  $\theta$ , forecast length  $L$ ,
        predictor vector  $x_t$ 

/* Construct library  $\{x_i\}$  */
1 for  $i = E : T$  do
2    $x_i = (x_i, x_{i-1}, \dots, x_{i-E+1})$ 

/* Construct mapping from library vectors
   to predictions */
3 for  $i = 1 : (T_E + 1)$  do
4   for  $j = 1 : E$  do
5      $A(i, j) = w(|x_i - x_t|)x_i(j)$ 
6   for  $i = 1 : (T_E + 1)$  do
7      $b(i) = w(|x_i - x_t|)y_i$ 


/* Use SVD to solve the mapping system,  $Ac = b$  */
8  $SVD(Ac = b)$ 

/* Compute forecast */
9  $\hat{y}_t = \sum_{j=0}^E c_t(j)x_t(j)$ 


/* Forecasted value in time series */
Output: Forecast  $\hat{y}_t$ 

```

## S-map forecast




## SIRS model forecasting error





## SIRS model forecasting runtimes



S-map: 316,000x faster than IF2, 61,800x faster than HMC MC

# Spatiotemporal Epidemics




## Stochastic Spatial SIR model

$$\begin{aligned}\frac{dS_i}{dt} &= - \left(1 - \phi \frac{M}{M+1}\right) \beta_i S_i I_i - \frac{\phi}{M+1} S_i \sum_{j=1}^M \beta_{ij} I_j \\ \frac{dI_i}{dt} &= \left(1 - \phi \frac{M}{M+1}\right) \beta_i S_i I_i + \frac{\phi}{M+1} S_i \sum_{j=1}^M \beta_{ij} I_j - \gamma I_i \\ \frac{dR_i}{dt} &= \gamma I_i\end{aligned}$$

+

$$\beta_{i,t+1} = \exp \left[ \beta_{i,t} + \eta \left( \bar{\beta} - \beta_{i,t} \right) + \mathcal{N}(0, \sigma_{\text{proc}}) \right]$$


## Stochastic spatial SIR model simulation (ring topology)




## Dewdrop Regression

- “Stitching” together multiple short time series into single library
- Requires scaling

## Stochastic SIR model forecasting error



## Spatial SIR model forecasting runtimes



# Parallelism & Future Directions






## Conclusions

- IF2 produces superior forecasts in all scenarios
- S-mapping runs orders of magnitude faster than other methods


## More than Moore

- Moore's Law is ceasing to hold
- Focus now on distributed computing
- MCMC-based methods are resistant to parallelization
  - Chain construction requires iterative dependence
- IF2 exhibits high parallel potential
  - Preliminary CUDA (GPU-accelerated) implementation - **cuIF2**

## Spatial SIR model fitting times



cuIF2: 9.33x faster than IF2, 617x faster than HMCMC



THE BEST THESIS DEFENSE IS A GOOD THESIS OFFENSE.