

ESTIMATION AND INFERENCE OF NONLINEAR
STOCHASTIC TIME SERIES

A COMPARATIVE STUDY OF TECHNIQUES FOR ESTIMATION AND
INFERENCE OF NONLINEAR STOCHASTIC TIME SERIES

By
DEXTER BARROWS, B.SC.

A Thesis Submitted to the School of Graduate Studies
in Partial Fulfilment of the Requirements for the Degree
Master of Science

McMaster University ©Copyright by Dexter Barrows, April 2016

MASTER OF SCIENCE (2016)
(Mathematics)

McMaster University
Hamilton, Ontario, Canada

TITLE	A Comparative Study of Techniques for Estimation and Inference of Nonlinear Stochastic Time Series
AUTHOR	Dexter Barrows, B.Sc. (Honours), Ryerson University
SUPERVISOR	Dr. Benjamin Bolker
NUMBER OF PAGES	x, 145

Abstract

Forecasting tools play an important role in public response to epidemics. Despite this, limited work has been done in comparing best-in-class techniques across the broad spectrum of time series forecasting methodologies. Forecasting frameworks were developed that utilised three methods designed to work with nonlinear dynamics: Iterated Filtering (IF) 2, Hamiltonian MCMC, and S-mapping. These were compared in several forecasting scenarios including a seasonal epidemic and a spatiotemporal epidemic. IF2 combined with parametric bootstrapping produced superior predictions in all scenarios. S-mapping combined with Dewdrop Regression produced forecasts slightly less-accurate than IF2 and HMCMC, but demonstrated vastly reduced running times. Hence, S-mapping with or without Dewdrop Regression should be used to glean initial insight into future epidemic behaviour, while IF2 and parametric bootstrapping should be used to refine forecast estimates in time.

Acknowledgements

There are many people I have to thank for their support over the last two years:

My supervisor Dr. Ben Bolker for his mentorship, advice, direction, and especially patience.

My defence committee members Dr. Jonathan Dushoff and Dr. David Earn who have taken the time to read my work and provide valuable input.

The Theobio lab for including me in stimulating discussions, even when they were over my head.

My Mom, Dad, Joel, and Sofia for being there for me through good times and trying ones.

Contents

1	Introduction	1
2	Hamiltonian MCMC	5
2.1	Markov Chains	5
2.2	Likelihood	6
2.3	Prior distribution	7
2.4	Proposal distribution	7
2.5	Algorithm	8
2.6	Burn-in	9
2.7	Thinning	9
2.8	Hamiltonian Monte Carlo	10
2.9	RStan Fitting	13
3	Iterated Filtering	18
3.1	Formulation	18
3.2	Algorithm	19
3.3	Particle Collapse	21
3.4	Iterated Filtering and Data Cloning	21
3.5	Iterated Filtering 2 (IF2)	22
3.6	IF2 Fitting	24
4	Parameter Fitting	26
4.1	Fitting Setup	26
4.2	Calibrating Samples	29
4.3	IF2 Fitting	30
4.4	IF2 Convergence	32
4.5	IF2 Densities	34
4.6	HMCMC Fitting	35
4.7	HMCMC Densities	35
4.8	HMCMC and Bootstrapping	36
4.9	Multi-trajectory Parameter Estimation	37

5	Forecasting Frameworks	39
5.1	Data Setup	39
5.2	IF2	40
5.2.1	Parametric Bootstrapping	41
5.2.2	IF2 Forecasts	41
5.3	HMCMC	43
5.4	Truncation vs. Error	44
6	S-map and SIRS	46
6.1	S-maps	46
6.2	S-map Algorithm	47
6.3	SIRS Model	49
6.4	SIRS Model Forecasting	51
7	Spatial Epidemics	54
7.1	Spatial SIR	54
7.2	Dewdrop Regression	56
7.3	Spatial Model Forecasting	57
8	Discussion and Future Directions	59
8.1	Parallel and Distributed Computing	59
8.2	IF2, Bootstrapping, and Forecasting Methodology	62
8.3	Fin	63
A	Hamiltonian MCMC	69
A.1	Full R code	69
A.2	Full Stan code	72
B	Iterated Filtering	74
B.1	Full R code	74
B.2	Full C++ code	76
C	Parameter Fitting	87
D	Forecasting Frameworks	88
D.1	IF2 Parametric Bootstrapping Function	88
D.2	RStan Forward Simulator	90
E	S-map and SIRS	92
E.1	SIRS R Function Code	92
E.2	SMAP Code	93
E.3	SMAP Parameter Optimization Code	95
E.4	RStan SIRS Code	97
E.5	IF2 SIRS Code	98

F	Spatial Epidemics	112
F.1	Spatial SIR R Function Code	112
F.2	RStan Spatial SIR Code	114
F.3	IF2 Spatial SIR Code	116
F.4	CUDA IF2 Spatial Fitting Code	127

List of Figures

2.1	A finite state machine. States are shown as graph nodes, and the probability of transitioning from one particular state to another is shown as a weighted graph edge. [2]	6
2.2	True SIR ODE solution infected counts, and with added observation noise.	14
2.3	Traceplot of samples drawn for parameter R_0 , excluding warmup.	15
2.4	Traceplot of samples drawn for parameter R_0 , including warmup.	16
2.5	Kernel density estimates produced by Stan. Dashed lines show true parameter values.	17
3.1	Kernel estimates for four essential system parameters. True values are indicated by dashed lines.	25
4.1	Simulated geometric autoregressive process show in Equation [4.1].	27
4.2	Density plot of values shown if Figure[4.1].	28
4.3	Stochastic SIR model simulated using an explicit Euler stepping scheme. The solid line is a single random trajectory, the dots show the data points obtained by adding in observation error defined as $\epsilon_{\text{obs}} = \mathcal{N}(0, 10)$, and the grey ribbon is centre 95th quantile from 100 random trajectories.	28
4.4	Fitting errors.	31
4.5	True system trajectory (solid line), observed data (dots), and IF2 estimated real state (dashed line).	31
4.6	The horizontal axis shows the IF2 pass number. The solid black lines show the evolution of the ML estimates, the solid grey lines show the true value, and the dashed grey lines show the mean parameter estimates from the particle swarm after the final pass.	33
4.7	The horizontal axis shows the IF2 pass number and the solid black lines show the evolution of the standard deviations of the particle swarm values.	33
4.8	As before, the solid grey lines show the true parameter values and the dashed grey lines show the density means.	35

4.9	As before, the solid grey lines show the true parameter values and the dashed grey lines show the density means.	36
4.10	Result from 100 HMC MC bootstrap trajectories. The solid line shows the true states, the dots show the data, the dotted line shows the average system behaviour, the dashed line shows the bootstrap mean, and the grey ribbon shows the centre 95th quantile of the bootstrap trajectories.	37
4.11	IF2 point estimate densities are shown in black and HMC MC point estimate densities are shown in grey. The vertical black lines show the true parameter values.	38
4.12	Fitting times for IF2 and HMC MC, in seconds. The centre box in each plot shows the centre 50th quantile, with the bold centre line showing the median.	38
5.1	Infection count data truncated at $T = 30$. The solid line shows the true underlying system states, and the dots show those states with added observation noise. Parameters used were $R_0 = 3.0$, $r = 0.1$, $\eta = .05$, $\sigma_{proc} = 0.5$, and additive observation noise was drawn from $\mathcal{N}(0, 10)$	40
5.2	Infection count data truncated at $T = 30$ from Figure [5.1]. The dashed line shows IF2's attempt to reconstruct the true underlying state from the observed data points.	41
5.3	Forecast produced by the IF2 / parametric bootstrapping framework. The dotted line shows the mean estimate of the forecasts, the dark grey ribbon shows the 95% confidence interval based on the 0.025 and 0.975 quantiles on the true state estimates, and the lighter grey ribbon shows the same confidence interval on the true state estimates with added observation noise drawn from $\mathcal{N}(0, \sigma)$	42
5.4	Forecast produced by the HMC MC / bootstrapping framework with $M = 200$ trajectories. The dotted line shows the mean estimate of the forecasts, and the grey ribbon shows the 95% confidence interval on the estimated true states as described in Figure [5.3].	44
5.5	Error growth as a function of data truncation amount. Both methods used 200 bootstrap trajectories. Note that the y-axis shows the natural log of the averaged SSE, not the total SSE.	45
6.1	Five cycles generated by the SIRS function. The solid line the the true number of cases, dots show case counts with added observation noise. The parameter values were $R_0 = 3.0$, $\gamma = 0.1$, $\eta = 1$, $\sigma = 5$, and 10 initial cases.	50
6.2	S-map applied to the data from the previous figure. The solid line shows the infection counts with observation noise from the previous plot, and the dotted line is the S-map forecast. Parameters chosen were $E = 14$ and $\theta = 3$	51

6.3	Error as a function of forecast length.	52
6.4	Runtimes for producing SIRS forecasts. The box shows the middle 50th quantile, the bold line is the median, and the dots are outliers. Note that these are not “true” outliers, simply ones outside a ranges based on the interquartile range.	53
7.1	Evolution of a spatial epidemic in a ring topology. The outbreak was started with 5 cases in Location 2. Parameters were $R_0 = 3.0$, $\gamma = 0.1$, $\eta = 0.5$, $\sigma_{err} = 0.5$, and $\phi = 0.5$	55
7.2	Evolution of a spatial epidemic as in Figure [7.1], with added observation noise drawn from $\mathcal{N}(0, 10)$	56
7.3	Average SSE (log scale) across each location and all trials as a function of the number of weeks ahead in the forecast.	58
7.4	Runtimes for producing spatial SIR forecasts. The box shows the middle 50th quantile, the bold line is the median, and the dots are outliers.	58
F.1	Running times for fitting the spatial SIR model to data.	145

Chapter 1

Introduction

Epidemic forecasting is an important tool that can help inform public policy and decision-making in the face of an infectious disease outbreak [9][27][39]. Successful intervention relies on accurate predictions of the number of cases, when they will occur, and where. Without this information it is difficult to efficiently allocate resources, a critical step in curbing the size and duration of an epidemic.

Despite the importance of reliable forecasts, obtaining them remains a challenge from both a theoretical and practical standpoint [27]. Mathematical models can capture the essential drivers in disease dynamics, and extend them past the present into the future. However, different epidemics may present with varying dynamics and require different model parameters to be accurately represented [7]. These parameters can be inferred by using statistical model fitting techniques, but this can become computationally intensive, and the modeller risks “overfitting” by attempting to capture too many drivers with too little data. Thus, the modeller must exercise restraint in model selection and fitting technique [5].

Securing precise, error-free data in the midst of an outbreak can be difficult if not impossible [33], thus observational uncertainty must be built into mathematical models of disease spread from the beginning. Models must differentiate between natural variation in the intensity of disease spread (process error) and error in data collection (observation error) in order to accurately determine the dynamics underlying a data set, adding another layer of complexity [22]. With these caveats and concerns acknowledged, we can turn to a discussion of technique.

Broadly, there are three primary categories of techniques used in forecasting: phenomenological, pure mechanistic, and semi-mechanistic.

Phenomenological methods operate purely on data, fitting models that do not try to reconstruct disease dynamics, but rather focus purely on trend. A long-standing and widely-used example is the Autoregressive Integrated Moving Average (ARIMA)

1 model. ARIMA assumes a linear underlying process and Gaussian error distributions.
 2 It uses three parameters representing the degree of autoregression (p), integration
 3 (trend removal) (d), and the moving average (q), where the orders of the autoregression
 4 and the moving average are determined through the use of an autocorrelation function
 5 (ACF) and partial autocorrelation function (PACF), respectively, applied to the the
 6 data *a priori* [40].

7 Pure mechanistic approaches simply try to capture the essential drivers in the disease
 8 spreading process and use the model alone to generate predictions. For example one
 9 could use a model in which individuals are divided into categories based on whether
 10 they are susceptible to infection or infected but not yet themselves infectious, infec-
 11 tious, or recovered. These are called compartmental models and are heavily used in
 12 epidemiological studies. Typically the transition between compartments is governed
 13 by a set of ordinary differential equations, such as

$$\begin{aligned}
 \frac{dS}{dt} &= -\beta IS \\
 \frac{dI}{dt} &= \beta IS - \gamma I \\
 \frac{dR}{dt} &= \gamma I,
 \end{aligned}
 \tag{1.1}$$

15 where S , I , and R are the number of individuals in each compartment, β is contact
 16 rate between susceptible and infected individuals, and γ is a recovery rate. We also
 17 let $\beta = R_0 r / N$, where R_0 is the number of secondary cases per infected individual
 18 in a wholly susceptible population, and N is the population size. As an outbreak
 19 progresses, individuals transition from the susceptible compartment, through the in-
 20 fectious compartment, then finish in the removed compartment where they no longer
 21 impact the system dynamics. Many extensions of the SIR model exist and are com-
 22 monly used, such as the SEIR model in which susceptible individuals pass through an
 23 exposed class (or several) where they have been infected but are not yet themselves
 24 infectious, and the SIRS model in which individuals become susceptible again after
 25 their immunity wanes [7].

26 Combining the phenomenological and mechanistic approaches are data integration
 27 schemes. These methods use a model to define the expected underlying dynamics
 28 of the system, but integrate data into the model in order to refine estimates of the
 29 model parameters and produce more accurate forecasts. Typically the first step in
 30 implementing such a technique is fitting the desired model to existing data. There are
 31 many ways to do this, most of which fall into two main categories: Sequential Monte
 32 Carlo (SMC) methods [3][31][39], and Markov chain Monte Carlo-based (MCMC)
 33 methods [2][26]. From there data can either be integrated into the model by refitting
 34 the model to the new longer data set, or in an “on-line” fashion in which data points

1 can be directly integrated without the need to refit the entire model. Normally,
2 MCMC-based machinery must refit the entire model whereas SMC-based approaches
3 can sometimes integrate data in an on-line fashion.

4 Another, broader, distinction among techniques can be drawn between those that rely
5 on assumptions of linearity, and those that make no such assumption. As epidemic
6 dynamics are highly non-linear, it can be questionable as to even consider linear ap-
7 proaches to epidemic forecasting at all. In particular, stalwart approaches such as
8 ARIMA and the venerable Kalman filter face a distinct (at least theoretical) disad-
9 vantage in the face of newer SMC-based methods [36][39]. Extensions of the Kalman
10 filter, such as the ensemble adjustment Kalman filter are designed to handle non-
11 linearity, but these methods are very well-studied, and further work showing their
12 viability would likely prove extraneous in the modern academic landscape.

13 Somewhat frustratingly, there exists no “gold standard” in forecasting [9][27][39]. As
14 methodology varies widely in theoretical justification, implementation, and operation,
15 it is difficult to compare the state of the art in forecasting methods from a first-
16 principles perspective. Further, published work making use of any of these methods
17 to forecast uses different prediction accuracy metrics, such as SSE, peak time/dura-
18 tion/intensity, correlation tests, or RMSE, among others [9][28]. Thus it is difficult
19 to select the best tool for the job when faced with a forecasting problem.

20 The primary focus of this work is to compare best-in-class methods for forecasting
21 in several epidemically-focused scenarios. These include the a “standard” one-shot
22 forecast outbreak in which the outbreak subsides and does not recur, a seasonal out-
23 break scenario such as the one we see with influenza each year, and a spatiotemporal
24 scenario in which multiple spatial locations are connected and disease is free to spread
25 from one to another.

26 From MCMC-based methods we have selected Hamiltonian MCMC [26], a (slightly)
27 less cutting-edge but nonetheless highly effective technique. We are using HMC
28 through an implementation in the R package RStan [8], which at its core uses HM-
29 CMC, but also contains implementations of several other innovative techniques. In-
30 terestingly, the original goal of this package was not to implement a statistical pro-
31 gramming language similar to Just Another Gibbs Sampler (JAGS) or Bayesian In-
32 ference Using Gibbs Sampling (BUGS), but with an HMC backend. In fact the
33 developers’ original goal was to implement any method that could fit multilevel hi-
34 erarchical models without halting as they were witnessing with BUGS and JAGS.
35 Only after experimenting with several options and starting to hear about it more and
36 more frequently did they attempt to work with HMC. In the end, the scope of the
37 project grew to include the development and subsequent integration of the No-U-Turn
38 Sampler (NUTS) [17], and an implementation of automatic differentiation machinery
39 [34].

40 For PF-based methods we have selected Iterated Filtering (IF) 2 [20], a very recently

1 developed approach that uses multiple particle filtering rounds to generate maximum
2 likelihood estimates (MLEs). It functions similarly to its predecessor, the Maximum
3 likelihood via Iterated Filtering (MIF) algorithm [19], but aims to be simpler, faster,
4 and more accurate. Theoretical justification and synthetic testing indicates that IF2
5 meets these goals, and as such the authors recommend skipping MIF and jumping
6 straight to IF2 if an algorithm of their variety is sought. And so, we are doing just
7 that. We wrote our own IF2 implementation in C++ and integrated it into R using
8 the Rcpp package [12]. The developers of MIF and IF2 have their own R package that
9 implements MIF and IF2, Partially Observed Markov Processes (POMP) [21][23], but
10 it didn't provide some of the diagnostic information we needed, so it was not used
11 here.

12 Finally, from the phenomenological methods we have selected the sequential locally
13 weighted global linear maps (S-map) [13][18][35][36], combined with Dewdrop Regres-
14 sion [13]. These methods stand on their own as a unique take on the forecasting
15 problem, and bear little resemblance to other methodology. The virtues of these tech-
16 niques have been long-extolled by their developers, but their efficacy when compared
17 to competing methods has not been well-studied. This work will mark one of the first
18 times this has been done.

19 This paper will begin with descriptions of HMCMC and IF2 with examples
20 of simple model fitting in Chapters 2 and 3. Chapter 4 explores parameter fitting of
21 a stochastic SIR model to synthetic data. Chapter 5 will establish the full forecasting
22 frameworks used with IF2 and HMCMC, and compare them in a simple scenario. All
23 three methods will be used to compare forecasts using a SIRS model in Chapter 6.
24 Chapter 7 will show forecasts using the aforementioned IF2 and HMCMC frameworks,
25 along with Dewdrop Regression combined with S-mapping. Finally, a summary of
26 these results, and a discussion of parallel computing and future directions will finish
27 the paper in Chapter 8.

Chapter 2

Hamiltonian MCMC

Markov Chain Monte Carlo (MCMC) is a general class of methods designed to sample from the posterior distribution of model parameters [2]. It is an algorithm used when we wish to fit a model M that depends on some parameter (or more typically vector of parameters) θ to observed data D . MCMC works by constructing a Markov chain whose stationary distribution converges to desired posterior distribution. The samples drawn using MCMC are used to numerically approximate the stationary distribution, and in turn the posterior [2].

2.1 Markov Chains

Figure [2.1] shows a finite state machine with 3 states $S = \{x_1, x_2, x_3\}$.

The transition probabilities can be summarized as a matrix as

$$T = \begin{bmatrix} 0 & 1 & 0 \\ 0 & 0.1 & 0.9 \\ 0.6 & 0.4 & 0 \end{bmatrix}. \quad (2.1)$$

The probability vector $\mu(x^{(1)})$ for a state $x^{(1)}$ can be evolved using T by evaluating $\mu(x^{(1)})T$, then again by evaluating $\mu(x^{(1)})T^2$, and so on. If we take the limit as the number of transitions approaches infinity, we find

$$\lim_{t \rightarrow \infty} \mu(x^{(1)})T^t = (27/122, 50/122, 45/122). \quad (2.2)$$

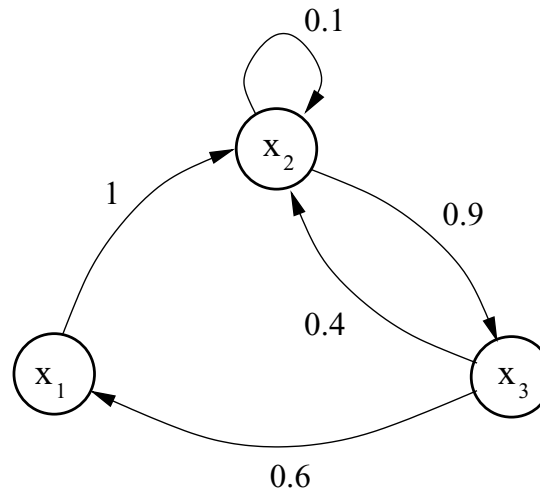


Figure 2.1: A finite state machine. States are shown as graph nodes, and the probability of transitioning from one particular state to another is shown as a weighted graph edge. [2]

1 This indicates that no matter what we pick for the initial probability distribution
 2 $\mu(x^{(1)})$, the chain will always stabilize at the equilibrium distribution.

3 This property holds when the chain satisfies the following conditions

- 4 • *Irreducible* Any state A can be reached from any other state B with non-zero
 5 probability
- 6 • *Positive Recurrent* The number of steps required for the chain to reach state A
 7 from state B must be finite
- 8 • *Aperiodic* The chain must be able to explore the parameter space without be-
 9 coming trapped in a cycle

10 Note that MCMC sampling generates a Markov chain $(\theta^{(1)}, \theta^{(2)}, \dots, \theta^{(N)})$ that does
 11 indeed satisfy these conditions, and uses the chain's equilibrium distribution to ap-
 12 proximate the posterior distribution of the parameter space [2].

13 2.2 Likelihood

14 MCMC and similar methods hinge on the idea that the weight or support bestowed
 15 upon a particular set of parameters θ should be proportional to the probability of
 16 observing the data D given the model output using that set of parameters $M(\theta)$. In
 17 order to do this we need a way to evaluate whether or not $M(\theta)$ is a good fit for D ;

1 this is done by specifying a likelihood function $\mathcal{L}(\theta)$ such that

$$2 \qquad \qquad \qquad \mathcal{L}(\theta) \propto P(D|\theta). \qquad \qquad \qquad (2.3)$$

3 In some standard Maximum Likelihood approaches, $\mathcal{L}(\theta)$ is searched to find a value
4 of θ that maximizes $\mathcal{L}(\theta)$, then this θ is taken to be the most likely true value. Here
5 our aim is to not just maximize the likelihood but to also explore the area around it
6 [2].

7 **2.3 Prior distribution**

8 Another significant component of MCMC is the user-specified prior distribution for
9 θ or distributions for the individual components of θ (priors). Priors serve as a way
10 for us to tell the MCMC algorithm what we think consist of good values for the
11 parameters. Note that if very little is known about the parameters, or we are worried
12 about biasing our estimate of the posterior, we can simply use a a wide uniform
13 distribution. We cannot, however, avoid this problem entirely. Bayesian frameworks,
14 such as MCMC, *require* priors to be specified; what the user must decide is how strong
15 to make priors.

16 Exceedingly weak priors can pose problematic in some circumstances. In the case of
17 MCMC, weak priors handicap the algorithm in two ways: convergence of the chain
18 may become exceedingly slow, and more pressure is put on the likelihood function to
19 be as good as possible – it will now be the only thing informing the algorithm of what
20 constitutes a “good” set of parameters, and what should be considered poor. In the
21 majority of cases this does not pose as much as problem as it would appear; if enough
22 samples are drawn, we should still obtain a good posterior estimate. We will only
23 really run into problems if an exceedingly weak prior, such as an unbounded uniform
24 distribution, or another unbounded distribution with a high standard deviation – in
25 those cases we may obtain poor posterior estimates if the data are weak [2].

26 **2.4 Proposal distribution**

27 As part of the MCMC algorithm, when we find a state in the parameter space that
28 is accepted as part of the Markov chain construction process, we need a good way
29 of generating a good next step to try. Unlike basic rejection sampling in which we
30 would just randomly sample from our prior distribution, MCMC attempts to optimise
31 our choices by choosing a step that is close enough to the last accepted step so as to

stand a decent chance of also being accepted, but far enough away that it doesn't get "trapped" in a particular region of the parameter space.

This is done through the use of a proposal or candidate distribution. This will usually be a distribution centred around our last accepted step and with a dispersion potential narrower than that of our prior distribution.

The choice of this distribution is theoretically not of the utmost importance, but in practice becomes important so as to not waste computer time [2].

2.5 Algorithm

Now that we have all the pieces necessary, we can discuss the details of the MCMC algorithm.

We will denote the previously discussed quantities as

- $p(\cdot)$ - the prior distribution
- $q(\cdot|\cdot)$ - the proposal distribution
- $\mathcal{L}(\cdot)$ - the Likelihood function
- $\mathcal{U}(\cdot, \cdot)$ - the uniform distribution

and then define the acceptance ratio, r , as

$$r = \frac{\mathcal{L}(\theta^*)p(\theta^*)q(\theta|\theta^*)}{\mathcal{L}(\theta)p(\theta)q(\theta^*|\theta)}, \quad (2.4)$$

where θ^* is the proposed sample to draw from the posterior, and θ is the last accepted sample. This is known as the Metropolis-Hastings rule.

In the special case of the Metropolis variation of MCMC, the proposal distribution is symmetric, meaning $q(\theta^*|\theta) = q(\theta|\theta^*)$, and so the acceptance ratio simplifies to

$$r = \frac{\mathcal{L}(\theta^*)p(\theta^*)}{\mathcal{L}(\theta)p(\theta)}. \quad (2.5)$$

Algorithm [1] shows the Metropolis MCMC algorithm.

In this way we are ensuring that steps that lead to better likelihood outcomes are likely to be accepted, but steps that do not will not be accepted as frequently. Note that these less "advantageous" moves will still occur but that this is by design – it ensures that as much of the parameter space as possible will be explored but more efficiently than using pure brute force [2].

Algorithm 1: Metropolis MCMC

```

/* Select a starting point                                     */
Input : Initialize  $\theta^{(1)}$ 
1 for  $i = 2 : N$  do
    /* Sample                                                 */
    2  $\theta^* \sim q(\cdot | \theta^{(i-1)})$ 
    3  $u \sim \mathcal{U}(0, 1)$ 
    /* Evaluate acceptance ratio                               */
    4  $r \leftarrow \frac{\mathcal{L}(\theta^*)p(\theta^*)}{\mathcal{L}(\theta)p(\theta)}$ 
    /* Step acceptance criterion                               */
    5 if  $u < \min\{1, r\}$  then
    6 |  $\theta^{(i)} = \theta^*$ 
    7 else
    8 |  $\theta^{(i)} = \theta^{(i-1)}$ 
    /* Samples from approximated posterior distribution       */
Output: Chain of samples  $(\theta^{(1)}, \theta^{(2)}, \dots, \theta^{(N)})$ 

```

1

2.6 Burn-in

2 One critical aspect of MCMC-based algorithms has yet to be discussed. The algorithm
3 requires an initial starting point θ to be selected, but as the proposal distribution
4 is supposed to restrict moves to an area close to the last accepted state, then the
5 posterior distribution will be biased towards this starting point. This issue is avoided
6 through the use of a Burn-in period.

7 Burning in a chain is the act of running the MCMC algorithm normally without saving
8 first M samples. As we are seeking a chain of length N , the total computation will
9 be equivalent to generating a chain of length $M + N$ [2].

10

2.7 Thinning

11 Some models will require very long chains to get a good approximation of the posterior,
12 which will consequently require a non-trivial amount of computer storage. One way
13 to reduce the burden of storing so many samples is by thinning. This involves saving
14 only every n^{th} step, which should still give a decent approximate of the posterior (since
15 the chain has time to explore a large portion of the parameter space), but require less
16 room to store [24].

2.8 Hamiltonian Monte Carlo

The Metropolis-Hastings algorithm has a primary drawback in that the parameter space may not be explored efficiently in some circumstances – a consequence of the rudimentary proposal mechanism. Instead, smarter moves can be proposed through the use of Hamiltonian dynamics, leading to a better exploration of the target distribution and a potential decrease in overall computational complexity. This algorithm is coined Hamiltonian MCMC (HMC or HMCMC) [26]. Prior to the advent of HMCMC, some work was conducted exploring adaptive step-sizing using MCMC-based methods, but lack strong theoretical justification, and can lead to some samples being drawn from an incorrect distribution [26].

From physics, we will borrow the ideas of potential and kinetic energy. Here potential energy is analogous to the negative log likelihood of the parameter selection given the data, formally

$$U(\theta) = -\log(\mathcal{L}(\theta)p(\theta)). \quad (2.6)$$

Kinetic energy will serve as a way to “nudge” the parameters along a different moment for each component of θ . We introduce n auxiliary variables $r = (r_1, r_1, \dots, r_n)$, where n is the number of components in θ . Note that the samples drawn for r are not of interest, they are only used to inform the evolution of the Hamiltonian dynamics of the system. We can now define the kinetic energy as

$$K(r) = \frac{1}{2}r^T M^{-1}r, \quad (2.7)$$

where M is an $n \times n$ matrix. In practice M can simply be chosen as the identity matrix of size n , however it can also be used to account for correlation between components of θ .

The Hamiltonian of the system is defined as

$$H(\theta, r) = U(\theta) + K(r), \quad (2.8)$$

where the Hamiltonian dynamics of the combined system can be simulated using the following system of ODEs:

$$\begin{aligned} \frac{d\theta}{dt} &= M^{-1}r \\ \frac{dr}{dt} &= -\nabla U(\theta). \end{aligned} \quad (2.9)$$

1 .

2 It is tempting to try to integrate this system using the standard Euler evolution
 3 scheme, but in practice this leads to instability. Instead the “Leapfrog” scheme is
 4 used. This scheme is very similar to Euler scheme, except instead of using a fixed
 5 step size h for all evolutions, a step size of ε is used for most evolutions, with a half
 6 step size of $\varepsilon/2$ for evolutions of $\frac{dr}{dt}$ at the first step, and last step L . In this way the
 7 evolution steps “leapfrog” over each other while using future values from the other
 8 set of steps, leading to the scheme’s name.

9 The end product of the Leapfrog steps are the new proposed parameters (θ^*, r^*) .
 10 These are either accepted or rejected using a mechanism similar to that of standard
 11 Metropolis-Hastings MCMC. Now, however, the acceptance ratio r is defined as

$$12 \quad r = \exp [H(\theta, r) - H(\theta^*, r^*)], \quad (2.10)$$

13 where (θ, r) are the last values in the chain.

14 Together, we have Algorithm [2].

15 Note that the parameters ε and L have to be tuned in order to maintain stability
 16 and maximize efficiency, a sometimes non-trivial process [26]. However, some recent
 17 algorithms, such as the No U-Turn sampler implemented in RStan, and adaptively
 18 select appropriate values automatically during the sampling process [17].

Algorithm 2: Hamiltonian MCMC

```

/* Select a starting point */
Input : Initialize  $\theta^{(1)}$ 
1 for  $i = 2 : N$  do
    /* Resample moments */
    2 for  $i = 1 : n$  do
    3    $r(i) \leftarrow \mathcal{N}(0, 1)$ 

    /* Leapfrog initialization */
    4    $\theta_0 \leftarrow \theta^{(i-1)}$ 
    5    $r_0 \leftarrow r - \nabla U(\theta_0) \cdot \varepsilon/2$ 

    /* Leapfrog intermediate steps */
    6   for  $j = 1 : L - 1$  do
    7      $\theta_j \leftarrow \theta_{j-1} + M^{-1}r_{j-1} \cdot \varepsilon$ 
    8      $r_j \leftarrow r_{j-1} - \nabla U(\theta_j) \cdot \varepsilon$ 

    /* Leapfrog last steps */
    9    $\theta^* \leftarrow \theta_{L-1} + M^{-1}r_{L-1} \cdot \varepsilon$ 
    10   $r^* \leftarrow \nabla U(\theta_L) \cdot \varepsilon/2 - r_{L-1}$ 

    /* Evaluate acceptance ratio */
    11   $r = \exp [H(\theta^{(i-1)}, r) - H(\theta^*, r^*)]$ 

    /* Sample */
    12   $u \sim \mathcal{U}(0, 1)$ 

    /* Step acceptance criterion */
    13  if  $u < \min \{1, r\}$  then
    14     $\theta^{(i)} = \theta^*$ 
    15  else
    16     $\theta^{(i)} = \theta^{(i-1)}$ 

/* Samples from approximated posterior distribution */
Output: Chain of samples  $(\theta^{(1)}, \theta^{(2)}, \dots, \theta^{(N)})$ 

```

2.9 RStan Fitting

Here we will examine a test case in which Hamiltonian MCMC will be used to fit a Susceptible-Infected-Removed (SIR) epidemic model to mock infectious count data.

The synthetic data was produced by taking the solution to a basic SIR ODE model, sampling it at regular intervals, and perturbing those values by adding in observation noise. The SIR model used was outlined in the introduction in Equation [1.1].

The solution to this system was obtained using the `ode()` function from the `deSolve` package. The required derivative array function in the format required by `ode()` was specified as the gradient in Equation [1.1].

The true parameter values were set to $R_0 = 3.0, r = 0.1, N = 500$. The initial conditions were set to 5 infectious individuals, 495 people susceptible to infection, and no one had yet recovered from infection and been removed. The system was integrated over $[0, 100]$ with infected counts drawn at each integer time step.

The observation error was taken to be $\varepsilon_{obs} \sim \mathcal{N}(0, \sigma)$, where individual values were drawn for each synthetic data point.

Figure [2.2] shows the system simulation results.

The Hamiltonian MCMC model fitting was done using Stan (<http://mc-stan.org/>), a program written in C++ that does Bayesian statistical inference using Hamiltonian MCMC. Stan's R interface (<http://mc-stan.org/interfaces/rstan.html>) was used to ease implementation.

Throughout this paper, the explicit Euler integration scheme was used to obtain solutions to our ODE-based models. While this scheme is not the most accurate or efficient one available, it was chosen for its ease of implementation in the required languages and transparency with regards to stochastic processes, which have been added into later models. Using a more advanced integrator such as Runge-Kutta makes it harder to properly specify how stochastic process evolution should be handled, and would have required significantly more implementation work to boot. Hence, we have opted for the lo-fi solution we know will function the way we require.

In order to use an Explicit Euler-like stepping method in the later Stan model, the synthetic observation counts were treated as weekly observations in which the counts on the other six days of the week were unobserved.

Figure [2.3] shows the traceplot for the the post-warmup chain data returned by the `stan()` function in the `fit` object. We see that the chains are mixing well and convergence has likely been reached.

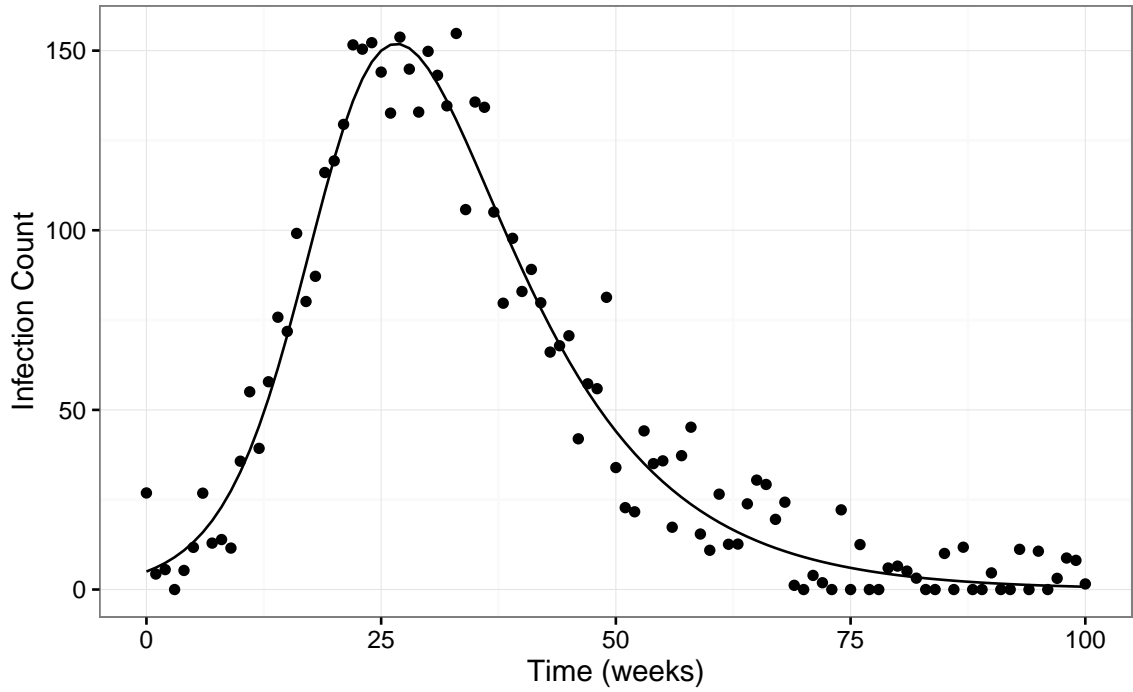


Figure 2.2: True SIR ODE solution infected counts, and with added observation noise.

- 1 Figure [2.4] shows the chain data including the warmup samples in. We can see why
- 2 is is wise to discard these samples (note the scale).
- 3 Figure [2.5] shows the the kernel density estimates for each of the model parameters
- 4 and the initial number of cases. We see that while the estimates are not perfect, they
- 5 are more than satisfactory.

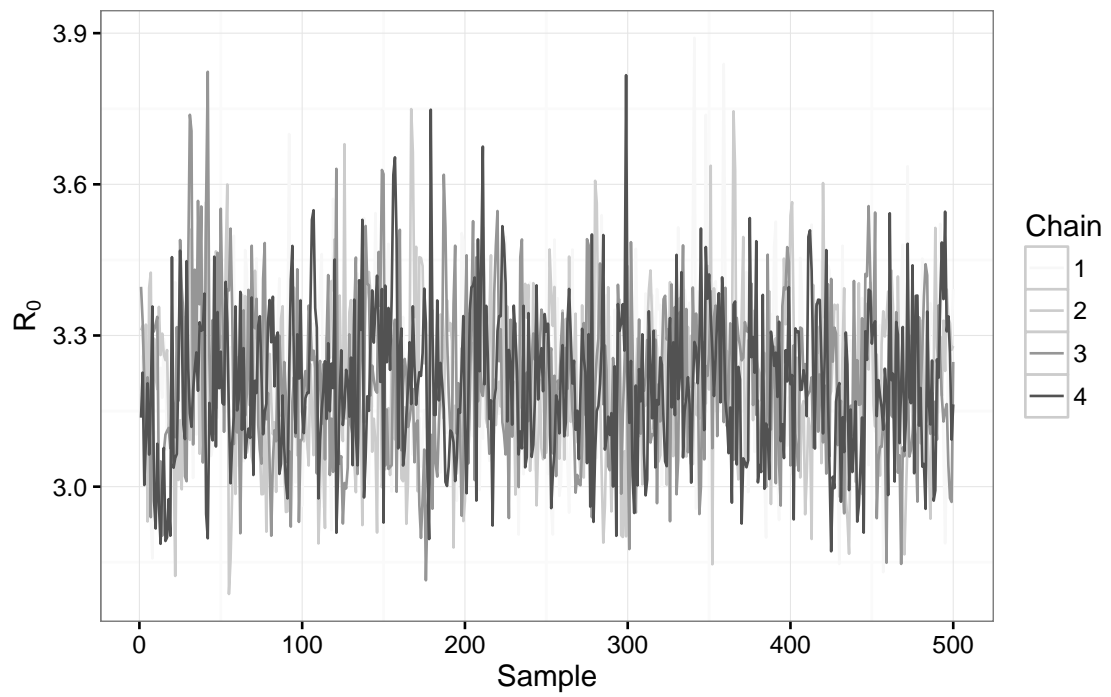


Figure 2.3: Traceplot of samples drawn for parameter R_0 , excluding warmup.

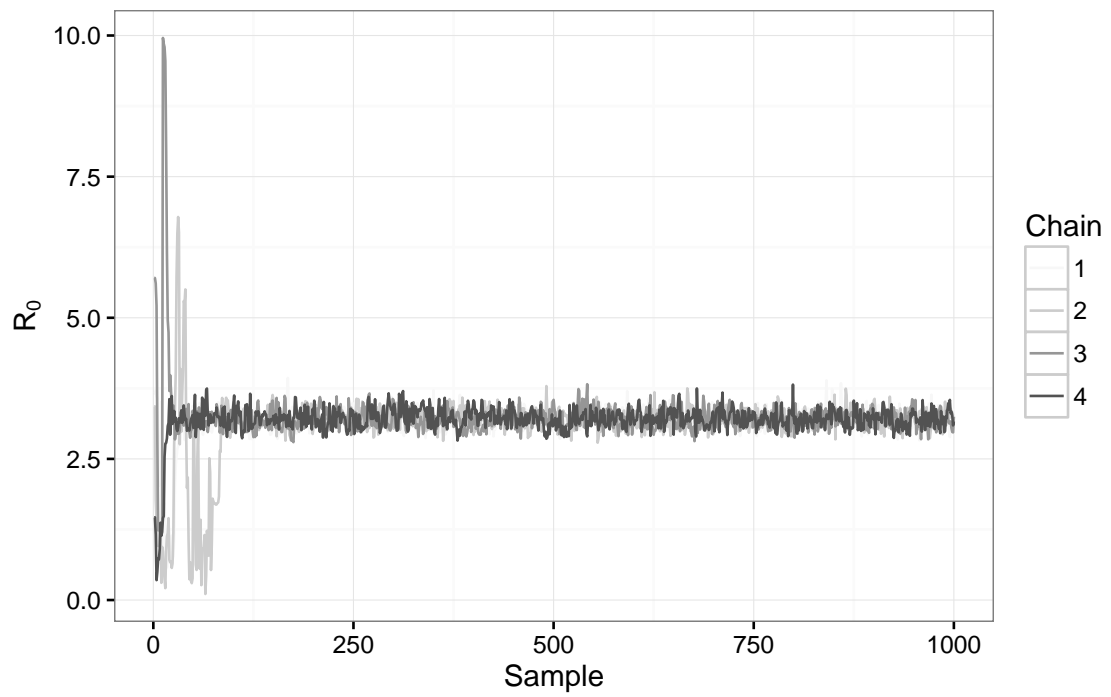


Figure 2.4: Traceplot of samples drawn for parameter R_0 , including warmup.

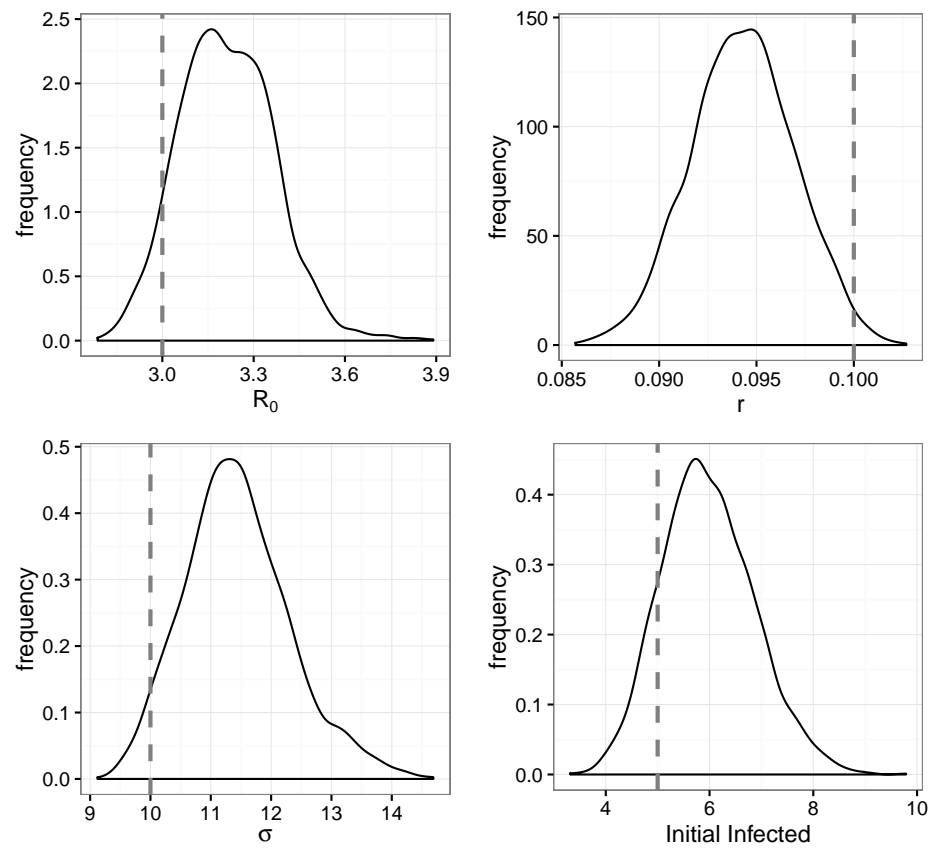


Figure 2.5: Kernel density estimates produced by Stan. Dashed lines show true parameter values.

Chapter 3

Iterated Filtering

Particle filters are similar to MCMC-based methods in that they use likelihoods to evaluate the validity of proposed parameter sets given observed data D , but differ in that they are largely trying to produce point estimates of the parameters instead of samples from the posterior distribution.

Instead of constructing a Markov chain and approximating its stationary distribution, a cohort of “particles” are used to move through the data in an on-line (sequential) fashion with the cohort being culled of poorly-performing particles at each iteration via importance sampling. If the culled particles are not replenished, this will be a Sequential Importance Sampling (SIS) particle filter. If the culled particles are replenished from surviving particles, in a sense setting up a process analogous to Darwinian selection, then this will be a Sequential Importance Resampling (SIR) particle filter [3].

3.1 Formulation

Particle filters, also called Sequential Monte-Carlo (SMC) filters, feature similar core functionality as the venerable Kalman Filter. As the algorithm moves through the data (sequence of observations), a prediction-update cycle is used to simulate the evolution of the model M with different particular parameter selections, track how closely these predictions approximate the new observed value, and update the current cohort appropriately [3].

Two separate functions are used to simulate the evolution and observation processes. The “true” state evolution is specified by

$$X_{t+1} \sim f_1(X_t, \theta), \tag{3.1}$$

1 And the observation process by

$$2 \qquad Y_t \sim f_2(X_t, \theta). \qquad (3.2)$$

3 Components of θ can contribute to both functions, but a typical formulation is to
 4 have some components contribute to $f_1(\cdot, \theta)$ and others to $f_2(\cdot, \theta)$.

5 The prediction part of the cycle uses $f_1(\cdot, \theta)$ to update each particle's current state
 6 estimate to the next time step, while $f_2(\cdot, \theta)$ is used to evaluate a weighting w for
 7 each particle which will be used to determine how closely that particle is estimating
 8 the true underlying state of the system. Note that $f_2(\cdot, \theta)$ could be thought of as a
 9 probability of observing a piece of data y_t given the particle's current state estimate
 10 and parameter set, $P(y_t|X_t, \theta)$. Then, the new cohort of particles is drawn from
 11 the old cohort proportional to the weights. This process is repeated until the set of
 12 observations D is exhausted.

13 **3.2 Algorithm**

14 Now we will formalize the particle filter.

15 We will denote each particle $p^{(j)}$ as the j^{th} particle consisting of a state estimate at
 16 time t , $X_t^{(j)}$, a parameter set $\theta^{(j)}$, and a weight $w^{(j)}$. Note that the state estimates
 17 will evolve with the system as the cohort traverses the data.

18 The algorithm for a Sequential Importance Resampling particle is shown in Algorithm
 19 [3].

Algorithm 3: SIR particle filter

```

/* Select a starting point */
Input : Observations  $D = y_1, y_2, \dots, y_T$ , initial particle distribution  $P_0$  of size
         $J$ 

/* Setup */
1 Initialize particle cohort by sampling  $(p^{(1)}, p^{(2)}, \dots, p^{(J)})$  from  $P_0$ 
2 for  $t = 1 : T$  do
    /* Evolve */
    3 for  $j = 1:J$  do
    4    $X_t^{(j)} \leftarrow f_1(X_{t-1}^{(j)}, \theta^{(j)})$ 

    /* Weight */
    5 for  $j = 1:J$  do
    6    $w^{(j)} \leftarrow P(y_t | X_t^{(j)}, \theta^{(j)}) = f_2(X_t^{(j)}, \theta^{(j)})$ 

    /* Normalize */
    7 for  $j = 1:J$  do
    8    $w^{(j)} \leftarrow w^{(j)} / \sum_1^J w^{(j)}$ 

    /* Resample */
    9  $p^{(1:J)} \leftarrow \text{sample}(p^{(1:J)}, \text{prob} = w, \text{replace} = \text{true})$ 

/* Samples from approximated posterior distribution */
Output: Cohort of posterior samples  $(\theta^{(1)}, \theta^{(2)}, \dots, \theta^{(J)})$ 

```

3.3 Particle Collapse

Often, a situation may arise in which a single particle is assigned a normalized weight very close to 1 and all the other particles are assigned weights very close to 0. When this occurs, the next generation of the cohort will overwhelmingly consist of descendants of the heavily-weighted particle, termed particle collapse or degeneracy [6][3].

Since the basic SIR particle filter does not perturb either the particle system states or system parameter values, the cohort will quickly consist solely of identical particles, effectively halting further exploration of the parameter space as new data is introduced.

A similar situation occurs when a small number of particles (but not necessarily a single particle) split almost all of the normalized weight between them, then jointly dominate the resampling process for the remainder of the iterations. This again halts the exploration of the parameter space with new data.

In either case, the hallmark feature used to detect collapse is the same – at some point the cohort will consist of particles with very similar or identical parameter sets which will consequently result in their assigned weights being extremely close together.

Mathematically, we are interested in the number of effective particles, N_{eff} , which represents the number of particles that are acceptably dissimilar. This is estimated by evaluating

$$N_{\text{eff}} = \frac{1}{\sum_1^J (w^{(j)})^2}. \quad (3.3)$$

This can be used to diagnose not only when collapse has occurred, but can also indicate when it is near [3].

3.4 Iterated Filtering and Data Cloning

A particle filter hinges on the idea that as it progresses through the data set D , its estimate of the posterior carried in the cohort of particles approaches maximum likelihood. However, this convergence may not be fast enough so that the estimate it produces is of quality before the data runs out. One way around this problem is to “clone” the data and make multiple passes through it as if it were a continuation of the original time series. Note that the system state contained in each particle will have to be reset with each pass.

1 Rigorous proofs have been developed [19][20] that show that by treating the param-
 2 eters as stochastic processes instead of fixed values, the multiple passes through the
 3 data will indeed force convergence of the process mean toward maximum likelihood,
 4 and the process variance toward 0.

5 **3.5 Iterated Filtering 2 (IF2)**

6 The successor to Iterated Filtering 1 [19], Iterated Filtering 2 [20] is simpler, faster,
 7 and demonstrated better convergence toward maximum likelihood. The core concept
 8 involves a two-pronged approach. First, a data cloning-like procedure is used to
 9 allow more time for the parameter stochastic process means to converge to maximum
 10 likelihood, and frequent cooled perturbation of the particle parameters allow better
 11 exploration of the parameter space while still allowing convergence to good point
 12 estimates.

13 IF2 is not designed to estimate the full posterior distribution, instead to produce
 14 a Maximum Likelihood (ML) point estimate. Further, IF2 thwarts the problem of
 15 particle collapse by keeping at least some perturbation in the system at all times. It
 16 is important to note that while true particle collapse will not occur, there is still risk
 17 of a pseudo-collapse in which all particles will be extremely close to one another so as
 18 to be virtually indistinguishable. However this will only occur with the use of overly-
 19 aggressive cooling strategies or by specifying an excessive number of passes through
 20 the data.

21 An important new quantity is the particle perturbation density denoted $h(\theta|\sigma)$. Typ-
 22 ically this is multivariate Normal with σ being a vector of variances proportional to
 23 the expected values of θ . In practice the proportionality can be derived from current
 24 means or specified ahead of time. Further, these intensities must decrease over time.
 25 This can be done via exponential or geometric cooling, a decreasing step function, a
 26 combination of these, or through some other similar scheme.

27 The algorithm for IF2 can be seen in Algorithm [4].

28

Algorithm 4: IF2

```

/* Select a starting point */
Input : Observations  $D = y_1, y_2, \dots, y_T$ , initial particle distribution  $P_0$  of size
         $J$ , decreasing sequence of perturbation intensity vectors  $\sigma_1, \sigma_2, \dots, \sigma_M$ 

/* Setup */
1 Initialize particle cohort by sampling  $(p^{(1)}, p^{(2)}, \dots, p^{(J)})$  from  $P_0$ 

/* Particle seeding distribution */
2  $\Theta \leftarrow P_0$ 
3 for  $m = 1 : M$  do
    /* Pass perturbation */
    4 for  $j = 1 : J$  do
    5      $p^{(j)} \sim h(\Theta^{(j)}, \sigma_m)$ 
    6 for  $t = 1 : T$  do
    7     for  $j = 1 : J$  do
    8         /* Iteration perturbation */
    9          $p^{(j)} \sim h(p^{(j)}, \sigma_m)$ 
    10        /* Evolve */
    11         $X_t^{(j)} \leftarrow f_1(X_{t-1}^{(j)}, \theta^{(j)})$ 
    12        /* Weight */
    13         $w^{(j)} \leftarrow P(y_t | X_t^{(j)}, \theta^{(j)}) = f_2(X_t^{(j)}, \theta^{(j)})$ 
    14        /* Normalize */
    15        for  $j = 1 : J$  do
    16             $w^{(j)} \leftarrow w^{(j)} / \sum_1^J w^{(j)}$ 
    17        /* Resample */
    18         $p^{(1:J)} \leftarrow \text{sample}(p^{(1:J)}, \text{prob} = w, \text{replace} = \text{true})$ 
    19        /* Collect particles for next pass */
    20        for  $j = 1 : J$  do
    21             $\Theta^{(j)} \leftarrow p^{(j)}$ 

/* Samples from approximated posterior distribution */
Output: Cohort of posterior samples  $(\theta^{(1)}, \theta^{(2)}, \dots, \theta^{(J)})$ 

```

1 3.6 IF2 Fitting

2 Here we will examine a test case in which IF2 will be used to fit a Susceptible-Infected-
 3 Removed (SIR) epidemic model to mock infectious count data.

4 As in the previous section, the model in Equation [1.1] was use to produce synthetic
 5 data. The same parameters and initial conditions were used, namely: parameter
 6 values were set to $R_0 = 3.0$, $r = 0.1$, $N = 500$, initial conditions were set to 5 infectious
 7 individuals, 495 people susceptible to infection, and no one had yet recovered from
 8 infection and been removed, and observation error was taken to be $\varepsilon_{obs} \sim \mathcal{N}(0, \sigma)$,
 9 where individual values were drawn for each synthetic data point.

10 Figure [2.2] in the previous section shows the true SIR ODE system solution and
 11 data.

12 The IF2 algorithm was implemented in C++ for speed, and integrated into the R
 13 workflow using the Rcpp package.

14 There are three primary reasons we implemented our own version of IF2 instead of
 15 using POMP. First, POMP does not provide final particle state distributions, making
 16 it difficult to calibrate the algorithm parameters against the parameters used in RStan
 17 (this procedure is described in the next chapter). Second, it is prudent to cross-check
 18 the validity of an algorithm using another implementation. Third, this code can then
 19 serve as a jumping-off point for further development using Graphics Processing Unit
 20 acceleration (outlined in Chapter 8). We must acknowledge the disadvantages as well:
 21 POMP has been extensively vetted with real-world usage, and using it would require
 22 far lees work as we would only need to specify the model. That being said, we believe
 23 the advantages outweigh the disadvantages in this case, and so have proceeded to
 24 develop our own implementation of IF2.

25 Figure [3.1] shows the final kernel estimates for four of the key parameters. As with
 26 HMCMC, the distributions are not perfect, but are promising. Unlike with HM-
 27 CMC, these distributions are not meant to consist of samples from the true posterior
 28 distribution, but rather serve a diagnostic role.

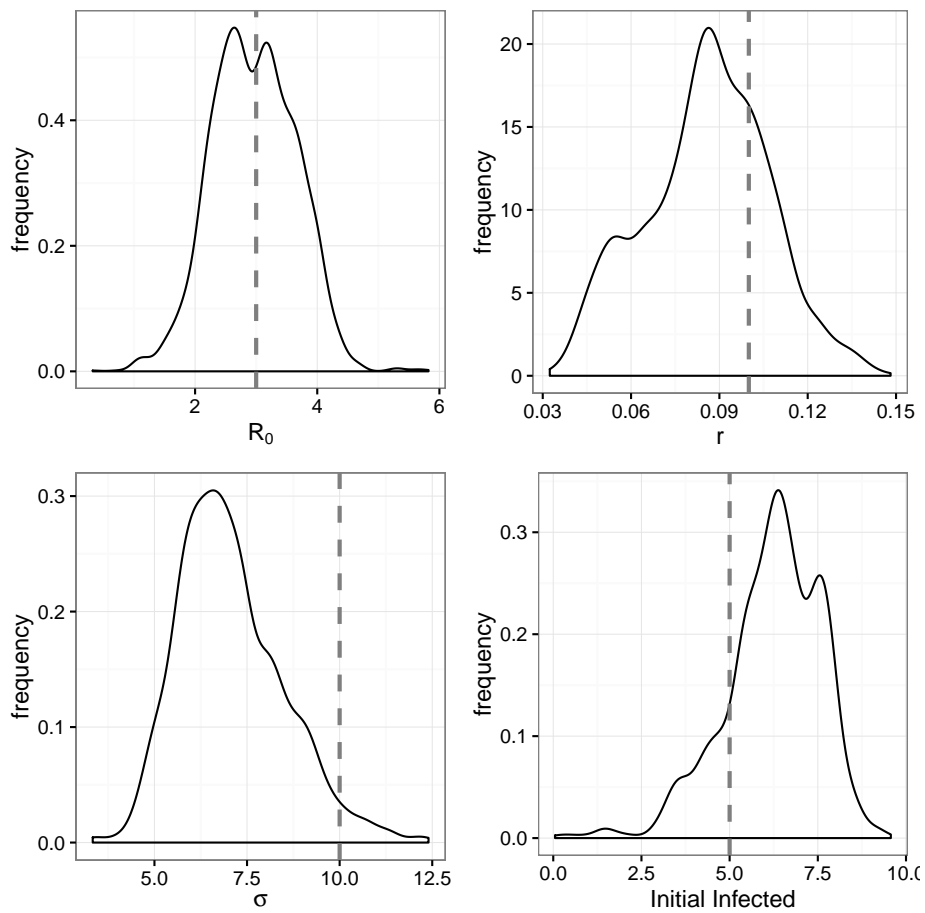


Figure 3.1: Kernel estimates for four essential system parameters. True values are indicated by dashed lines.

Chapter 4

Parameter Fitting

4.1 Fitting Setup

Now that we have established which methods we wish to evaluate the efficacy of for epidemic forecasting, it is prudent to see how they perform when fitting parameters for a known epidemic model. We have already seen how they perform when fitting parameters for a model with a deterministic evolution process and observation noise, but a more realistic model will have both process and observation noise.

To form such a model, we will take a deterministic SIR ODE model specified in Equation [1.1] and add process noise by allowing β to follow a geometric random walk given by

$$\beta_{t+1} = \exp(\log(\beta_t) + \eta(\log(\bar{\beta}) - \log(\beta_t)) + \epsilon_t). \quad (4.1)$$

We will take ϵ_t to be normally distributed with standard deviation ρ^2 such that $\epsilon_t \sim \mathcal{N}(0, \rho^2)$. The geometric attraction term constrains the random walk, the force of which is $\eta \in [0, 1]$. If we take $\eta = 0$ then the walk will be unconstrained; if we let $\eta = 1$ then all values of β_t will be independent from the previous value (and consequently all other values in the sequence).

When $\eta \in (0, 1)$, we have an autoregressive process of order 1 on the logarithmic scale of the form

$$X_{t+1} = c + \rho X_t + \epsilon_t, \quad (4.2)$$

where ϵ_t is normally distributed noise with mean 0 and standard deviation σ_E . This process has a theoretical expected mean of $\mu = c/(1-\rho)$ and variance $\sigma = \sigma_E^2/(1-\rho^2)$.

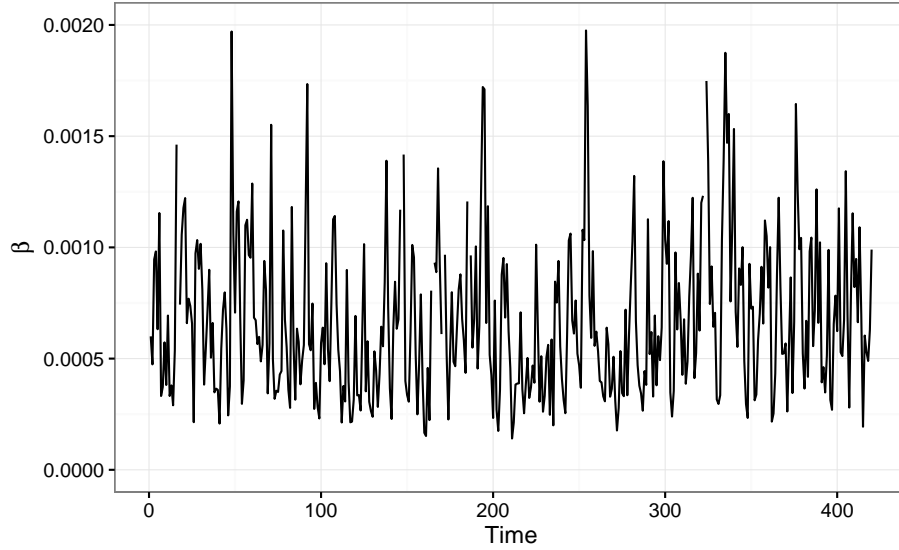


Figure 4.1: Simulated geometric autoregressive process show in Equation [4.1].

- 1 If we choose $\eta = 0.5$, the resulting log-normal distribution has a mean of 6.80×10^{-4}
- 2 and standard deviation of 4.46×10^{-4} .
- 3 Figure [4.1] shows the result of simulating the process in Equation [4.1] with $\eta =$
- 4 0.5.
- 5 Figure [4.2] shows the density plot corresponding to the values in Figure [4.1].
- 6 We see a density plot similar in shape to the desired density, and the geometric random
- 7 walk displays dependence on previous values. Further the mean of this distribution
- 8 was calculated to be 6.92×10^{-4} and standard the deviation to be 3.99×10^{-4} , which
- 9 are very close to the theoretical values.
- 10 If we take the full stochastic SIR system and evolve it using an Euler stepping scheme
- 11 with a step size of $h = 1/7$, for 1 step per day, we obtain the plot in Figure [4.3].

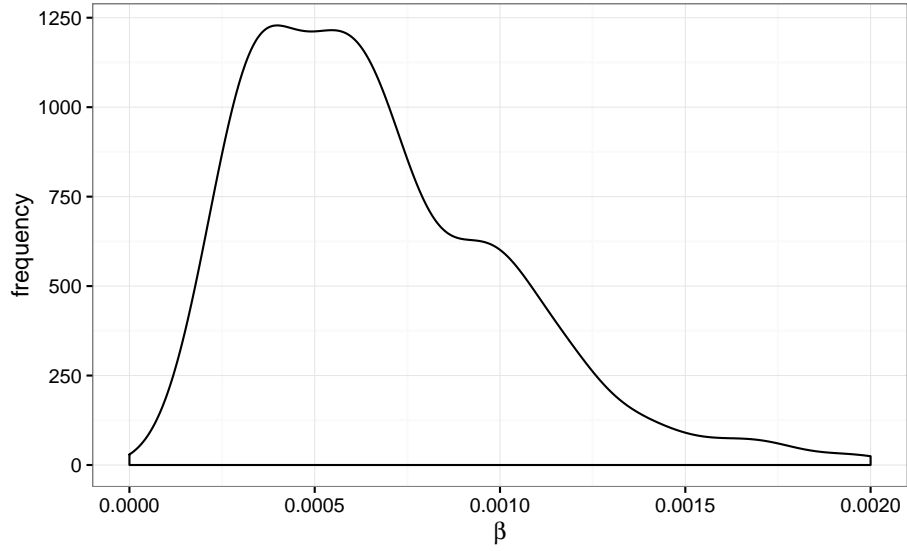


Figure 4.2: Density plot of values shown in Figure 4.1.




Figure 4.3: Stochastic SIR model simulated using an explicit Euler stepping scheme. The solid line is a single random trajectory, the dots show the data points obtained by adding in observation error defined as $\epsilon_{\text{obs}} = \mathcal{N}(0, 10)$, and the grey ribbon is centre 95th quantile from 100 random trajectories.

1 4.2 Calibrating Samples

2 In order to compare HMCMC and IF2 we need to set up a fair and theoretically
 3 justified way to select the number of samples to draw for the HMCMC iterations and
 4 the number of particles to use for IF2. As we wish to compare, among other things,
 5 approximate computational cost using runtimes, we need to determine how many
 6 sample draws for each method are required to obtain a certain accuracy. Sample
 7 draws are typically not comparable in terms of quality when considering multiple
 8 methods. For example, vanilla MCMC draws are computationally cheap compared
 9 to those from HMCMC, but HMCMC produces draws that more efficiently cover the
 10 sampling space. Thus we cannot just set the number of HMCMC draws equal to the
 11 number of particles used in IF2 – we must calibrate both quantities based on a desired
 12 target error. We assume that we are working with a problem that has an unknown
 13 real solution, so we use the Monte Carlo Standard Error (MCSE) [15].

14 Suppose we are using a Monte-Carlo based method to obtain a mean estimate $\hat{\mu}_n$ for
 15 a quantity μ , where n is the number of samples. Then the Law of Large Numbers
 16 says that $\hat{\mu}_n \rightarrow \mu$ as $n \rightarrow \infty$. Further, the Central Limit Theorem says that the error
 17 $\hat{\mu}_n - \mu$ should shrink with number of samples such that $\sqrt{n}(\hat{\mu}_n - \mu) \rightarrow \mathcal{N}(0, \sigma^2)$ as
 18 $n \rightarrow \infty$, where σ^2 is the variance of the samples drawn.

19 We of course do not know μ , but the above allows us to obtain an estimate $\hat{\sigma}_n$ for σ
 20 given a number of samples n as

$$21 \quad \hat{\sigma}_n = \sqrt{\frac{1}{n} \sum_{i=1}^n (X_i - \hat{\mu})^2}, \quad (4.3)$$

22 which is known as the Monte Carlo Standard Error.

23 We can modify this formula to account for multiple, potentially correlated, variables
 24 by replacing the single variance measure sum by

$$25 \quad \Theta^* V (\Theta^*)^T \quad (4.4)$$

26 where Θ^* is a row vector containing the reciprocals of the means of the parameters of
 27 interest, and V is the variance-covariance matrix with respect to the same parameters.
 28 This in effect scales the variances with respect to their magnitudes and accounts for
 29 covariation between parameters in one fell swoop. We also divide by the number of
 30 parameters, yielding

$$31 \quad \hat{\sigma}_n = \sqrt{\frac{1}{n} \frac{1}{P} \Theta^* V (\Theta^*)^T} \quad (4.5)$$

1 where P is the number of particles.

2 The goal here is to then pick the number of HMCMC samples and IF2 particles to
 3 yield similar MCSE values. To do this we picked a combination of parameters for
 4 RStan that yielded decent results when applied to the stochastic SIR model specified
 5 above, calculated the resulting mean MCSE across several model fits, and isolated the
 6 expected number of IF2 particles needed to obtain the same value. This was used as
 7 a starting value to “titrate” the IF2 iterations to the same point.

8 The resulting values were 1000 HMCMC warm-up iterations with 1000 samples drawn
 9 post-warm-up, and 2500 IF2 particles sent through 50 passes, each method giving an
 10 approximate MCSE of 0.0065.

11 4.3 IF2 Fitting

12 Now we will use an implementation of the IF2 algorithm to attempt to fit the stochas-
 13 tic SIR model to the previous data. The goal here is just parameter inference, but
 14 since IF2 works by applying a series of particle filters we essentially get the average
 15 system state estimates for a very small additional computational cost. Hence, we will
 16 will also look at that estimated behaviour in addition the parameter estimates.

17 The code used here is a mix of R and C++ implemented using Rcpp. The fitting
 18 was undertaken using 2500 particles with 50 IF2 passes and a cooling schedule given
 19 by a reduction in particle spread determined by 0.975^p , where p is the pass number
 20 starting with 0. This geometric cooling scheme is standard for use with IF2 [21][23],
 21 with the cooling rate chosen to neatly scale the perturbation factor from 1 to 0.02
 22 (almost 0) over 50 passes.

23 The MLE parameter estimates, taken to be the mean of the particle swarm values
 24 after the final pass, are shown in the table in Figure [4.4], along with the true values
 25 and the relative error.

26 From last IF2 particle filtering iteration, the mean state values from the particle
 27 swarm at each time step are shown with the true underlying state and data in the
 28 plot in Figure [4.5].

Name	True	IF2		HMCMC	
		Fit	Error	Fit	Error
R_0	3.0	3.27	9.08×10^{-2}	3.12	1.05×10^{-1}
r	10^{-1}	1.04×10^{-1}	3.61×10^{-2}	9.99×10^{-2}	-7.56×10^{-4}
Initial Infected	5	7.90	5.80×10^{-1}	6.64	3.28×10^{-1}
σ	10	8.84	-1.15×10^{-1}	8.5	-1.50×10^{-1}
η	5×10^{-1}	5.87×10^{-1}	1.73×10^{-1}	4.57×10^{-1}	-8.27×10^{-2}
ε_{err}	5×10^{-1}	1.63×10^{-1}	-6.73×10^{-1}	1.60×10^{-1}	-6.80×10^{-1}

Figure 4.4: Fitting errors.

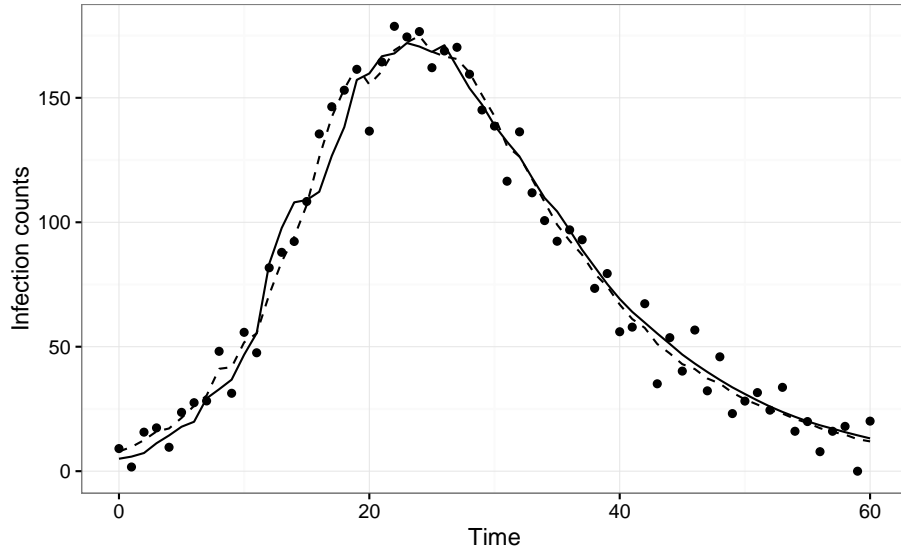


Figure 4.5: True system trajectory (solid line), observed data (dots), and IF2 estimated real state (dashed line).

1 4.4 IF2 Convergence

2 Since IF2 is an iterative algorithm where each pass through the data is expected to
3 push the parameter estimates towards the MLE, we can see the evolution of these es-
4 timates as a function of the pass number. We expect near-convergence in the param-
5 eter estimates as IF2 nears the maximum number of passes specified. Unconvincing
6 convergence plots may signal suboptimal algorithm parameters. If sensible algorithm
7 parameters have been chosen, we should see the convergence plots display “flattening”
8 over time.

9 Figure [4.6] shows evolution of the mean estimates for the six most critical paramete-
10 rs.

11 Figure [4.7] shows the evolution of the standard deviations of the parameter estimates
12 from the particle swarm as a function of the pass number. We should expect to see
13 asymptotic convergence to zero if the filter is converging.

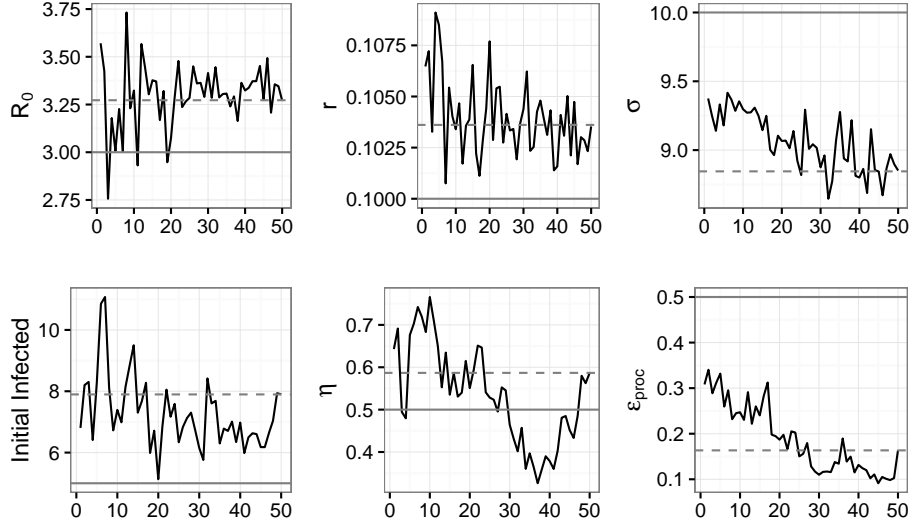


Figure 4.6: The horizontal axis shows the IF2 pass number. The solid black lines show the evolution of the ML estimates, the solid grey lines show the true value, and the dashed grey lines show the mean parameter estimates from the particle swarm after the final pass.

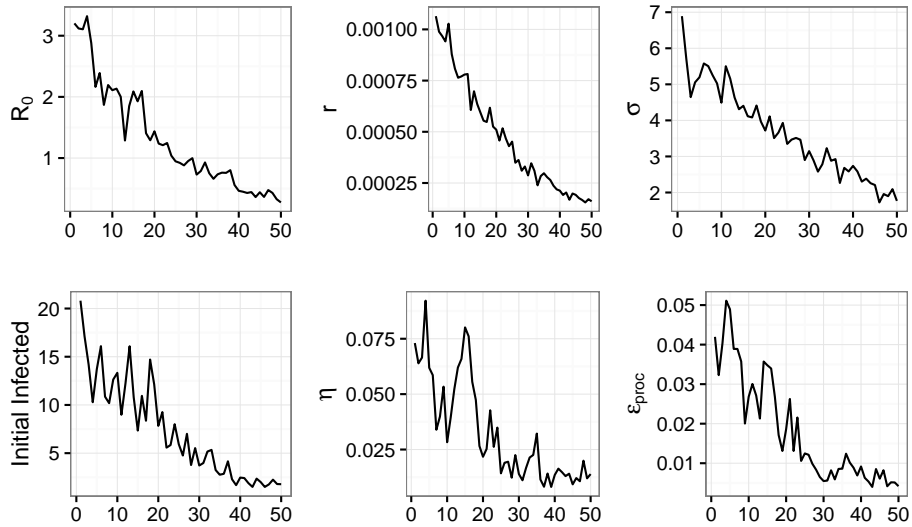


Figure 4.7: The horizontal axis shows the IF2 pass number and the solid black lines show the evolution of the standard deviations of the particle swarm values.

1 4.5 IF2 Densities

2 Of diagnostic importance are the densities of the parameter estimates given by the
3 final parameter swarm. If the swarm has collapsed, these densities will be extremely
4 narrow, almost resembling a vertical line. A “healthy” swarm should display relatively
5 smooth kernels of reasonable breadth.

6 Figure [4.8] shows the parameter sample distributions from the final parameter swarm.

7 The IF2 parameters chosen were in part chosen so as to not artificially narrow these
8 densities; a more aggressive cooling schedule and/or an increased number of passes
9 would have resulted in much narrower densities, and indeed have the potential to
10 collapse them to point estimates. This is undesirable as it may indicate instability –
11 the particles may have become “trapped” in a region of the sampling space.

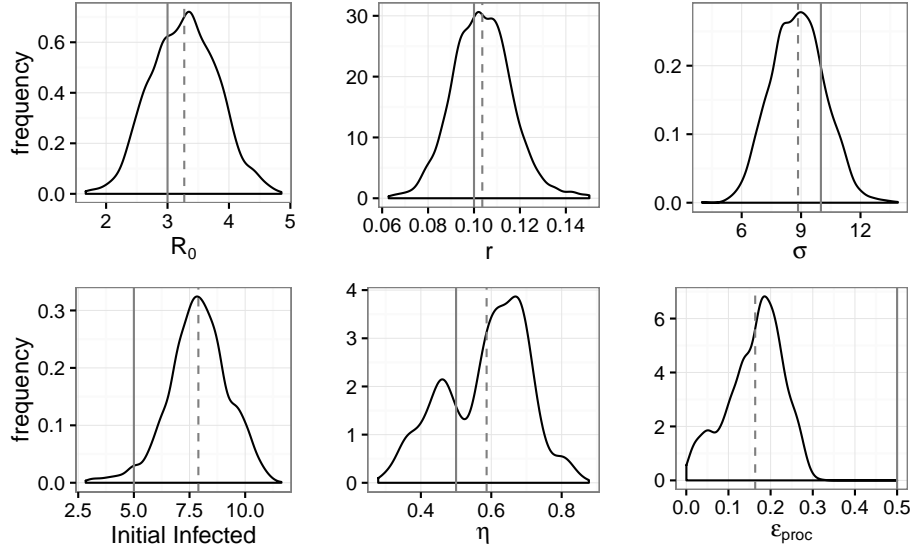


Figure 4.8: As before, the solid grey lines show the true parameter values and the dashed grey lines show the density means.

4.6 HMCMC Fitting

We can use the Hamiltonian Monte Carlo algorithm implemented in the ‘Rstan’ package to fit the stochastic SIR model as above. This was done with a single HMC chain of 2000 iterations with 1000 of those being warm-up iterations.

The MLE parameter estimates, taken to be the means of the samples in the chain, were shown in the table in Figure [4.4] along with the true values and relative error.

4.7 HMCMC Densities

Figure [4.9] shows the parameter estimation densities from the Stan HMCMC fitting.

the densities shown here represent a “true” MLE density estimate in that they represent HMC’s attempt to directly sample from the parameter space according to the likelihood surface, unlike IF2 which is in theory only trying to get a ML point estimate. Hence, these densities are potentially more robust than those produced by the IF2 implementation.

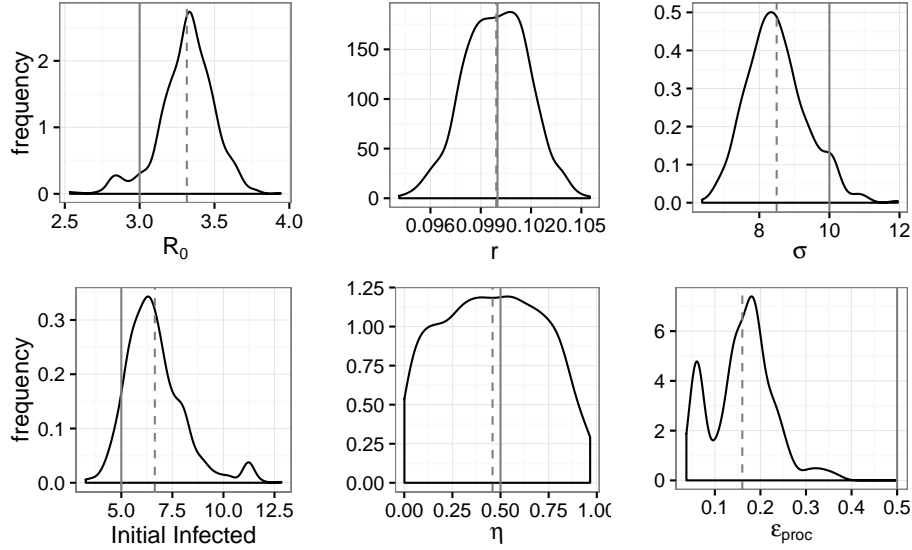


Figure 4.9: As before, the solid grey lines show the true parameter values and the dashed grey lines show the density means.

1 4.8 HMC MC and Bootstrapping

2 Unlike in some models, our RStan epidemic model does not keep track of state esti-
 3 mates directly, but does keep track of the autoregressive process latent variable draws,
 4 which allow us to reconstruct states. This was done to ease implementation as RStan
 5 places some restrictions on how interactions between parameters and states can be
 6 specified.

7 Figure [4.10] shows the results of 100 bootstrap trajectories generated from the RStan
 8 HMC MC samples.

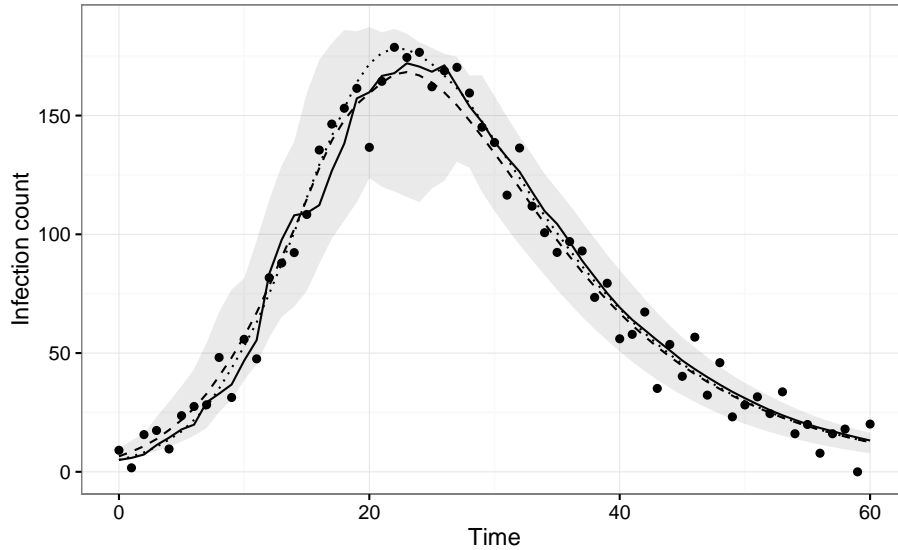


Figure 4.10: Result from 100 HMCBC bootstrap trajectories. The solid line shows the true states, the dots show the data, the dotted line shows the average system behaviour, the dashed line shows the bootstrap mean, and the grey ribbon shows the centre 95th quantile of the bootstrap trajectories.

1 4.9 Multi-trajectory Parameter Estimation

2 Here we fit the stochastic SIR model to 200 random independent trajectories using
 3 each method and examine the density of the point estimates produced.

4 Figure [4.11] shows the results of the mean parameter fits from IF2 and HMCBC for
 5 200 independent data sets generated using the previously described model.

6 The densities by and large display similar coverage, with the IF2 densities for r and
 7 ε_{proc} showing slightly wider coverage than the HMCBC densities for the same param-
 8 eters.

9 Figure [4.12] summarizes the running times for each algorithm.

10 The average running times were approximately 45.5 seconds and 257.4 seconds for IF2
 11 and HMCBC respectively, representing a 5.7x speedup for IF2 over HMCBC. While
 12 IF2 may be able to fit the model to data faster than HMCBC, we are obtaining less
 13 information; this will become important in the next section. Further, the results in
 14 Figure [4.12] show that while the running time for IF2 is relatively fixed, the times
 15 for HMCBC are anything but, showing a wide spread of potential times.

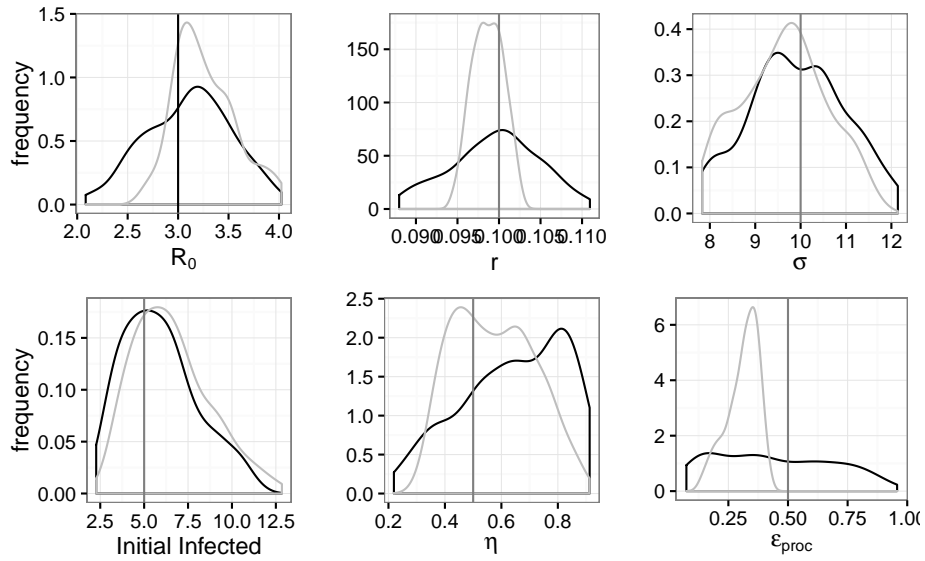


Figure 4.11: IF2 point estimate densities are shown in black and HMCMC point estimate densities are shown in grey. The vertical black lines show the true parameter values.

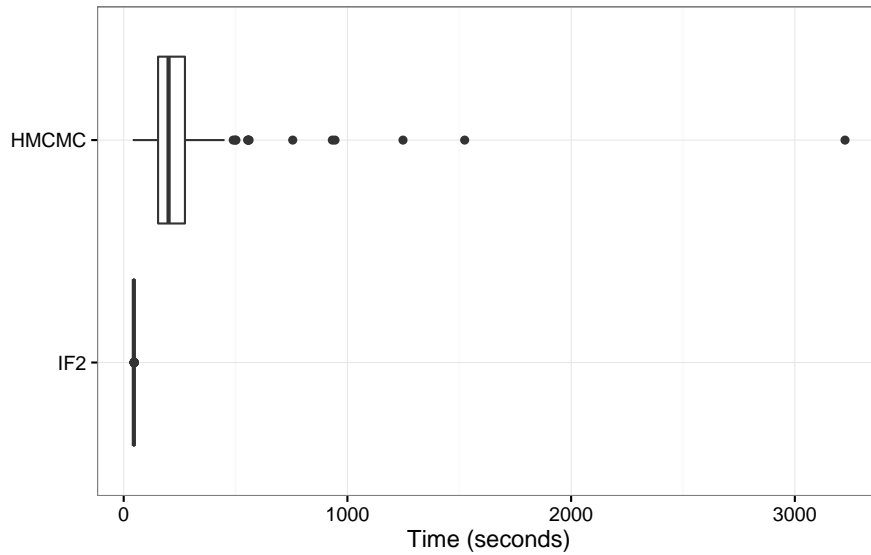


Figure 4.12: Fitting times for IF2 and HMCMC, in seconds. The centre box in each plot shows the centre 50th quantile, with the bold centre line showing the median.

Chapter 5

Forecasting Frameworks

5.1 Data Setup

This section will focus on taking the stochastic SIR model from the previous section, truncating the synthetic data output from realizations of that model, and seeing how well IF2 and HMCMC can reconstruct out-of-sample forecasts.

Figure [5.1] shows an example of a simulated system with truncated data.

In essence, we want to be able to give either IF2 or HMCMC only the data points and have it reconstruct the entirety of the true system states.

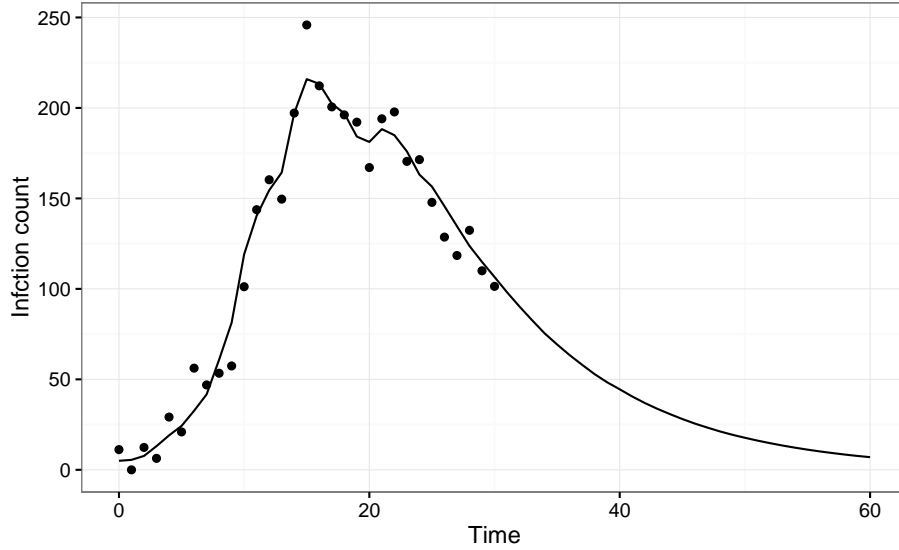


Figure 5.1: Infection count data truncated at $T = 30$. The solid line shows the true underlying system states, and the dots show those states with added observation noise. Parameters used were $R_0 = 3.0$, $r = 0.1$, $\eta = .05$, $\sigma_{proc} = 0.5$, and additive observation noise was drawn from $\mathcal{N}(0, 10)$.

1 5.2 IF2

2 For IF2, we will take advantage of the fact that the particle filter will produce state
 3 estimates for every datum in the time series given to it, as well as producing ML
 4 point estimates for the parameters. Both of these sources of information will be used
 5 to produce forecasts by parametric bootstrapping using the final parameter estimates
 6 from the particle swarm after the last IF2 pass, then using the newly generated
 7 parameter sets along with the system state point estimates from the first fitting to
 8 simulate the systems forward into the future.

9 We will truncate the data at half the original time series length (to $T = 30$), and fit
 10 the model as previously described.

11 Figure 2 shows [5.2] the state estimates for each time point produced by the last IF2
 12 pass.

13 Recall that IF2 is not trying to generate posterior probability densities, but rather
 14 produce a point estimate. Since we wish to determine the approximate distribution of
 15 each of the parameters in addition to the point estimate, we must add another layer
 16 atop the IF2 machinery, parametric bootstrapping.

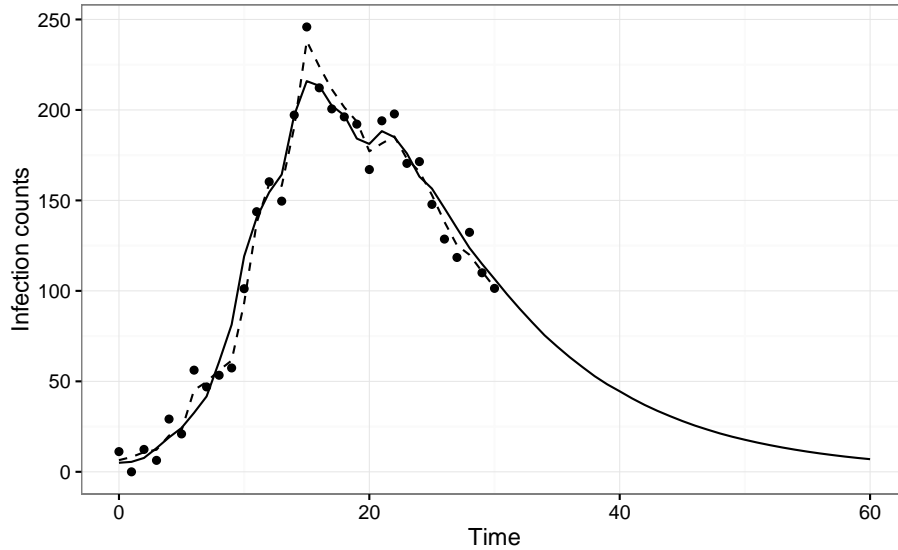


Figure 5.2: Infection count data truncated at $T = 30$ from Figure [5.1]. The dashed line shows IF2's attempt to reconstruct the true underlying state from the observed data points.

1 5.2.1 Parametric Bootstrapping

2 The goal of the parametric bootstrap is use an initial density sample θ^* to generate
 3 further samples $\theta_1, \theta_2, \dots, \theta_M$ from the sampling distribution of θ . It works by using θ to
 4 generate artificial data sets D_1, D_2, \dots, D_M to which we can refit our model of interest
 5 and generate new parameter sets. The literature suggests the most straightforward
 6 way of doing this is to fit the model to obtain θ^* , then use the model's forward
 7 simulator to generate new data sets, in essence treating our original estimate θ^* as
 8 the “truth” set.

9 An algorithm for parametric bootstrapping using IF2 and our stochastic SIR model
 10 is shown in Algorithm [5].

11 5.2.2 IF2 Forecasts

12 Using the parameter sets $\theta_1, \theta_2, \dots, \theta_M$ and the point estimate of the state provided by
 13 the initial IF2 fit, we can use use forward simulations from the last estimated state
 14 to produce estimates of the future state.

15 Figure [5.3] shows a projection of the data from the previous plots in Figures [5.1]
 16 and [5.2].

17 We can define a metric to gauge overall forecast effectiveness by calculating the SSE

Algorithm 5: Parametric Bootstrap

Input : Forward simulator $S(\theta)$, data set D

```

/* Initial fit */
1  $\theta^* \leftarrow IF2(D)$ 
/* Generate artificial data sets */
2 for  $i = 1 : M$  do
3    $D_i \leftarrow S(\theta^*)$ 
/* Fit to new data sets */
4 for  $i = 1 : M$  do
5    $\theta_i \leftarrow IF2(D_i)$ 

```

Output: Distribution samples $\theta_1, \theta_2, \dots, \theta_M$

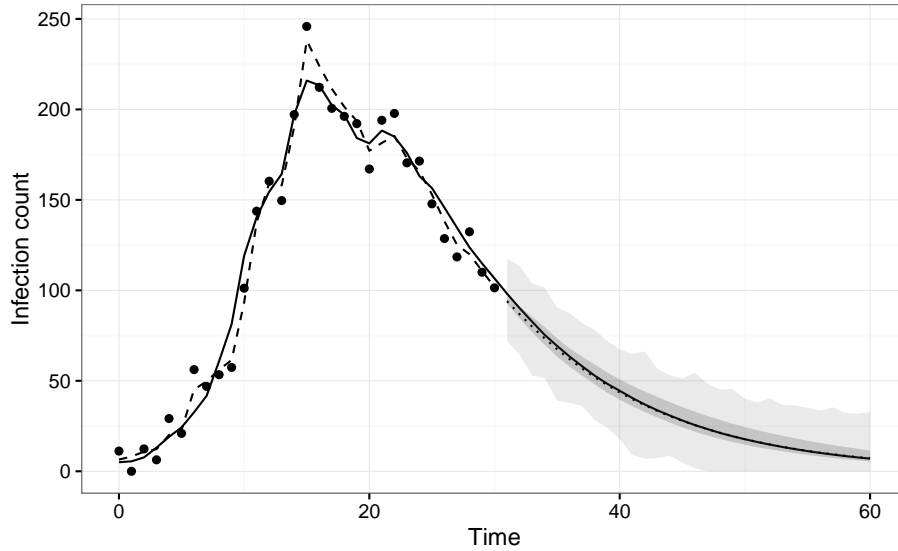


Figure 5.3: Forecast produced by the IF2 / parametric bootstrapping framework. The dotted line shows the mean estimate of the forecasts, the dark grey ribbon shows the 95% confidence interval based on the 0.025 and 0.975 quantiles on the true state estimates, and the lighter grey ribbon shows the same confidence interval on the true state estimates with added observation noise drawn from $\mathcal{N}(0, \sigma)$.

1 and dividing that value by the number of values predicted to get the average squared
2 error per point. For the data in Figure [5.3] the value was $\overline{SSE} = 1.67$. Normally
3 we would also want to address questions of forecast coverage, but this would require
4 at least a 100-fold increase in computational cost. This is potentially an avenue of
5 future investigation.

6 **5.3 HMC MC**

7 For HMC MC we can use a simpler approach to approach forecasting. We do not get
8 state estimates directly from the RStan fitting due to the way we implemented the
9 model, but we can construct them using the process noise latent variables as described
10 in Chapter 2. Once we've done this we can forward simulate the system from the state
11 estimate into the future.

12 Figure [5.4] shows the result of the HMC MC forecasting framework as applied to the
13 data from Figure [5.1].

14 And as before we can evaluate the averaged SSE of the forecast for the data shown,
15 giving $\overline{SSE} = 20.27$.

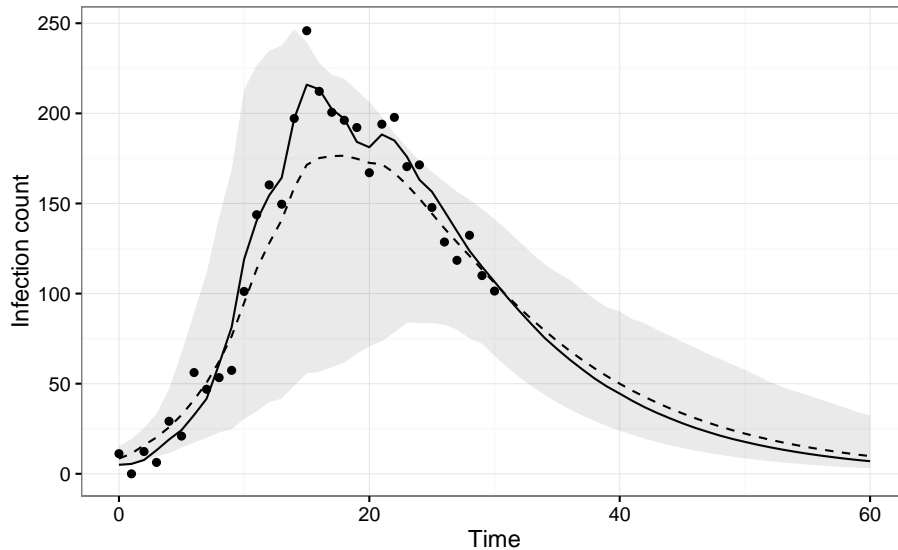


Figure 5.4: Forecast produced by the HMCMC / bootstrapping framework with $M = 200$ trajectories. The dotted line shows the mean estimate of the forecasts, and the grey ribbon shows the 95% confidence interval on the estimated true states as described in Figure [5.3].

1 5.4 Truncation vs. Error

2 Of course the above mini-comparison only shows one truncation value for one trajec-
 3 tory. Really, we need to know how each method performs on average given different
 4 trajectories and truncation amounts. In effect we wish to “starve” each method of data
 5 and see how poor the estimates become with each successive data point loss.

6 Using each method, we can fit the stochastic SIR model to successively smaller time
 7 series to see the effect of truncation on forecast averaged SSE. This was performed
 8 with 10 new trajectories drawn for each of the desired lengths. The results are shown
 9 in Figure [5.5].

10 IF2 and HMCMC perform very closely, with IF2 maintaining a small advantage up
 11 to a truncation of about 25-30 data points.

12 Since the parametric bootstrapping approach used by IF2 requires a significant num-
 13 ber of additional fits, its computational cost is significantly higher than the simpler
 14 bootstrapping approach used by the HMCMC framework, about 35.5x as expensive.
 15 However the now much longer running time can somewhat alleviated by parallelizing
 16 the parametric bootstrapping process; as each of the parametric bootstrap fittings
 17 in entirely independent, this can be done without a great deal of additional effort.

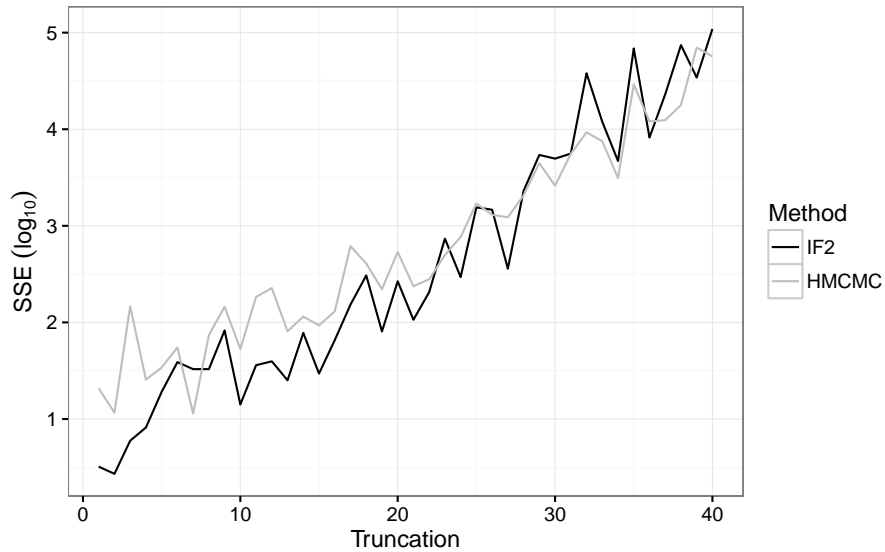


Figure 5.5: Error growth as a function of data truncation amount. Both methods used 200 bootstrap trajectories. Note that the y-axis shows the natural log of the averaged SSE, not the total SSE.

- 1 The code used here has this capability, but it was not utilised in the comparison so
- 2 as to accurately represent total computational cost, rather than potential running
- 3 time.

Chapter 6

S-map and SIRS

6.1 S-maps

A family of forecasting methods that shy away from the mechanistic model-based approaches outlined in the previous sections have been developed by George Sugihara and collaborators [35][36][18][13] over the last several decades. As these methods do not include a mechanistic model in their forecasting process, they also do not attempt to perform parameter estimation or inference. Instead they attempt to reconstruct the underlying dynamical process as a weighted linear model from a time series.

One such method, the sequential locally weighted global linear maps (S-map), builds a global linear map model and uses it to produce forecasts directly. Despite relying on a linear mapping, the S-map does not assume the time series on which it is operating is the product of linear system dynamics, and in fact was developed to accommodate non-linear dynamics. The linear component of the method only comes into play when combining forecast components together to produce a single estimate

The S-map works by first constructing a time series embedding of length E , known as the library and denoted $\{\mathbf{x}_i\}$. Consider a time series of length T denoted x_1, x_2, \dots, x_T . Each element in the time series with indices in the range $E, E + 1, \dots, T$ will have a corresponding entry in the library such that a given element x_t will correspond to a library vector of the form $\mathbf{x}_i = (x_t, x_{t-1}, \dots, x_{t-E+1})$. Next, given a forecast length L (representing L time steps into the future), each library vector \mathbf{x}_i is assigned a prediction from the time series $y_i = x_{t+L}$, where x_t is the first entry in \mathbf{x}_i . Finally, a forecast \hat{y}_t for specified predictor vector \mathbf{x}_t (usually from the library itself), is generated using an exponentially weighted function of the library $\{\mathbf{x}_i\}$, predictions $\{y_i\}$, and predictor vector \mathbf{x}_t .

This function is defined as follows:

1 First construct a matrix A and vector b defined as

$$\begin{aligned} A(i, j) &= w(\|\mathbf{x}_i - \mathbf{x}_t\|) \mathbf{x}_i(j) \\ b(i) &= w(\|\mathbf{x}_i - \mathbf{x}_t\|) y_i \end{aligned} \tag{6.1}$$

3 where $\|\cdot\|$ is the Euclidean norm, i ranges over 1 to the length of the library, and j
4 ranges over $[0, E]$. In the above equations and the ones that follow, we set $x_t(0) \equiv 1$
5 to account for the linear term in the map.

6 The weighting function w is defined as

$$w(d) = \exp\left(\frac{-\theta d}{\bar{d}}\right), \tag{6.2}$$

8 where d is the euclidean distance between the predictor vector and library vectors in
9 Equation [6.1] and \bar{d} is the average of these distances. We can then see that θ serves
10 as a way to specify the appropriate level of penalization applied to poorly-matching
11 library vectors – if θ is 0 all weights are the same (no penalization), and increasing θ
12 increases the level of penalization.

13 Now we solve the system $Ac = b$ to obtain the linear weightings used to generate the
14 forecast according to

$$\hat{y}_t = \sum_{j=0}^E c_t(j) \mathbf{x}_t(j). \tag{6.3}$$

16 In this way we have produced a forecast value for a single time. This process can
17 be repeated for a sequence of times $T + 1, T + 2, \dots$ to project a time series into the
18 future.

19 In essence what we are doing is generating a series of forecasts from every vector in
20 the library, weighting those forecasts based on how similar the corresponding library
21 vector is to our predictor vector, obtaining a solution to the system that maps com-
22 ponents of a predictor vector to its library vector's forecasted point (the mapping),
23 then applying that mapping to our predictor variable to obtain a forecast.

24 **6.2 S-map Algorithm**

25 The above description can be summarized in Algorithm [6].

Algorithm 6: S-map

```

/* Select a starting point */
Input : Time series  $x_1, x_2, \dots, x_T$ , embedding dimension  $E$ , distance
        penalization  $\theta$ , forecast length  $L$ , predictor vector  $\mathbf{x}_t$ 

/* Construct library  $\{\mathbf{x}_i\}$  */
1 for  $i = E : T$  do
2    $\mathbf{x}_i = (x_i, x_{i-1}, \dots, x_{i-E+1})$ 

/* Construct mapping from library vectors to predictions */
3 for  $i = 1 : (T_E + 1)$  do
4   for  $j = 1 : E$  do
5      $A(i, j) = w(\|\mathbf{x}_i - \mathbf{x}_t\|)\mathbf{x}_i(j)$ 
6 for  $i = 1 : (T_E + 1)$  do
7    $b(i) = w(\|\mathbf{x}_i - \mathbf{x}_t\|)y_i$ 

/* Use SVD to solve the mapping system,  $Ac = b$  */
8  $SVD(Ac = b)$ 

/* Compute forecast */
9  $\hat{y}_t = \sum_{j=0}^E c_t(j)\mathbf{x}_t(j)$ 

/* Forecasted value in time series */
Output: Forecast  $\hat{y}_t$ 

```

6.3 SIRS Model

In an epidemic or infectious disease context, the S-map algorithm will only really work on time series that appear cyclic. While there is nothing mechanically that prevents it from operating on a time series that do not appear cyclic, S-mapping requires a long time series in order to build a quality library. Without one the forecasting process would produce unreliable data.

Given, the S-map's data requirements, we need to specify a modified version of the SIR model. As IF2 and HMCMC in principle should be able operate on any reasonably well-specified model, the easiest way to compare the efficacy of S-mapping to IF2 or HMCMC is to generate data from a SIRS model with a seasonal component, and have all methods operate on the resulting time series.

The basic skeleton of the SIRS model is similar to the stochastic SIR model described previously, with one small addition. The deterministic ODE component of the model is as follows.

$$\begin{aligned}\frac{dS}{dt} &= -\Gamma(t)\beta SI + \eta R \\ \frac{dI}{dt} &= \Gamma(t)\beta SI - \gamma I \\ \frac{dR}{dt} &= \gamma I - \eta R,\end{aligned}\tag{6.4}$$

There are two new features here. We have a rate of waning immunity η through which people become able to be reinfected, and a seasonality factor function $\Gamma(t)$ defined as

$$\Gamma(t) = \exp \left[2 \left(\cos \left(\frac{2\pi}{365} t \right) - 1 \right) \right].\tag{6.5}$$

This function oscillates between 1 and e^{-4} (close to 0) and is meant to represent transmission damping during the off-season, for example summer for influenza. Further, it displays flatter troughs and sharper peaks to exaggerate its effect in peak season.

As before, β is allowed to walk restricted by a geometric mean, described by

$$\beta_{t+1} = \exp \left(\log(\beta_t) + \eta(\log(\bar{\beta}) - \log(\beta_t)) + \epsilon_t \right).\tag{6.6}$$

Figure [6.1] shows the SIRS model simulated for the equivalent of 5 years (260 weeks) and adding noise drawn from $\mathcal{N}(0, \sigma)$.




Figure 6.1: Five cycles generated by the SIRS function. The solid line the the true number of cases, dots show case counts with added observation noise. The parameter values were $R_0 = 3.0$, $\gamma = 0.1$, $\eta = 1$, $\sigma = 5$, and 10 initial cases.

1 Figure [6.2] shows how the S-map can reconstruct the next cycle in the time se-
 2 ries.

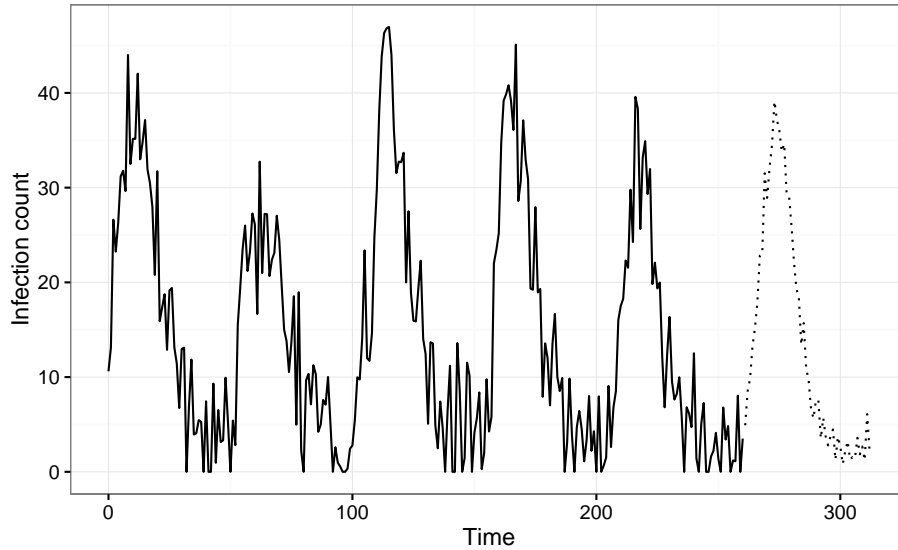


Figure 6.2: S-map applied to the data from the previous figure. The solid line shows the infection counts with observation noise from the previous plot, and the dotted line is the S-map forecast. Parameters chosen were $E = 14$ and $\theta = 3$.

- 1 The parameters used in the S-map algorithm to obtain the forecast used in Figure
- 2 [6.2] were obtained using a grid search of potential parameters outlined in [13]. The
- 3 script to perform this optimisation procedure is included in the appendices.

4 6.4 SIRS Model Forecasting

5 Naturally we wish to compare the efficacy of this comparatively simple technique
 6 against the more complex and more computationally taxing frameworks we have es-
 7 tablished to perform forecasting using IF2 and HMCMC.

8 To do this we generated a series of artificial time series of length 260 meant to represent
 9 5 years of weekly incidence counts and used each method to forecast up to 2 years into
 10 the future. Our goal here was to determine how forecast error changed with forecast
 11 length.

12 Figure [6.3] shows the results of the simulation.

13 Interestingly, all methods produce roughly the same result, which is to say the spikes
 14 in each outbreak cycle are difficult to accurately predict. IF2 produces better results
 15 than either HMCMC and the S-map for the majority of forecast lengths, with the

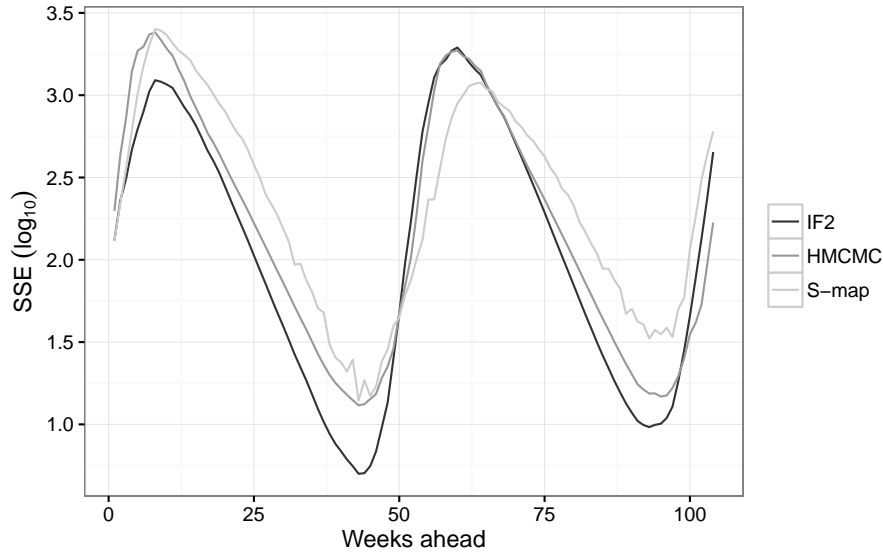


Figure 6.3: Error as a function of forecast length.

- 1 S-map producing the poorest results with the exception of the second rise in infection
- 2 rates where it outperforms the other methods.
- 3 While the S-map may not provide the same fidelity or forecast as IF2 or HMC MC, it
- 4 shines when it comes to running time. Figure [6.4] shows the running times over 20
- 5 simulations.
- 6 It is clear from Figure [6.4] that the S-map running times are minute compared to the
- 7 other methods, but to emphasize the degree: The average running time for the S-map
- 8 is about 0.15 seconds, for IF2 it is about 47,000, and for HMC MC it is about 9,200.
- 9 This is a speed-up of over 316,000x compared to IF2 and over 61,800x compared to
- 10 HMC MC.

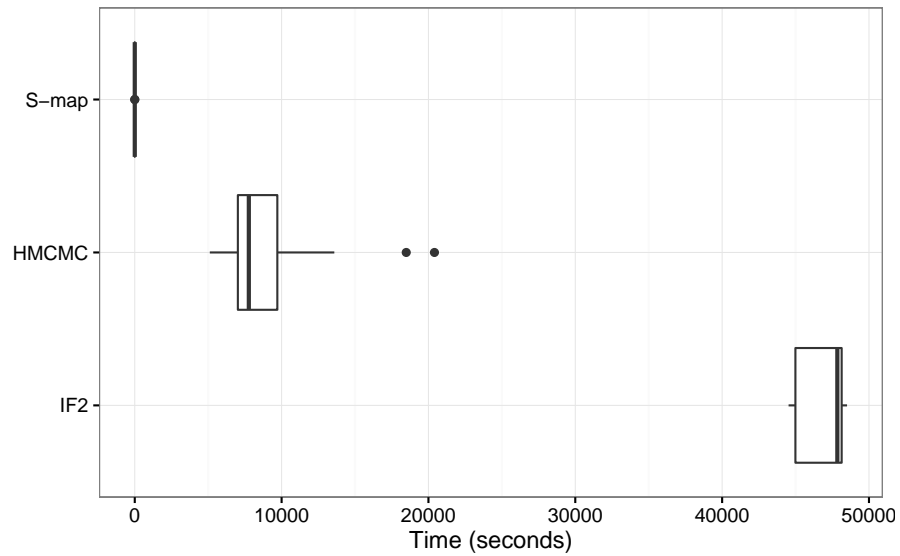


Figure 6.4: Runtimes for producing SIRS forecasts. The box shows the middle 50th quantile, the bold line is the median, and the dots are outliers. Note that these are not “true” outliers, simply ones outside a ranges based on the interquartile range.

Chapter 7

Spatial Epidemics

7.1 Spatial SIR

Spatial epidemic models provide a way to capture not just the temporal trend in an epidemic, but to also integrate spatial data and infer how the infection is spreading in both space and time. One such model we can use is a dynamic spatiotemporal SIR model.

We wish to construct a model build upon the stochastic SIR compartment model described previously but one that consists of several connected spatial locations, each with its own set of compartments. Consider a set of locations numbered $i = 1, \dots, N$, where N is the number of locations. Further, let N_i be the number of neighbours location i has. The model is then

$$\begin{aligned} \frac{dS_i}{dt} &= - \left(1 - \phi \frac{N_i}{N_i + 1}\right) \beta_i S_i I_i - \left(\frac{\phi}{N_i + 1}\right) S_i \sum_{j=0}^{N_i} \beta_j I_j \\ \frac{dI_i}{dt} &= \left(1 - \phi \frac{N_i}{N_i + 1}\right) \beta_i S_i I_i + \left(\frac{\phi}{N_i + 1}\right) S_i \sum_{j=0}^{N_i} \beta_j I_j - \gamma I_i \\ \frac{dR_i}{dt} &= \gamma I_i, \end{aligned} \tag{7.1}$$

Neighbours for a particular location are numbered $j = 1, \dots, N_i$. We have a new parameter, $\phi \in [0, 1]$, which is the degree of connectivity. If we let $\phi = 0$ we have total spatial isolation, and the dynamics reduce to the basic SIR model. If we let $\phi = 1$ then each of the neighbouring locations will have weight equivalent to the parent location.

As before we let β embark on a geometric random walk defined as




Figure 7.1: Evolution of a spatial epidemic in a ring topology. The outbreak was started with 5 cases in Location 2. Parameters were $R_0 = 3.0$, $\gamma = 0.1$, $\eta = 0.5$, $\sigma_{err} = 0.5$, and $\phi = 0.5$.

$$\beta_{i,t+1} = \exp \left(\log(\beta_{i,t}) + \eta(\log(\bar{\beta}) - \log(\beta_{i,t})) + \epsilon_t \right). \quad (7.2)$$

Note that as β is a state variable, each location has its own stochastic process driving the evolution of its β state.

If we imagine a circular topology in which each of 10 locations is connected to exactly two neighbours (i.e. location 1 is connected to locations N and 2, location 2 is connected to locations 1 and 3, etc.), and we start each location with completely susceptible populations except for a handful of infected individuals in one of the locations, we obtain a plot of the outbreak progression in Figure [7.1].

If we add noise to the data from Figure [7.1], we obtain Figure [7.2].

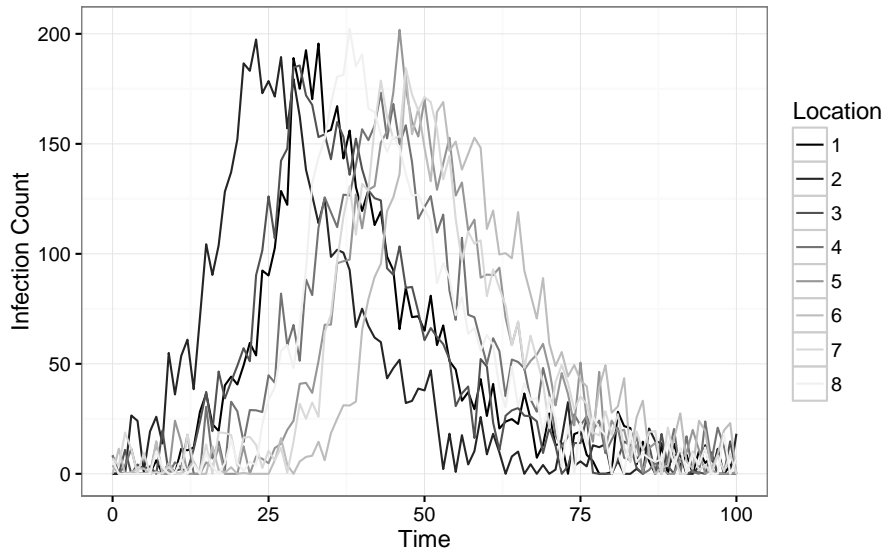


Figure 7.2: Evolution of a spatial epidemic as in Figure [7.1], with added observation noise drawn from $\mathcal{N}(0, 10)$.

7.2 Dewdrop Regression

Dewdrop regression [18] aims to overcome the primary disadvantage suffered by methods such as the S-map or its cousin Simplex Projection: the requirement of long time series from which to build a library. Suggested by Sugihara’s group in 2008, Dewdrop Regression works by stitching together shorter, related, time series, in order to give the S-map or similar methods enough data to operate on. The underlying idea is that as long as the underlying dynamics of the time series display similar behaviour (such as potentially collapsing to the same attractor), they can be treated as part of the same overarching system.

It is not enough to simply concatenate the shorter time series together – several procedures must be carried out and a few caveats observed. First, as the individual time series can be or drastically differing scales and breadths, they all must be rescaled to unit mean and variance. Then the library is constructed as before with an embedding dimension E , but any library vectors that span any of the seams joining the time series are discarded. Further, and predictions stemming from a library vector must stay within the time series from which they originated. In this way we are allowing the “shadow” of the underlying dynamics of the separate time series to infer the forecasts for segments of other time series. Once the library has been constructed, S-mapping can be carried out as previously specified.

This procedure is especially well-suited to the spatial model we are using. While the dynamics are stochastic, they still display very similar means and variances.

1 This means the rescaling process in Dewdrop Regression is not necessary and can
 2 be skipped. Further, the overall variation between the epidemic curves in each loca-
 3 tion is on the smaller side, meaning the S-map will have a high-quality library from
 4 which to build forecasts.

5 **7.3 Spatial Model Forecasting**

6 In order to compare the forecasting efficacy of Dewdrop Regression with S-mapping
 7 against IF2 and HMCMC, we generated 20 independent spatial data sets up to time
 8 $T = 50$ weeks in each of $L = 10$ locations and forecasted 10 weeks into the future.
 9 Forecasts were compared to that of the true model evolution, and the average SSE
 10 for each week ahead in the forecast were computed. The number of bootstrapping
 11 trajectories used by IF2 and HMCMC was reduced from 200 to 50 to curtail running
 12 times.

13 The results are shown in Figure [7.3].

14 The results show a clear delineation in forecast fidelity between methods. IF2 main-
 15 tains an advantage regardless of how long the forecast produced. Interestingly, Dew-
 16 drop Regression with S-mapping performs almost as well as IF2, and outperforms
 17 HMCMC. HMCMC lags behind both methods by a healthy margin.

18 If we examine the runtimes for each forecast framework, we obtain the data in Figure
 19 [7.4].

20 As before, the S-map with Dewdrop Regression runs faster than the other two methods
 21 with a huge margin. It is again hard to see exactly how large the margin is from the
 22 figure due to the scale, but we can examine the average values: the average running
 23 time for S-mapping with Dewdrop Regression was about 249 seconds, whereas the
 24 average times for IF2 and HMCMC were about 2.90×10^4 and 3.88×10^4 , respectively.
 25 This is a speed-up of just over 116x over IF2 and 156x over HMCMC.

26 Considering how well S-mapping performed with regards to forecast error, it shows a
 27 significant advantage over HMCMC in particular – it outperforms it in both forecast
 28 error and running times.

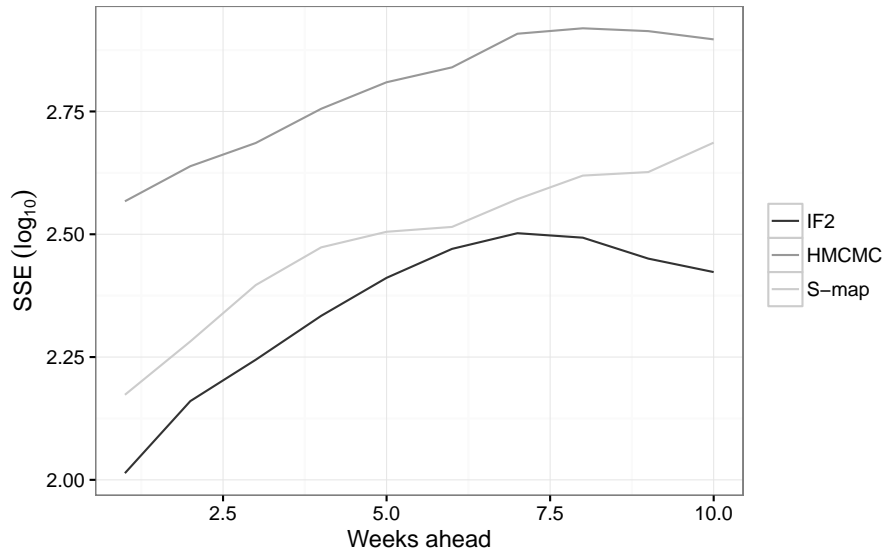


Figure 7.3: Average SSE (log scale) across each location and all trials as a function of the number of weeks ahead in the forecast.

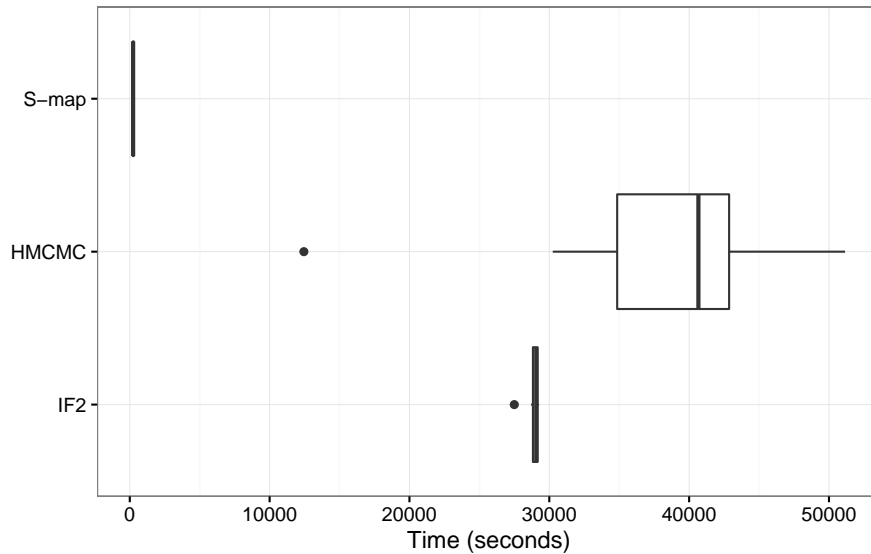


Figure 7.4: Runtimes for producing spatial SIR forecasts. The box shows the middle 50th quantile, the bold line is the median, and the dots are outliers.

Chapter 8

Discussion and Future Directions

A summary of the the results of forecasts in previous four chapters follows.

Immediately, we can see that the IF2 / parametric bootstrapping framework produces great results. This framework consistently out-performs both the HMCMC framework and S-mapping by itself or with Dewdrop Regression. This is not to say that the results produced by the other methods are poor, but rather that the ones produced by IF2 are noticeably better. This is true in every scenario we have explored here, and is particularly pronounced in the SIRS and spatial forecasting set-ups.

A surprise has been how well S-mapping has performed. Given the almost ludicrously shorter running times exhibited by S-mapping, it is almost shocking how well it performs. In the SIRS scenario it produces results only slightly less accurate than the other two methods, and is even the most accurate at predicting the rise to the second outbreak peak. In the spatial scenario it performs almost as well as IF2, and much better than HMCMC. The critical point here is that S-mapping, with its relative ease of implementation, efficiency, and accuracy would make a great “first-blush” forecasting tool that could be run and give a good prediction well before one would be able to even code the model specification for either of the other two methods. While S-mapping does require an up-front computational cost in “tuning” the two algorithm parameters, it is still negligible when compared to the costs incurred by IF2 and HMCMC.

8.1 Parallel and Distributed Computing

Whenever running times are discussed, we must consider the current computing landscape and hardware boundaries. In 1965, Intel co-founder Gordon E. Moore published a paper in which he observed that the number of transistors per unit area in integrated

1 circuits double roughly every year. The consequence of this growth is the approximate
2 year-over-year doubling of clock speeds (maximum number of sequential calculations
3 performed per second), equivalent to raw performance of the chip. This forecast was
4 updated in 1975 to double every 2 years and has held steady until the very recent
5 past [38].

6 Recently, transistor growth has begin to falter. This is due to several physical factors
7 preventing tighter “packing” of transistors into a single processor. To compensate
8 for these limitations, chip manufacturers have instead redesigned the internal chip
9 structures to consists to smaller “cores” within a single CPU die. The resulting
10 processing power per processor then stays on track with Moore’s Law, but keeps the
11 clock speeds of each individual core under control.

12 Of course this raises many problems on the software and algorithm side of computing.
13 Using several smaller cores instead of a single large one has the distinct disadvan-
14 tage of lack of cohesion – the cores must execute instructions completely decoupled
15 from each other. This means algorithms have to be redesigned, or at least rewritten
16 at the software level to consists of multiple independent pieces that can be run in
17 parallel. This practice is known as parallelization, and has become critical in taking
18 full advantage of machines of all scales – from mobile phones which overwhelmingly
19 favour multi-core CPU architectures, to large clusters and supercomputers which rely
20 on distributed computing “nodes”.

21 When working with computationally intensive algorithms, particularly iterative meth-
22 ods such those used in this paper, the question of parallelism naturally arises. It may
23 come as no surprise that the potential degrees of parallelism varies between meth-
24 ods.

25 Hamiltonian MCMC is cursed with high dependence between iterations. While HM-
26 CMC has an advantage over “vanilla” MCMC formulations in terms of efficiency of
27 step acceptance and ease of exploration of the parameter per number of samples, each
28 sample still depends entirely on the preceding one, and at a conceptual level the con-
29 struction of a Markov Chain *requires* iterative dependence. We cannot simply take
30 an accepted step, compute several proposed steps accept/reject them independently –
31 doing so would break the chain construction and could potentially bias our posterior
32 estimate to boot. We can, however, process multiple chains simultaneously and merge
33 the resulting samples, which has the added benefit of providing data from which to
34 assess convergence. If the required number of samples for a problem were large and
35 the required burn-in time were low, this methods could prove effective. However, the
36 parallel burn-in sampling is still inefficient as it is a duplication of effort with limited
37 pay-off – in the sense that the saved sample to discarded burn-in sample ratio would
38 not be as efficient as running a single long chain. Thus while parallelism via multiple
39 independent chains would help with a reduction in wall clock running times, it would
40 result in an *increase* in total computer time.

1 With regards to the bootstrapping process we used here, it should be clear that each
 2 bootstrap trajectory is completely independent, and thus this component of the fore-
 3 casting framework can be considered “embarrassingly” parallel or distributed. Un-
 4 fortunately, however, this is the least computationally demanding part of the process
 5 by several orders of magnitude, and so working to parallelize it would provide little
 6 advantage.

7 In the case of IF2, we have a decidedly different picture. In IF2 we have 5 primary
 8 steps in each data point integration:

- 9 • Forward evolution of the particles’ internal system state using their parameter
 10 state
- 11 • Weighting those state estimates against the data point using the observation
 12 function
- 13 • Particle weight normalizations
- 14 • Resampling from the particle weight distribution
- 15 • Particle parameter perturbations

16 Luckily, 4 of the 5 steps can be individually parallelized and run on a per-particle
 17 basis. The particle weight normalizations, however, cannot. Summation “reductions”
 18 are a well-known problem for parallel algorithms; they can be parallelized to a degree
 19 using binary reduction, but that only reduces the approximate running time from
 20 $\mathcal{O}(n)$ to $\mathcal{O}(\log(n))$. The normalization process requires the particles’ weight sum
 21 to be determined, hence the unavoidable obstacle of summation reductions rears its
 22 head. However this is in practice a less-taxing step, and its more demanding siblings
 23 are more amenable to parallelization.

24 Further, the full parametric bootstrapping process is highly computationally demand-
 25 ing, and also completely parallelizable. Each trajectory requires a fair bit of time to
 26 generate, on the order of of the original fitting time, and can be computed completely
 27 independently. Hence, IF2 is a very good candidate for a good parallel implementa-
 28 tion.

29 A future offshoot of this project would be a good parallel implementation of both the
 30 IF2 fitting process and the parametric bootstrapping framework. An ideal platform for
 31 this work would be NVIDIA’s Compute Unified Device Architecture (CUDA) Graph-
 32 ics Processing Unit (GPU) computing framework. While a CUDA implementation of
 33 a spatial epidemic IF2 parameter fitting algorithm was implemented, it lacked a good
 34 front-end implementation, R integration, and a parametric bootstrapping framework
 35 and so was not included in the main results of this paper. However, the code and
 36 some preliminary results are included in the appendices.

37 S-mapping, like the other two methods, is parallelizable to a degree. However, the

1 S-map is already a great deal faster than the other two methods, and in the worst case
 2 (paired with Dewdrop Regression and applied to a spatiotemporal data set) still only
 3 takes a few minutes to run. Setting this observation aside, if one were investing in
 4 developing a faster S-map implementation, this is certainly possible. By far the most
 5 computationally expensive component of the algorithm is the SVD decomposition, and
 6 algorithms exist to accelerate it via parallelization. Further, each point in the forecast
 7 can be computed separately; in the cases similar to the one here with application to
 8 spatiotemporal prediction, there can be a significant number of these points.

9 Further work developing parallel implementations of forecasting frameworks could be
 10 advantageous if the goal was to generate accurate forecasts under more stringent time
 11 limitations. IF2 seems to have emerged as a leader in forecast accuracy, if not in
 12 efficient running times, and demonstrates high potential for parallelism. Expansion
 13 of the CUDA IF2 (cuIF2) implementation to include a parallel bootstrapping layer
 14 and R integration could prove very promising.

15 **8.2 IF2, Bootstrapping, and Forecasting Method-** 16 **ology**

17 The parametric bootstrapping approach used to generate additional parameter pos-
 18 terior samples and produce forecasts has proven effective, but not necessarily compu-
 19 tationally efficient.

20 A recent paper utilising IF2 for forecasting [22] generated trajectories using IF2,
 21 parameter likelihood profiles, weighted quantiles, and the basic particle filter. The
 22 parameter profiles were used to construct a bounding box to search for good parameter
 23 sets, within which combinations of parameters to generate forecasts were selected
 24 using a Sobol sequence. Finally the forecasts were combined using a weighted quantile,
 25 taking into account the likelihood of the parameter sets used. Whether this approach
 26 would result in higher quality forecasts or lower running times is of interest, and could
 27 serve as a future research direction.

28 Expanding on this, there are other bootstrapping approaches that could be used to
 29 produce forecasts. A paper focusing solely on using IF2 with varied bootstrapping
 30 approaches and determining a forecast accuracy versus computational time trade-off
 31 curve of sorts would be useful, and would be another step towards establishing which
 32 tools are best for which jobs.

1 **8.3 Fin**

2 The overarching theme in this paper, from the theoretical considerations to the results
3 to the discussion, is that there still exists no “silver bullet” for forecasting problems.
4 Largely you can decide, as the user, how accurate you need your results to be, how
5 much computer time you have at your disposal, and how fast you need your results,
6 and select the method that best satisfies your needs. If speed is the priority, then you
7 can use S-mapping to get very quick and relatively accurate results. If you require
8 accuracy above all else, you must turn to heavier methods such as IF2, HMC MC,
9 and parametric bootstrapping in order to produce the cleanest forecast possible. And
10 this only represents three data points in a larger picture. There are a wide variety
11 of methods that are similar but not identical to methods explored here, each with
12 their own positive and negative attributes, their own advantages and disadvantages,
13 and that are ultimately likely to fill out our spectrum of methods more completely.
14 Thus future work should focus on attempting further direct comparison across a wider
15 swath of techniques, and implementing those techniques in a parallel fashion to take
16 advantage of the current and future landscape of high-performance computing.

Bibliography

- [1] Mohammad Ali et al. “Time Series Analysis of Cholera in Matlab, Bangladesh, during 1988-2001”. In: *Journal of health, population, ...* 31.1 (2013), pp. 11–19. URL: <http://www.ncbi.nlm.nih.gov/pmc/articles/PMC3702354/>.
- [2] Christophe Andrieu et al. “An introduction to MCMC for machine learning”. In: *Machine Learning* 50.1-2 (2003), pp. 5–43. ISSN: 08856125. DOI: 10.1023/A:1020281327116. arXiv: 1109.4435v1.
- [3] M.S. Arulampalam et al. “A tutorial on particle filters for online nonlinear/non-Gaussian Bayesian tracking”. In: *IEEE Transactions on Signal Processing* 50.2 (2002), pp. 174–188. ISSN: 1053587X. DOI: 10.1109/78.978374. URL: <http://ieeexplore.ieee.org/lpdocs/epic03/wrapper.htm?arnumber=978374>.
- [4] Jacob Bock Axelsen et al. “Multiannual forecasting of seasonal influenza dynamics reveals climatic and evolutionary drivers.” In: *Proceedings of the National Academy of Sciences of the United States of America* 111.26 (July 2014), pp. 9538–42. ISSN: 1091-6490. DOI: 10.1073/pnas.1321656111. URL: <http://www.pubmedcentral.nih.gov/articlerender.fcgi?artid=4084473%7B%5C%26%7Dttool=pmcentrez%7B%5C%26%7Drendertype=abstract>.
- [5] M Babyak. “What You See May Not Be What You Get: A Brief, Nontechnical Introduction to Overfitting in Regression Type Models”. In: *J. Bio. Med.* 66.3 (2004), pp. 411–421. ISSN: 0033-3174. DOI: 10.1097/01.psy.0000127692.23278.a9.
- [6] Thomas Bengtsson, Peter Bickel, and Bo Li. “Curse-of-dimensionality revisited: Collapse of the particle filter in very large scale systems”. In: *Probability and Statistics* 2 (2008), pp. 316–334. DOI: 10.1214/193940307000000518. arXiv: 0805.3034. URL: <http://arxiv.org/abs/0805.3034>.
- [7] a. Camacho et al. “Explaining rapid reinfections in multiple-wave influenza outbreaks: Tristan da Cunha 1971 epidemic as a case study”. In: *Proceedings of the Royal Society B: Biological Sciences* 278 (2011), pp. 3635–3643. ISSN: 0962-8452. DOI: 10.1098/rspb.2011.0300.

- 1 [8] Bob Carpenter et al. “Stan: A Probabilistic Programming Language”. In: *Journal of Statistical Software* (2016). URL: <http://www.stat.columbia.edu/%7B%7Dgelman/research/unpublished/stan-resubmit-JSS1293.pdf>.
- 2
- 3
- 4 [9] Jean-Paul Chretien et al. “Influenza forecasting in human populations: a scoping review.” In: *PloS one* 9.4 (Jan. 2014), e94130. ISSN: 1932-6203. DOI: 10.1371/journal.pone.0094130. URL: <http://www.pubmedcentral.nih.gov/articlerender.fcgi?artid=3979760%7B%5C%7Dtool=pmcentrez%7B%5C%7Drendertype=abstract>.
- 5
- 6
- 7
- 8
- 9 [10] Samantha Cook et al. “Assessing Google Flu trends performance in the United States during the 2009 influenza virus A (H1N1) pandemic”. In: *PLoS ONE* 6.8 (2011), pp. 1–8. ISSN: 19326203. DOI: 10.1371/journal.pone.0023610.
- 10
- 11
- 12 [11] Andrea Freyer Dugas et al. “Influenza forecasting with Google Flu Trends.” In: *PloS one* 8.2 (Jan. 2013), e56176. ISSN: 1932-6203. DOI: 10.1371/journal.pone.0056176. URL: <http://www.pubmedcentral.nih.gov/articlerender.fcgi?artid=3572967%7B%5C%7Dtool=pmcentrez%7B%5C%7Drendertype=abstract>.
- 13
- 14
- 15
- 16 [12] Dirk Eddelbuettel and Romain Fran. “Rcpp: Seamless R and C ++ Integration”. In: *Journal Of Statistical Software* 40.8 (2011), pp. 1–18. ISSN: 15487660. DOI: 10.1007/978-1-4614-6868-4. arXiv: arXiv:1011.1669v3. URL: <http://www.jstatsoft.org/v40/i08/>.
- 17
- 18
- 19
- 20 [13] Sarah M. Glaser, Hao Ye, and George Sugihara. “A nonlinear, low data requirement model for producing spatially explicit fishery forecasts”. In: *Fisheries Oceanography* 23 (2014), pp. 45–53. ISSN: 10546006. DOI: 10.1111/fog.12042.
- 21
- 22
- 23 [14] Andrea L Graham et al. “Explaining rapid reinfections in multiple-wave influenza outbreaks : Tristan da Cunha 1971 epidemic as a case study”. In: *Proc. R. Soc. B* (2016). DOI: 10.1098/rspb.2011.0300.
- 24
- 25
- 26 [15] Bradley Harding. “Standard errors : A review and evaluation of standard error estimators using Monte Carlo simulations”. In: (2014), pp. 107–123.
- 27
- 28 [16] Florian Hartig et al. “Statistical inference for stochastic simulation models - theory and application”. In: *Ecology Letters* 14.8 (2011), pp. 816–827. ISSN: 1461023X. DOI: 10.1111/j.1461-0248.2011.01640.x.
- 29
- 30
- 31 [17] Matthew D. Hoffman and Andrew Gelman. “The No-U-Turn Sampler: Adaptively Setting Path Lengths in Hamiltonian Monte Carlo”. In: *Journal of Machine Learning Research* 15.April (2014), pp. 1593–1623. arXiv: 1111.4246. URL: <http://mc-stan.org/>.
- 32
- 33
- 34
- 35 [18] Chih-hao Hsieh, Christian Anderson, and George Sugihara. “Extending nonlinear analysis to short ecological time series.” In: *The American naturalist* 171.1 (2008), pp. 71–80. ISSN: 0003-0147. DOI: 10.1086/524202.
- 36
- 37

- 1 [19] E L Ionides, C Bret, and A. A. King. “Inference for nonlinear dynamical sys-
 2 tems.” In: *Proceedings of the National Academy of Sciences of the United States*
 3 *of America* 103.49 (Dec. 2006), pp. 18438–43. ISSN: 0027-8424. DOI: 10.1073/
 4 pnas.0603181103. URL: [http://www.pubmedcentral.nih.gov/articlerender.
 5 fcgi?artid=3020138%7B%5C%7Dtool=pmcentrez%7B%5C%7Drendertype=
 6 abstract](http://www.pubmedcentral.nih.gov/articlerender.fcgi?artid=3020138%7B%5C%7Dtool=pmcentrez%7B%5C%7Drendertype=abstract).
- 7 [20] Edward L. Ionides et al. “Inference for dynamic and latent variable models
 8 via iterated, perturbed Bayes maps”. In: *Proceedings of the National Academy*
 9 *of Sciences* 112.3 (2015), pp. 719–724. ISSN: 0027-8424. DOI: 10.1073/pnas.
 10 1410597112. URL: [http://www.pnas.org/lookup/doi/10.1073/pnas.
 11 1410597112](http://www.pnas.org/lookup/doi/10.1073/pnas.1410597112).
- 12 [21] Aaron A King, Dao Nguyen, and Edward L. Ionides. “Statistical Inference for
 13 Partially Observed Markov Processes via the R Package pomp”. In: *Journal of*
 14 *Statistical Software* 59.10 (2015). arXiv: arXiv:1509.00503v1.
- 15 [22] Aaron A King et al. “Avoidable errors in the modelling of outbreaks of emerg-
 16 ing pathogens, with special reference to Ebola”. In: *Proceedings of the Royal*
 17 *Society B: Biological Sciences* 282.1806 (2015), pp. 20150347–20150347. ISSN:
 18 0962-8452. DOI: 10.1098/rspb.2015.0347. arXiv: 1412.0968. URL: [http:
 19 //rspb.royalsocietypublishing.org/cgi/doi/10.1098/rspb.2015.0347](http://rspb.royalsocietypublishing.org/cgi/doi/10.1098/rspb.2015.0347).
- 20 [23] Aaron A King et al. *pomp: Statistical inference for partially observed Markov*
 21 *processes*. 2016. URL: <http://kingaa.github.io/pomp>.
- 22 [24] William A. Link and Mitchell J. Eaton. “On thinning of chains in MCMC”.
 23 In: *Methods in Ecology and Evolution* 3.1 (2012), pp. 112–115. ISSN: 2041210X.
 24 DOI: 10.1111/j.2041-210X.2011.00131.x.
- 25 [25] Kanti V. Mardia et al. “The Kriged Kalman filter”. In: *Test* 7.2 (1998), pp. 217–
 26 282. ISSN: 11330686. DOI: 10.1007/BF02565111.
- 27 [26] Radford M Neal. “Handbook of Markov Chain Monte Carlo”. In: *Handbook*
 28 *of Markov Chain Monte Carlo* 20116022 (2011), pp. 113–162. DOI: 10.1201/
 29 b10905. arXiv: arXiv:1206.1901v1. URL: [http://www.crcnetbase.com/doi/
 30 book/10.1201/b10905](http://www.crcnetbase.com/doi/book/10.1201/b10905).
- 31 [27] Elaine O Nsoesie et al. “A systematic review of studies on forecasting the dy-
 32 namics of influenza outbreaks.” In: *Influenza and other respiratory viruses* 8.3
 33 (May 2014), pp. 309–16. ISSN: 1750-2659. DOI: 10.1111/irv.12226. URL: [http:
 34 //www.pubmedcentral.nih.gov/articlerender.fcgi?artid=4181479%7B%5C%
 35 %7Dtool=pmcentrez%7B%5C%7Drendertype=abstract](http://www.pubmedcentral.nih.gov/articlerender.fcgi?artid=4181479%7B%5C%7Dtool=pmcentrez%7B%5C%7Drendertype=abstract).
- 36 [28] Elaine Nsoesie, Madhav Marathe, and John Brownstein. “Forecasting peaks of
 37 seasonal influenza epidemics.” In: *PLoS currents* 5 (Jan. 2013), pp. 1–14. ISSN:
 38 2157-3999. DOI: 10.1371/currents.outbreaks.bb1e879a23137022ea79a8c508b030bc.
 39 URL: [http://www.pubmedcentral.nih.gov/articlerender.fcgi?artid=
 40 3712489%7B%5C%7Dtool=pmcentrez%7B%5C%7Drendertype=abstract](http://www.pubmedcentral.nih.gov/articlerender.fcgi?artid=3712489%7B%5C%7Dtool=pmcentrez%7B%5C%7Drendertype=abstract).

- 1 [29] Robert C Reiner et al. “Highly localized sensitivity to climate forcing drives
2 endemic cholera in a megacity.” In: *Proceedings of the National Academy of*
3 *Sciences of the United States of America* 109.6 (Feb. 2012), pp. 2033–6. ISSN:
4 1091-6490. DOI: 10.1073/pnas.1108438109. URL: [http://www.pubmedcentral.](http://www.pubmedcentral.nih.gov/articlerender.fcgi?artid=3277579%7B%5C%7Dtool=pmcentrez%7B%5C%7Drendertype=abstract)
5 [nih.gov/articlerender.fcgi?artid=3277579%7B%5C%7Dtool=pmcentrez%](http://www.pubmedcentral.nih.gov/articlerender.fcgi?artid=3277579%7B%5C%7Dtool=pmcentrez%7B%5C%7Drendertype=abstract)
6 [7B%5C%7Drendertype=abstract](http://www.pubmedcentral.nih.gov/articlerender.fcgi?artid=3277579%7B%5C%7Dtool=pmcentrez%7B%5C%7Drendertype=abstract).
- 7 [30] Sujit K Sahu. “A Bayesian Kriged-Kalman model for short-term fore- casting of
8 air pollution levels”. In: *Appl. Statist* 54 (2005), pp. 223–244. ISSN: 0035-9254.
9 DOI: 10.1111/j.1467-9876.2005.00480.x.
- 10 [31] Jacques Sau et al. “Particle filter-based real-time estimation and prediction of
11 traffic conditions Modeling framework”. In: *Information Systems* (1918), pp. 1–
12 8.
- 13 [32] Jeffrey Shaman and Alicia Karspeck. “Forecasting seasonal outbreaks of in-
14 fluenza.” In: *Proceedings of the National Academy of Sciences of the United*
15 *States of America* 109.50 (Dec. 2012), pp. 20425–30. ISSN: 1091-6490. DOI:
16 10.1073/pnas.1208772109. URL: [http://www.pubmedcentral.nih.gov/](http://www.pubmedcentral.nih.gov/articlerender.fcgi?artid=3528592%7B%5C%7Dtool=pmcentrez%7B%5C%7Drendertype=abstract)
17 [articlerender.fcgi?artid=3528592%7B%5C%7Dtool=pmcentrez%7B%5C%](http://www.pubmedcentral.nih.gov/articlerender.fcgi?artid=3528592%7B%5C%7Dtool=pmcentrez%7B%5C%7Drendertype=abstract)
18 [7Drendertype=abstract](http://www.pubmedcentral.nih.gov/articlerender.fcgi?artid=3528592%7B%5C%7Dtool=pmcentrez%7B%5C%7Drendertype=abstract).
- 19 [33] Jeffrey Shaman, Wan Yang, and Sasikiran Kandula. “Inference and Forecast of
20 the Current West African Ebola Outbreak in Guinea, Sierra Leone and Liberia”.
21 In: *PLoS Currents* 6 (2014), ecurrents.outbreaks.3408774290b1a0f2dd7cae877c8b8f.
22 ISSN: 2157-3999. DOI: 10.1371/currents.outbreaks.3408774290b1a0f2dd7cae877c8b8ff6.
- 23 [34] Stan Development Team. “Stan Modeling Language Users Guide and Reference
24 Manual”. In: 2.9.0 (2015).
- 25 [35] G Sugihara and R M May. *Nonlinear forecasting as a way of distinguishing*
26 *chaos from measurement error in time series*. 1990. DOI: 10.1038/344734a0.
- 27 [36] George Sugihara. “Nonlinear Forecasting for the Classification of Natural Time
28 Series”. In: *Philosophical Transactions of the Royal Society A: Mathematical,*
29 *Physical and Engineering Sciences* 348 (1994), pp. 477–495. ISSN: 1364-503X.
30 DOI: 10.1098/rsta.1994.0106.
- 31 [37] Carmen L Vidal Rodeiro and Andrew B Lawson. “Online updating of space-
32 time disease surveillance models via particle filters.” In: *Statistical methods in*
33 *medical research* 15.5 (2006), pp. 423–444. ISSN: 0962-2802. DOI: 10.1177 /
34 0962280206071640.
- 35 [38] Mitchell Waldrop. “More then Moore”. In: *Nature* 530.11. February (2016),
36 p. 145.

- 1 [39] Wan Yang, Alicia Karspeck, and Jeffrey Shaman. “Comparison of filtering meth-
2 ods for the modeling and retrospective forecasting of influenza epidemics.” In:
3 *PLoS computational biology* 10.4 (Apr. 2014), e1003583. ISSN: 1553-7358. DOI:
4 10.1371/journal.pcbi.1003583. URL: [http://www.pubmedcentral.nih.gov/
5 articlerender.fcgi?artid=3998879%7B%5C%7Dtool=pmcentrez%7B%5C%
6 %7Drendertype=abstract](http://www.pubmedcentral.nih.gov/articlerender.fcgi?artid=3998879%7B%5C%7Dtool=pmcentrez%7B%5C%7Drendertype=abstract).
- 7 [40] Xingyu Zhang et al. “Comparative study of four time series methods in fore-
8 casting typhoid fever incidence in China.” In: *PloS one* 8.5 (Jan. 2013), e63116.
9 ISSN: 1932-6203. DOI: 10.1371/journal.pone.0063116. URL: [http://www.
10 pubmedcentral.nih.gov/articlerender.fcgi?artid=3641111%7B%5C%
11 %7Dtool=pmcentrez%7B%5C%7Drendertype=abstract](http://www.pubmedcentral.nih.gov/articlerender.fcgi?artid=3641111%7B%5C%7Dtool=pmcentrez%7B%5C%7Drendertype=abstract).

1 Appendix A

2 Hamiltonian MCMC

3 A.1 Full R code

4 This code will run all the indicated analysis and produce all plots.

```
5  
6 1 ## Dexter Barrows  
7 2 ## McMaster University  
8 3 ## 2016  
9 4  
10 5 library(deSolve)  
11 6 library(rstan)  
12 7 library(shinystan)  
13 8 library(ggplot2)  
14 9 library(RColorBrewer)  
15 10 library(reshape2)  
16 11  
17 12 SIR ← function(Time, State, Pars) {  
18 13  
19 14     with(as.list(c(State, Pars)), {  
20 15  
21 16         B ← R0*r/N  
22 17         BSI ← B*S*I  
23 18         rI ← r*I  
24 19  
25 20         dS = -BSI  
26 21         dI = BSI - rI  
27 22         dR = rI  
28 23  
29 24         return(list(c(dS, dI, dR)))  
30 25  
31 26     })  
32 27  
33 28 }  
34 29
```



```

1 30 pars <- c(R0 <- 3.0,      # average number of new infected individuals
2      per infectious person
3 31      r <- 0.1,          # recovery rate
4 32      N <- 500)          # population size
5 33
6 34 T <- 100
7 35 y_ini <- c(S = 495, I = 5, R = 0)
8 36 times <- seq(0, T, by = 1)
9 37
10 38 odeout <- ode(y_ini, times, SIR, pars)
11 39
12 40 set.seed(1001)
13 41 sigma <- 10
14 42 infec_counts_raw <- odeout[,3] + rnorm(T+1, 0, sigma)
15 43 infec_counts <- ifelse(infec_counts_raw < 0, 0, infec_counts_raw)
16 44
17 45 g <- qplot(0:T, odeout[,3], geom = "line", xlab = "Time (weeks)", ylab
18      = "Infection Count") +
19 46      geom_point(aes(y = infec_counts)) +
20 47      theme_bw()
21 48
22 49 print(g)
23 50 ggsave(g, filename="dataplot.pdf", height=4, width=6.5)
24 51
25 52 sPw <- 7
26 53 datlen <- (T-1)*7 + 1
27 54
28 55 data <- matrix(data = -1, nrow = T+1, ncol = sPw)
29 56 data[,1] <- infec_counts
30 57 standata <- as.vector(t(data))[1:datlen]
31 58
32 59 sir_data <- list( T = datlen,      # simulation time
33 60                  y = standata,    # infection count data
34 61                  N = 500,         # population size
35 62                  h = 1/sPw )      # step size per day
36 63
37 64 rstan_options(auto_write = TRUE)
38 65 options(mc.cores = parallel::detectCores())
39 66 stan_options <- list( chains = 4,      # number of chains
40 67                      iter  = 2000,   # iterations per chain
41 68                      warmup = 1000,   # warmup iterations
42 69                      thin   = 2)      # thinning number
43 70 fit <- stan(file = "d_siode_euler.stan",
44 71            data   = sir_data,
45 72            chains  = stan_options$chains,
46 73            iter    = stan_options$iter,
47 74            warmup  = stan_options$warmup,
48 75            thin    = stan_options$thin )
49 76
50 77 exfit <- extract(fit, permuted = TRUE, inc_warmup = FALSE)
51 78

```

```

1 79 R0points <- exfit$R0
2 80 R0kernel <- qplot(R0points, geom = "density", xlab = expression(R[0]),
3      ylab = "frequency") +
4 81     geom_vline(aes(xintercept=R0), linetype="dashed", size=1,
5      color="grey50") +
6 82     theme_bw()
7 83
8 84 print(R0kernel)
9 85 ggsave(R0kernel, filename="kernelR0.pdf", height=3, width=3.25)
10 86
11 87 rpoints <- exfit$r
12 88 rkernell <- qplot(rpoints, geom = "density", xlab = "r", ylab = "
13     frequency") +
14 89     geom_vline(aes(xintercept=r), linetype="dashed", size=1,
15     color="grey50") +
16 90     theme_bw()
17 91
18 92 print(rkernell)
19 93 ggsave(rkernell, filename="kernelr.pdf", height=3, width=3.25)
20 94
21 95 sigmapoints <- exfit$sigma
22 96 sigmakernell <- qplot(sigmapoints, geom = "density", xlab = expression(
23     sigma), ylab = "frequency") +
24 97     geom_vline(aes(xintercept=sigma), linetype="dashed", size=1,
25     color="grey50") +
26 98     theme_bw()
27 99
28 100 print(sigmakernell)
29 101 ggsave(sigmakernell, filename="kernelsigma.pdf", height=3, width=3.25)
30 102
31 103 infecpoints <- exfit$y0[,2]
32 104 infeckernell <- qplot(infecpoints, geom = "density", xlab = "Initial
33     Infected", ylab = "frequency") +
34 105     geom_vline(aes(xintercept=y_ini[['I']]), linetype="dashed",
35     size=1, color="grey50") +
36 106     theme_bw()
37 107
38 108 print(infeckernell)
39 109 ggsave(infeckernell, filename="kernelinfec.pdf", height=3, width=3.25)
40 110
41 111 exfit <- extract(fit, permuted = FALSE, inc_warmup = FALSE)
42 112 plotdata <- melt(exfit[,,"R0"])
43 113 tracefitR0 <- ggplot() +
44 114     geom_line(data = plotdata,
45     aes(x = iterations,
46     y = value,
47     color = factor(chains, labels = 1:stan_
48     options$chains))) +
49 118     labs(x = "Sample", y = expression(R[0]), color = "Chain
50     ") +
51 119     scale_color_brewer(palette="Greys") +

```

```

1 120         theme_bw()
2 121
3 122 print(tracefitR0)
4 123 ggsave(tracefitR0, filename="traceplotR0.pdf", height=4, width=6.5)
5 124
6 125 exfit ← extract(fit, permuted = FALSE, inc_warmup = TRUE)
7 126 plotdata ← melt(exfit[,,"R0"])
8 127 tracefitR0 ← ggplot() +
9 128         geom_line(data = plotdata,
10 129                 aes(x = iterations,
11 130                     y = value,
12 131                     color = factor(chains, labels = 1:stan_
13                                     options$chains))) +
14 132         labs(x = "Sample", y = expression(R[0]), color = "Chain
15              ") +
16 133         scale_color_brewer(palette="Greys") +
17 134         theme_bw()
18 135
19 136 print(tracefitR0)
20 137 ggsave(tracefitR0, filename="traceplotR0_inc.pdf", height=4, width
21              =6.5)
22 138
23 139 sso ← as.shinystan(fit)
24 140 sso ← launch_shinystan(sso)
25

```

26 A.2 Full Stan code

27 Stan model code to be used with the preceding R code.

```

28
29 1 ## Dexter Barrows
30 2 ## McMaster University
31 3 ## 2016
32 4
33 5 data {
34 6
35 7     int      <lower=1>    T;      // total integration steps
36 8     real     y[T];       // observed number of cases
37 9     int      <lower=1>    N;      // population size
38 10    real     h;          // step size
39 11
40 12 }
41 13
42 14 parameters {
43 15
44 16     real <lower=0, upper=10> R0;    // R0
45 17     real <lower=0, upper=10> r;     // recovery rate
46 18     real <lower=0, upper=20> sigma; // observation error
47 19     real <lower=0, upper=500> y0[3]; // initial conditions

```

```

1  20
2  21 }
3  22
4  23 model {
5  24
6  25     real S[T];
7  26     real I[T];
8  27     real R[T];
9  28
10 29     S[1] <- y0[1];
11 30     I[1] <- y0[2];
12 31     R[1] <- y0[3];
13 32
14 33     y[1] ~ normal(y0[2], sigma);
15 34
16 35     for (t in 2:T) {
17 36
18 37         S[t] <- S[t-1] + h*( - S[t-1]*I[t-1]*R0*r/N );
19 38         I[t] <- I[t-1] + h*( S[t-1]*I[t-1]*R0*r/N - I[t-1]*r );
20 39         R[t] <- R[t-1] + h*( I[t-1]*r );
21 40
22 41         if (y[t] > 0) {
23 42             y[t] ~ normal( I[t], sigma );
24 43         }
25 44
26 45     }
27 46
28 47     y0[1] ~ normal(N - y[1], sigma);
29 48     y0[2] ~ normal(y[1], sigma);
30 49
31 50     R0      ~ lognormal(1,1);
32 51     r        ~ lognormal(1,1);
33 52     sigma    ~ lognormal(1,1);
34 53
35 54 }
36

```

1 Appendix B

2 Iterated Filtering

3 B.1 Full R code

4 This code will run all the indicated analysis and produce all plots.

```
5  
6 1 ## Author: Dexter Barrows  
7 2 ## Github: dbarrows.github.io  
8 3  
9 4 library(deSolve)  
10 5 library(ggplot2)  
11 6 library(reshape2)  
12 7 library(gridExtra)  
13 8 library(Rcpp)  
14 9  
15 10 SIR ← function(Time, State, Pars) {  
16 11  
17 12     with(as.list(c(State, Pars)), {  
18 13  
19 14         B ← R0*r/N  
20 15         BSI ← B*S*I  
21 16         rI ← r*I  
22 17  
23 18         dS = -BSI  
24 19         dI = BSI - rI  
25 20         dR = rI  
26 21  
27 22         return(list(c(dS, dI, dR)))  
28 23  
29 24     })  
30 25  
31 26 }  
32 27  
33 28 T ← 100  
34 29 N ← 500
```

```

1 30 sigma ← 10
2 31 i_infec ← 5
3 32
4 33 ## Generate true trajecory and synthetic data
5 34 ##
6 35
7 36 true_init_cond ← c(S = N - i_infec,
8 37                   I = i_infec,
9 38                   R = 0)
10 39
11 40 true_pars ← c(R0 = 3.0,
12 41              r = 0.1,
13 42              N = 500.0)
14 43
15 44 odeout ← ode(true_init_cond, 0:T, SIR, true_pars)
16 45 trueTraj ← odeout[,3]
17 46
18 47 set.seed(1001)
19 48
20 49 infec_counts_raw ← odeout[,3] + rnorm(T+1, 0, sigma)
21 50 infec_counts ← ifelse(infec_counts_raw < 0, 0, infec_counts_raw)
22 51
23 52 g ← qplot(0:T, odeout[,3], geom = "line", xlab = "Time (weeks)", ylab =
24 53       "Infection Count") +
25 54       geom_point(aes(y = infec_counts)) +
26 55       theme_bw()
27 56
28 56 print(g)
29 57 ggsave(g, filename="dataplot.pdf", height=4, width=6.5)
30 58
31 59 ## Rcpp stuff
32 60 ##
33 61
34 62 sourceCpp(paste(getwd(),"d_if2.cpp",sep="/"))
35 63
36 64 paramdata ← data.frame(if2(infec_counts, T+1, N))
37 65 colnames(paramdata) ← c("R0", "r", "sigma", "Sinit", "Iinit", "Rinit"
38 66       )
39 67
40 67 ## Parameter density kernels
41 68 ##
42 69
43 70 R0points ← paramdata$R0
44 71 R0kernel ← qplot(R0points, geom = "density", xlab = expression(R[0]),
45 72       ylab = "frequency") +
46 73       geom_vline(aes(xintercept=true_pars[["R0"]]), linetype="
47 74       dashed", size=1, color="grey50") +
48 75       theme_bw()
49 76
50 75 print(R0kernel)
51 76 ggsave(R0kernel, filename="kernelR0.pdf", height=3, width=3.25)

```

```

1 77
2 78 rpoints <- paramdata$r
3 79 rkernell <- qplot(rpoints, geom = "density", xlab = "r", ylab = "
4 frequency") +
5 80   geom_vline(aes(xintercept=true_pars[["r"]]), linetype="dashed
6   ", size=1, color="grey50") +
7 81   theme_bw()
8 82
9 83 print(rkernell)
10 84 ggsave(rkernell, filename="kernelr.pdf", height=3, width=3.25)
11 85
12 86 sigmapoints <- paramdata$sigma
13 87 sigmakernell <- qplot(sigmapoints, geom = "density", xlab = expression(
14   sigma), ylab = "frequency") +
15 88   geom_vline(aes(xintercept=sigma), linetype="dashed", size=1,
16   color="grey50") +
17 89   theme_bw()
18 90
19 91 print(sigmakernell)
20 92 ggsave(sigmakernell, filename="kernelsigma.pdf", height=3, width=3.25)
21 93
22 94 infecpoints <- paramdata$Iinit
23 95 infeckernell <- qplot(infecpoints, geom = "density", xlab = "Initial
24   Infected", ylab = "frequency") +
25 96   geom_vline(aes(xintercept=true_init_cond[["I"]]), linetype="
26   dashed", size=1, color="grey50") +
27 97   theme_bw()
28 98
29 99 print(infeckernell)
30 100 ggsave(infeckernell, filename="kernelinfec.pdf", height=3, width=3.25)
31 101
32 102 # show grid
33 103 grid.arrange(R0kernell, rkernell, sigmakernell, infeckernell, ncol = 2,
34   nrow = 2)
35 104
36 105 pdf("if2kernels.pdf", height = 6.5, width = 6.5)
37 106 grid.arrange(R0kernell, rkernell, sigmakernell, infeckernell, ncol = 2,
38   nrow = 2)
39 107 dev.off()
40 108 #ggsave(filename="if2kernels.pdf", g2, height=6.5, width=6.5)
41

```

42 B.2 Full C++ code

43 Stan model code to be used with the preceding R code.

```

44
45 1 /* Author: Dexter Barrows
46 2   Github: dbarrows.github.io
47 3

```

```

1  4  */
2  5
3  6 #include <stdio.h>
4  7 #include <math.h>
5  8 #include <sys/time.h>
6  9 #include <time.h>
7 10 #include <stdlib.h>
8 11 #include <string>
9 12 #include <cmath>
10 13 #include <cstdlib>
11 14 #include <fstream>
12 15
13 16 // #include "rand.h"
14 17 // #include "timer.h"
15 18
16 19 #define Treal    100           // time to simulate over
17 20 #define R0true   3.0           // infectiousness
18 21 #define rtrue    0.1           // recovery rate
19 22 #define Nreal    500.0         // population size
20 23 #define merr     10.0          // expected measurement error
21 24 #define I0       5.0           // Initial infected individuals
22 25
23 26 #include <Rcpp.h>
24 27 using namespace Rcpp;
25 28
26 29
27 30 struct Particle {
28 31     double R0;
29 32     double r;
30 33     double sigma;
31 34     double S;
32 35     double I;
33 36     double R;
34 37     double Sinit;
35 38     double Iinit;
36 39     double Rinit;
37 40 };
38 41
39 42 struct ParticleInfo {
40 43     double R0mean;    double R0sd;
41 44     double rmean;     double rsd;
42 45     double sigmamean; double sigmasd;
43 46     double Sinitmean; double Sinitsd;
44 47     double Iinitmean; double Iinitsd;
45 48     double Rinitmean; double Rinitsd;
46 49 };
47 50
48 51
49 52 int timeval_subtract (double *result, struct timeval *x, struct
50 53     timeval *y);
51 54 int check_double(double x, double y);

```



```

1 54 void exp_euler_SIR(double h, double t0, double tn, int N, Particle *
2     particle);
3 55 void copyParticle(Particle * dst, Particle * src);
4 56 void perturbParticles(Particle * particles, int N, int NP, int
5     passnum, double coolrate);
6 57 bool isCollapsed(Particle * particles, int NP);
7 58 void particleDiagnostics(ParticleInfo * partInfo, Particle *
8     particles, int NP);
9 59 NumericMatrix if2(NumericVector * data, int T, int N);
10 60 double randu();
11 61 double randn();
12 62
13 63 // [[Rcpp::export]]
14 64 NumericMatrix if2(NumericVector data, int T, int N) {
15 65
16 66     int      NP          = 2500;
17 67     int      nPasses     = 50;
18 68     double   coolrate    = 0.975;
19 69
20 70     int      i_infec     = I0;
21 71
22 72     NumericMatrix paramdata(NP, 6);
23 73
24 74     srand(time(NULL));    // Seed PRNG with system time
25 75
26 76     double w[NP];        // particle weights
27 77
28 78     Particle particles[NP]; // particle estimates for current
29     step
30 79     Particle particles_old[NP]; // intermediate particle states for
31     resampling
32 80
33 81     printf("Initializing particle states\n");
34 82
35 83     // initialize particle parameter states (seeding)
36 84     for (int n = 0; n < NP; n++) {
37 85
38 86         double R0can, rcan, sigmacan, Iinitcan;
39 87
40 88         do {
41 89             R0can = R0true + R0true*randn();
42 90         } while (R0can < 0);
43 91         particles[n].R0 = R0can;
44 92
45 93         do {
46 94             rcan = rtrue + rtrue*randn();
47 95         } while (rcan < 0);
48 96         particles[n].r = rcan;
49 97
50 98         do {
51 99             sigmacan = merr + merr*randn();

```

```

1 100     } while (sigmacan < 0);
2 101     particles[n].sigma = sigmacan;
3 102
4 103     do {
5 104         Iinitcan = i_infec + i_infec*randn();
6 105     } while (Iinitcan < 0 || N < Iinitcan);
7 106     particles[n].Sinit = N - Iinitcan;
8 107     particles[n].Iinit = Iinitcan;
9 108     particles[n].Rinit = 0.0;
10 109
11 110 }
12 111
13 112 // START PASSES THROUGH DATA
14 113
15 114 printf("Starting filter\n");
16 115 printf("-----\n");
17 116 printf("Pass\n");
18 117
19 118
20 119 for (int pass = 0; pass < nPasses; pass++) {
21 120
22 121     printf("...%d / %d\n", pass, nPasses);
23 122
24 123     perturbParticles(particles, N, NP, pass, coolrate);
25 124
26 125     // initialize particle system states
27 126     for (int n = 0; n < NP; n++) {
28 127
29 128         particles[n].S = particles[n].Sinit;
30 129         particles[n].I = particles[n].Iinit;
31 130         particles[n].R = particles[n].Rinit;
32 131
33 132     }
34 133
35 134     // between-pass perturbations
36 135
37 136     for (int t = 1; t < T; t++) {
38 137
39 138         // between-iteration perturbations
40 139         perturbParticles(particles, N, NP, pass, coolrate);
41 140
42 141         // generate individual predictions and weight
43 142         for (int n = 0; n < NP; n++) {
44 143
45 144             exp_euler_SIR(1.0/10.0, 0.0, 1.0, N, &particles[n]);
46 145
47 146             double merr_par = particles[n].sigma;
48 147             double y_diff   = data[t] - particles[n].I;
49 148
50 149             w[n] = 1.0/(merr_par*sqrt(2.0*M_PI)) * exp( - y_diff*
51             y_diff / (2.0*merr_par*merr_par) );

```

```

1 150
2 151     }
3 152
4 153     // cumulative sum
5 154     for (int n = 1; n < NP; n++) {
6 155         w[n] += w[n-1];
7 156     }
8 157
9 158     // save particle states to resample from
10 159     for (int n = 0; n < NP; n++){
11 160         copyParticle(&particles_old[n], &particles[n]);
12 161     }
13 162
14 163     // resampling
15 164     for (int n = 0; n < NP; n++) {
16 165
17 166         double w_r = randu() * w[NP-1];
18 167         int i = 0;
19 168         while (w_r > w[i]) {
20 169             i++;
21 170         }
22 171
23 172         // i is now the index to copy state from
24 173         copyParticle(&particles[n], &particles_old[i]);
25 174
26 175     }
27 176
28 177     }
29 178
30 179 }
31 180
32 181 ParticleInfo pInfo;
33 182 particleDiagnostics(&pInfo, particles, NP);
34 183
35 184 printf("Parameter results (mean | sd)\n");
36 185 printf("-----\n");
37 186 printf("R0      %f %f\n", pInfo.R0mean, pInfo.R0sd);
38 187 printf("r      %f %f\n", pInfo.rmean, pInfo.rsd);
39 188 printf("sigma   %f %f\n", pInfo.sigamean, pInfo.sigmasd);
40 189 printf("S_init  %f %f\n", pInfo.Sinitmean, pInfo.Sinit_sd);
41 190 printf("I_init  %f %f\n", pInfo.Iinitmean, pInfo.Iinit_sd);
42 191 printf("R_init  %f %f\n", pInfo.Rinitmean, pInfo.Rinit_sd);
43 192
44 193 printf("\n");
45 194
46 195
47 196
48 197 // Get particle results to pass back to R
49 198
50 199 for (int n = 0; n < NP; n++) {
51 200

```

```

1 201     paramdata(n, 0) = particles[n].R0;
2 202     paramdata(n, 1) = particles[n].r;
3 203     paramdata(n, 2) = particles[n].sigma;
4 204     paramdata(n, 3) = particles[n].Sinit;
5 205     paramdata(n, 4) = particles[n].Iinit;
6 206     paramdata(n, 5) = particles[n].Rinit;
7 207
8 208 }
9 209
10 210     return paramdata;
11 211
12 212 }
13 213
14 214
15 215 /* Use the Explicit Euler integration scheme to integrate SIR model
16 216     forward in time
17 216     double h      - time step size
18 217     double t0     - start time
19 218     double tn     - stop time
20 219     double * y    - current system state; a three-component vector
21 219                   representing [S I R], susceptible-infected-recovered
22 220
23 221     */
24 222 void exp_euler_SIR(double h, double t0, double tn, int N, Particle *
25 223     particle) {
26 223
27 224     int num_steps = floor( (tn-t0) / h );
28 225
29 226     double S = particle->S;
30 227     double I = particle->I;
31 228     double R = particle->R;
32 229
33 230     double R0    = particle->R0;
34 231     double r      = particle->r;
35 232     double B      = R0 * r / N;
36 233
37 234     for(int i = 0; i < num_steps; i++) {
38 235         // get derivatives
39 236         double dS = - B*S*I;
40 237         double dI = B*S*I - r*I;
41 238         double dR = r*I;
42 239         // step forward by h
43 240         S += h*dS;
44 241         I += h*dI;
45 242         R += h*dR;
46 243     }
47 244
48 245     particle->S = S;
49 246     particle->I = I;
50 247     particle->R = R;
51 248

```

```

1 249 }
2 250
3 251
4 252 /* Particle pertubation function to be run between iterations and
5      passes
6 253
7 254 */
8 255 void perturbParticles(Particle * particles, int N, int NP, int
9      passnum, double coolrate) {
10 256
11 257     double coolcoef = pow(coolrate, passnum);
12 258
13 259     double spreadR0      = coolcoef * R0true / 10.0;
14 260     double spreadr       = coolcoef * rtrue  / 10.0;
15 261     double spreadsigma   = coolcoef * merr   / 10.0;
16 262     double spreadIinit   = coolcoef * I0     / 10.0;
17 263
18 264     double R0can, rcan, sigmacan, Iinitcan;
19 265
20 266     for (int n = 0; n < NP; n++) {
21 267
22 268         do {
23 269             R0can = particles[n].R0 + spreadR0*randn();
24 270         } while (R0can < 0);
25 271         particles[n].R0 = R0can;
26 272
27 273         do {
28 274             rcan = particles[n].r + spreadr*randn();
29 275         } while (rcan < 0);
30 276         particles[n].r = rcan;
31 277
32 278         do {
33 279             sigmacan = particles[n].sigma + spreadsigma*randn();
34 280         } while (sigmacan < 0);
35 281         particles[n].sigma = sigmacan;
36 282
37 283         do {
38 284             Iinitcan = particles[n].Iinit + spreadIinit*randn();
39 285         } while (Iinitcan < 0 || Iinitcan > 500);
40 286         particles[n].Iinit = Iinitcan;
41 287         particles[n].Sinit = N - Iinitcan;
42 288
43 289     }
44 290
45 291 }
46 292
47 293
48 294 /* Convinience function for particle resampling process
49 295
50 296 */
51 297 void copyParticle(Particle * dst, Particle * src) {

```

```

1 298
2 299     dst->R0      = src->R0;
3 300     dst->r       = src->r;
4 301     dst->sigma   = src->sigma;
5 302     dst->S       = src->S;
6 303     dst->I       = src->I;
7 304     dst->R       = src->R;
8 305     dst->Sinit   = src->Sinit;
9 306     dst->Iinit   = src->Iinit;
10 307     dst->Rinit  = src->Rinit;
11 308
12 309 }
13 310
14 311
15 312 /* Checks to see if particles are collapsed
16 313    This is done by checking if the standard deviations between the
17    particles' parameter
18 314    values are significantly close to one another. Spread threshold
19    may need to be tuned.
20 315
21 316 */
22 317 bool isCollapsed(Particle * particles, int NP) {
23 318
24 319     bool retVal;
25 320
26 321     double R0mean = 0, rmean = 0, sigmamean = 0, Sinitmean = 0,
27     Iinitmean = 0, Rinitmean = 0;
28 322
29 323     // means
30 324
31 325     for (int n = 0; n < NP; n++) {
32 326
33 327         R0mean      += particles[n].R0;
34 328         rmean       += particles[n].r;
35 329         sigmamean   += particles[n].sigma;
36 330         Sinitmean   += particles[n].Sinit;
37 331         Iinitmean   += particles[n].Iinit;
38 332         Rinitmean   += particles[n].Rinit;
39 333
40 334     }
41 335
42 336     R0mean      /= NP;
43 337     rmean       /= NP;
44 338     sigmamean   /= NP;
45 339     Sinitmean   /= NP;
46 340     Iinitmean   /= NP;
47 341     Rinitmean   /= NP;
48 342
49 343     double R0sd = 0, rsd = 0, sigmasd = 0, Sinitsd = 0, Iinitsd = 0,
50     Rinitsd = 0;
51 344

```

```

1 345     for (int n = 0; n < NP; n++) {
2 346
3 347         R0sd      += ( particles[n].R0 - R0mean ) * ( particles[n].R0 -
4         R0mean );
5 348         rsd       += ( particles[n].r - rmean ) * ( particles[n].r -
6         rmean );
7 349         sigmasd   += ( particles[n].sigma - sigmamean ) * ( particles[n]
8         ].sigma - sigmamean );
9 350         Sinitd    += ( particles[n].Sinit - Sinitmean ) * ( particles[n]
10        ].Sinit - Sinitmean );
11 351         Iinitd    += ( particles[n].Iinit - Iinitmean ) * ( particles[n]
12        ].Iinit - Iinitmean );
13 352         Rinitd    += ( particles[n].Rinit - Rinitmean ) * ( particles[n]
14        ].Rinit - Rinitmean );
15 353
16 354     }
17 355
18 356     R0sd          /= NP;
19 357     rsd           /= NP;
20 358     sigmasd       /= NP;
21 359     Sinitd        /= NP;
22 360     Iinitd        /= NP;
23 361     Rinitd        /= NP;
24 362
25 363     if ( (R0sd + rsd + sigmasd) < 1e-5)
26 364         retVal = true;
27 365     else
28 366         retVal = false;
29 367
30 368     return retVal;
31 369
32 370 }
33 371
34 372 void particleDiagnostics(ParticleInfo * partInfo, Particle *
35     particles, int NP) {
36 373
37 374     double  R0mean      = 0.0,
38 375            rmean       = 0.0,
39 376            sigmamean    = 0.0,
40 377            Sinitmean    = 0.0,
41 378            Iinitmean    = 0.0,
42 379            Rinitmean    = 0.0;
43 380
44 381     // means
45 382
46 383     for (int n = 0; n < NP; n++) {
47 384
48 385         R0mean      += particles[n].R0;
49 386         rmean       += particles[n].r;
50 387         sigmamean   += particles[n].sigma;
51 388         Sinitmean   += particles[n].Sinit;

```

```

1 389         Iinitmean    += particles[n].Iinit;
2 390         Rinitmean    += particles[n].Rinit;
3 391
4 392     }
5 393
6 394     R0mean        /= NP;
7 395     rmean         /= NP;
8 396     sigmamean     /= NP;
9 397     Sinitmean     /= NP;
10 398     Iinitmean     /= NP;
11 399     Rinitmean     /= NP;
12 400
13 401     // standard deviations
14 402
15 403     double  R0sd      = 0.0,
16 404             rsd       = 0.0,
17 405             sigmasd   = 0.0,
18 406             Sinitsd   = 0.0,
19 407             Iinitd    = 0.0,
20 408             Rinitd    = 0.0;
21 409
22 410     for (int n = 0; n < NP; n++) {
23 411
24 412         R0sd      += ( particles[n].R0 - R0mean ) * ( particles[n].R0 -
25                    R0mean );
26 413         rsd       += ( particles[n].r - rmean ) * ( particles[n].r -
27                    rmean );
28 414         sigmasd += ( particles[n].sigma - sigmamean ) * ( particles[n]
29                    ].sigma - sigmamean );
30 415         Sinitsd += ( particles[n].Sinit - Sinitmean ) * ( particles[n]
31                    ].Sinit - Sinitmean );
32 416         Iinitd  += ( particles[n].Iinit - Iinitmean ) * ( particles[n]
33                    ].Iinit - Iinitmean );
34 417         Rinitd  += ( particles[n].Rinit - Rinitmean ) * ( particles[n]
35                    ].Rinit - Rinitmean );
36 418
37 419     }
38 420
39 421     R0sd          /= NP;
40 422     rsd           /= NP;
41 423     sigmasd       /= NP;
42 424     Sinitsd       /= NP;
43 425     Iinitd        /= NP;
44 426     Rinitd        /= NP;
45 427
46 428     partInfo->R0mean    = R0mean;
47 429     partInfo->R0sd      = R0sd;
48 430     partInfo->sigmamean = sigmamean;
49 431     partInfo->sigmasd   = sigmasd;
50 432     partInfo->rmean     = rmean;
51 433     partInfo->rsd       = rsd;

```



```
1 434     partInfo->Sinitmean = Sinitmean;
2 435     partInfo->Sinitstd  = Sinitstd;
3 436     partInfo->Iinitmean = Iinitmean;
4 437     partInfo->Iinitstd  = Iinitstd;
5 438     partInfo->Rinitmean = Rinitmean;
6 439     partInfo->Rinitstd  = Rinitstd;
7 440
8 441 }
9 442
10 443 double randu() {
11 444
12 445     return (double) rand() / (double) RAND_MAX;
13 446
14 447 }
15 448
16 449
17 450 /* Return a normally distributed random number with mean 0 and
18 451    standard deviation 1
19 452    Uses the polar form of the Box-Muller transformation
20 453    From http://www.design.caltech.edu/erik/Misc/Gaussian.html
21 454    */
22 455 double randn() {
23 456
24 457     double x1, x2, w, y1;
25 458
26 459     do {
27 460         x1 = 2.0 * randu() - 1.0;
28 461         x2 = 2.0 * randu() - 1.0;
29 462         w = x1 * x1 + x2 * x2;
30 463     } while ( w >= 1.0 );
31 464
32 465     w = sqrt( (-2.0 * log( w ) ) / w );
33 466     y1 = x1 * w;
34 467
35 468     return y1;
36 469 }
37 470
38 471 }
```

¹ **Appendix C**

² **Parameter Fitting**

1 Appendix D

2 Forecasting Frameworks

3 D.1 IF2 Parametric Bootstrapping Function

4 The parametric bootstrapping machinery used to produce forecasts.

```
5 1 # Dexter Barrows
6 2 #
7 3 # IF2 parametric bootstrapping function
8 4
9 5 library(foreach)
10 6 library(parallel)
11 7 library(doParallel)
12 8 library(Rcpp)
13 9
14 10 if2_paraboot ← function(if2data_parent, T, Tlim, steps, N, nTrials,
15 11 if2file, if2_s_file, stoc_sir_file, NP, nPasses, coolrate) {
16 12
17 13   source(stoc_sir_file)
18 14
19 15   if (nTrials < 2)
20 16     ntrials ← 2
21 17
22 18   # unpack if2 first fit data
23 19   # ...parameters
24 20   paramdata_parent ← data.frame( if2data_parent$paramdata )
25 21   names(paramdata_parent) ← c("R0", "r", "sigma", "eta", "berr", "
26 22     Sinit", "Iinit", "Rinit")
27 23   parmeans_parent ← colMeans(paramdata_parent)
28 24   names(parmmeans_parent) ← c("R0", "r", "sigma", "eta", "berr", "
29 25     Sinit", "Iinit", "Rinit")
30 26   # ...states
31 27   statedata_parent ← data.frame( if2data_parent$statedata )
32 28   names(statedata_parent) ← c("S", "I", "R", "B")
33 29   statemeans_parent ← colMeans(statedata_parent)
```

```

1 27 names(statemeans_parent) ← c("S", "I", "R", "B")
2 28
3 29
4 30 ## use parametric bootstrapping to generate forecasts
5 31 ##
6 32 trajectories ← foreach( i = 1:nTrials, .combine = rbind, .packages
7      = "Rcpp") %dopar% {
8 33
9 34   source(stoc_sir_file)
10 35
11 36   ## draw new data
12 37   ##
13 38
14 39   pars ← with( as.list(parmeans_parent),
15 40               c(R0 = R0,
16 41                 r = r,
17 42                 N = N,
18 43                 eta = eta,
19 44                 berr = berr) )
20 45
21 46   init_cond ← with( as.list(parmeans_parent),
22 47                     c(S = Sinit,
23 48                       I = Iinit,
24 49                       R = Rinit) )
25 50
26 51   # generate trajectory
27 52   sdeout ← StocSIR(init_cond, pars, Tlim + 1, steps)
28 53   colnames(sdeout) ← c('S', 'I', 'R', 'B')
29 54
30 55   # add noise
31 56   counts_raw ← sdeout[, 'I'] + rnorm(dim(sdeout)[1], 0, parmeans_
32      parent[['sigma']])
33 57   counts      ← ifelse(counts_raw < 0, 0, counts_raw)
34 58
35 59   ## refit using new data
36 60   ##
37 61
38 62   rm(if2) # because stupid things get done in packages
39 63   sourceCpp(if2file)
40 64   if2time ← system.time( if2data ← if2(counts, Tlim+1, N, NP,
41      nPasses, coolrate) )
42 65
43 66   paramdata ← data.frame( if2data$paramdata )
44 67   names(paramdata) ← c("R0", "r", "sigma", "eta", "berr", "Sinit",
45      "Iinit", "Rinit")
46 68   parmeans ← colMeans(paramdata)
47 69   names(parmeans) ← c("R0", "r", "sigma", "eta", "berr", "Sinit", "
48      Iinit", "Rinit")
49 70
50 71   ## generate the rest of the trajectory
51 72   ##

```

```

1  73
2  74   # pack new parameter estimates
3  75   pars ← with( as.list(parmmeans),
4  76                 c(R0 = R0,
5  77                   r = r,
6  78                   N = N,
7  79                   eta = eta,
8  80                   berr = berr) )
9  81   init_cond ← c(S = statemeans_parent[['S']],
10 82                I = statemeans_parent[['I']],
11 83                R = statemeans_parent[['R']])
12 84
13 85   # generate remaining trajectory part
14 86   sdeout_future ← StocSIR(init_cond, pars, T-Tlim, steps)
15 87   colnames(sdeout_future) ← c('S','I','R','B')
16 88
17 89   return ( c( counts = unname(sdeout_future[, 'I']),
18 90              parmeans,
19 91              time = if2time[['user.self']]) )
20 92
21 93
22 94 }
23 95
24 96 return(trajectories)
25 97
26 98 }

```

28 D.2 RStan Forward Simulator

29 The code used to reconstruct the state estimates, then project the trajectory forward
30 past data.

```

31 1 StocSIRstan ← function(y, pars, T, steps, berrvec, bveclim) {
32 2
33 3   out ← matrix(NA, nrow = (T+1), ncol = 4)
34 4
35 5   R0 ← pars[['R0']]
36 6   r ← pars[['r']]
37 7   N ← pars[['N']]
38 8   eta ← pars[['eta']]
39 9   berr ← pars[['berr']]
40 10
41 11   S ← y[['S']]
42 12   I ← y[['I']]
43 13   R ← y[['R']]
44 14
45 15   B0 ← R0 * r / N
46 16   B ← B0

```

```

1 17
2 18 out[1,] ← c(S,I,R,B)
3 19
4 20 h ← 1 / steps
5 21
6 22 for ( i in 1:(T*steps) ) {
7 23
8 24     if (i <= bveclim) {
9 25         B ← exp( log(B) + eta*(log(B0) - log(B)) + berrvec[i])
10 26     } else {
11 27         B ← exp( log(B) + eta*(log(B0) - log(B)) + rnorm(1, 0, berr
12 28         ))
13 29     }
14 30
15 31     BSI ← B*S*I
16 32     rI ← r*I
17 33
18 34     dS ← -BSI
19 35     dI ← BSI - rI
20 36     dR ← rI
21 37
22 38     S ← S + h*dS #newInf
23 39     I ← I + h*dI #newInf - h*dR
24 40     R ← R + h*dR #h*dR
25 41
26 42     if (i %% steps == 0)
27 43         out[i/steps+1,] ← c(S,I,R,B)
28 44
29 45 }
30 46
31 47 return(out)
32 48
33 49 }

```

1 Appendix E

2 S-map and SIRS

3 E.1 SIRS R Function Code

4 R code to simulate the outlines SIRS function.

```
5  
6 1 StocSIRS ← function(y, pars, T, steps) {  
7 2  
8 3   out ← matrix(NA, nrow = (T+1), ncol = 4)  
9 4  
10 5   R0 ← pars[['R0']]  
11 6   r ← pars[['r']]  
12 7   N ← pars[['N']]  
13 8   eta ← pars[['eta']]  
14 9   berr ← pars[['berr']]  
15 10   re ← pars[['re']]  
16 11  
17 12   S ← y[['S']]  
18 13   I ← y[['I']]  
19 14   R ← y[['R']]  
20 15  
21 16   B0 ← R0 * r / N  
22 17   B ← B0  
23 18  
24 19   out[1,] ← c(S,I,R,B)  
25 20  
26 21   h ← 1 / steps  
27 22  
28 23   for ( i in 1:(T*steps) ) {  
29 24  
30 25       #Bfac ← 1/2 - cos((2*pi/365)*i)/2  
31 26       Bfac ← exp(2*cos((2*pi/365)*i) - 2)  
32 27  
33 28       B ← exp( log(B) + eta*(log(B0) - log(B)) + rnorm(1, 0, berr) )  
34 29
```

```

1 30 BSI ← Bfac*B*S*I
2 31 rI ← r*I
3 32 reR ← re*R
4 33
5 34 dS ← -BSI + reR
6 35 dI ← BSI - rI
7 36 dR ← rI - reR
8 37
9 38 S ← S + h*dS #newInf
10 39 I ← I + h*dI #newInf - h*dR
11 40 R ← R + h*dR #h*dR
12 41
13 42 if (i %% steps == 0)
14 43   out[i/steps+1,] ← c(S,I,R,B)
15 44
16 45 }
17 46
18 47 colnames(out) ← c("S","I","R","B")
19 48 return(out)
20 49
21 50 }
22 51
23 52 ### suggested parameters
24 53 #
25 54 # T      ← 200
26 55 # i_infec ← 10
27 56 # steps  ← 7
28 57 # N      ← 500
29 58 # sigma  ← 5
30 59 #
31 60 # pars ← c(R0 = 3.0, # new infected people per infected person
32 61 #           r = 0.1, # recovery rate
33 62 #           N = 500, # population size
34 63 #           eta = 0.5, # geometric random walk
35 64 #           berr = 0.5, # Beta geometric walk noise
36 65 #           re = 1) # resuceptibility rate

```

38 E.2 SMAP Code

39 This code implements an SMAP function on a user-provided time series.

```

40
41 1 library(pracma)
42 2
43 3 smap ← function(data, E, theta, stepsAhead) {
44 4
45 5   # construct library
46 6   tseries ← as.vector(data)
47 7   liblen ← length(tseries) - E + 1 - stepsAhead

```



```

1  8      lib      ← matrix(NA, liblen, E)
2  9
3  10     for (i in 1:E) {
4  11         lib[,i] ← tseries[(E-i+1):(liblen+E-i)]
5  12     }
6  13
7  14     # predict from the last index
8  15     tslen ← length(tseries)
9  16     predictee ← rev(t(as.matrix(tseries[(tslen-E+1):tslen])))
10 17     predictions ← numeric(stepsAhead)
11 18
12 19     #allPredictees ← matrix(NA, stepsAhead, E)
13 20
14 21     # for each prediction index (number of steps ahead)
15 22     for(i in 1:stepsAhead) {
16 23
17 24         # set up weight calculation
18 25         predmat ← repmat(predictee, liblen, 1)
19 26         distances ← sqrt( rowSums( abs(lib - predmat)^2 ) )
20 27         meanDist ← mean(distances)
21 28
22 29         # calculate weights
23 30         weights ← exp( - (theta * distances) / meanDist )
24 31
25 32         # construct A, B
26 33
27 34         preds ← tseries[(E+i):(liblen+E+i-1)]
28 35
29 36         A ← cbind( rep(1.0, liblen), lib ) * repmat(as.matrix(weights
30 37             ), 1, E+1)
31 38         B ← as.matrix(preds * weights)
32 39
33 40         # solve system for C
34 41
35 42         Asvd ← svd(A)
36 43         C ← Asvd$v %*% diag(1/Asvd$d) %*% t(Asvd$u) %*% B
37 44
38 45         # get prediction
39 46
40 47         predsum ← sum(C * c(1,predictee))
41 48
42 49         # save
43 50
44 51         predictions[i] ← predsum
45 52
46 53         # next predictee
47 54
48 55         #predictee ← c( predsum, predictee[-E] )
49 56         #allPredictees[i,] ← predictee
50 57
51 58     }

```

```

1 58
2 59     return(predictions)
3 60
4 61 }

```

6 E.3 SMAP Parameter Optimization Code

7 This code determines the optimal parameter values to be used by the S-map algo-
8 rithm.

```

9
10 1 library(deSolve)
11 2 library(ggplot2)
12 3 library(RColorBrewer)
13 4 library(pracma)
14 5
15 6 set.seed(1010)
16 7
17 8 ## external files
18 9 ##
19 10 stoc_sirs_file ← paste(getwd(), "../sir-functions", "StocSIRS.r",
20      sep = "/")
21 11 smap_file      ← paste(getwd(), "smap.r", sep = "/")
22 12 source(stoc_sirs_file)
23 13 source(smap_file)
24 14
25 15
26 16
27 17 ## parameters
28 18 ##
29 19 T      ← 6*52
30 20 Tlim   ← T - 52
31 21 i_infec ← 10
32 22 steps  ← 7
33 23 N      ← 500
34 24 sigma  ← 5
35 25
36 26 true_pars ← c( R0 = 3.0, # new infected people per infected person
37 27               r = 0.1, # recovery rate
38 28               N = 500, # population size
39 29               eta = 0.5, # geometric random walk
40 30               berr = 0.5, # Beta geometric walk noise
41 31               re = 1) # resuseptibility rate
42 32
43 33 true_init_cond ← c(S = N - i_infec,
44 34                  I = i_infec,
45 35                  R = 0)
46 36
47 37 ## trial parameter values to check.options

```

```

1 38 ##
2 39 Elist ← 1:20
3 40 thetalist ← 10*exp(-(seq(0,9.5,0.5)))
4 41 nTrials ← 100
5 42
6 43 ssemat ← matrix(NA, 20, 20)
7 44
8 45 for (i in 1:length(Elist)) {
9 46   for (j in 1:length(thetalist)) {
10 47
11 48     ssemean ← 0
12 49
13 50     for (k in 1:nTrials) {
14 51
15 52       E ← Elist[i]
16 53       theta ← thetalist[j]
17 54
18 55       ## get true trajectory
19 56       ##
20 57       sdeout ← StocSIRS(true_init_cond, true_pars, T, steps)
21 58
22 59       ## perturb to get data
23 60       ##
24 61       infec_counts_raw ← sdeout[1:(Tlim+1),'I'] + rnorm(Tlim+1,0,
25 62         sigma)
26 62       infec_counts ← ifelse(infec_counts_raw < 0, 0, infec_counts_
27 63         raw)
28 63
29 64       predictions ← smap(infec_counts, E, theta, 52)
30 65
31 66       err ← sdeout[(Tlim+2):dim(sdeout)[1],'I'] - predictions
32 67       sse ← sum(err^2)
33 68
34 69       ssemean ← ssemean + (sse / nTrials)
35 70
36 71     }
37 72
38 73     ssemat[i,j] ← ssemean
39 74
40 75   }
41 76 }
42 77 }
43 78
44 79 quartz()
45 80 image(-ssemat)
46 81 quartz()
47 82 filled.contour(-ssemat)
48 83
49 84 #print(ssemat)
50 85 #cms ← colMeans(ssemat)
51 86 #rms ← rowMeans(ssemat)

```

```

1  87
2  88 #Emin ← Elist[which.min(rms)]
3  89 #thetamin ← thetalist[which.min(cms)]
4  90 #print(Emin)
5  91 #print(thetamin)
6  92
7  93 mininds ← which(ssemat==min(ssemat),arr.ind=TRUE)
8  94
9  95 Emin ← Elist[mininds[, 'row']]
10 96 thetamin ← thetalist[mininds[, 'col']]
11 97
12 98 print(Emin)
13 99 print(thetamin)

```

15 E.4 RStan SIRS Code

16 This code implements a periodic SIRS model in Rstan.

```

17 1
18 2 data {
19 3
20 4     int      <lower=1>    T;      // total integration steps
21 5     real     y[T];       // observed number of cases
22 6     int      <lower=1>    N;      // population size
23 7     real     h;          // step size
24 8
25 9 }
26 10
27 11 parameters {
28 12
29 13     real <lower=0, upper=10>    R0;      // R0
30 14     real <lower=0, upper=10>    r;       // recovery rate
31 15     real <lower=0, upper=10>    re;      // resusceptibility rate
32 16     real <lower=0, upper=20>    sigma;   // observation error
33 17     real <lower=0, upper=30>    Iinit;   // initial infected
34 18     real <lower=0, upper=1>     eta;     // geometric walk
35 19     attraction strength
36 20     real <lower=0, upper=1>     berr;    // beta walk noise
37 21     real <lower=-1.5, upper=1.5> Bnoise[T]; // Beta vector
38 22
39 23 }
40 24
41 25 //transformed parameters {
42 26 //     real B0 ← R0 * r / N;
43 27 //}
44 28
45 29 model {
46 30
47 31     real S[T];

```

```

1 30 real I[T];
2 31 real R[T];
3 32 real B[T];
4 33 real B0;
5 34
6 35 real pi;
7 36 real Bfac;
8 37
9 38 pi ← 3.1415926535;
10 39
11 40 B0 ← R0 * r / N;
12 41
13 42 B[1] ← B0;
14 43
15 44 S[1] ← N - Iinit;
16 45 I[1] ← Iinit;
17 46 R[1] ← 0.0;
18 47
19 48 for (t in 2:T) {
20 49
21 50     Bnoise[t] ~ normal(0,berr);
22 51     Bfac ← exp(2*cos((2*pi/365)*t) - 2);
23 52     B[t] ← exp( log(B0) + eta * ( log(B[t-1]) - log(B0) ) +
24 53         Bnoise[t] );
25 54
26 54     S[t] ← S[t-1] + h*( - Bfac*B[t]*S[t-1]*I[t-1] + re*R[t-1] );
27 55     I[t] ← I[t-1] + h*( Bfac*B[t]*S[t-1]*I[t-1] - I[t-1]*r );
28 56     R[t] ← R[t-1] + h*( I[t-1]*r - re*R[t-1] );
29 57
30 58     if (y[t] > 0) {
31 59         y[t] ~ normal( I[t], sigma );
32 60     }
33 61
34 62 }
35 63
36 64 R0 ~ lognormal(1,1);
37 65 r ~ lognormal(1,1);
38 66 sigma ~ lognormal(1,1);
39 67 re ~ lognormal(1,1);
40 68 Iinit ~ normal(y[1], sigma);
41 69
42 70 }

```

44 E.5 IF2 SIRS Code

45 This code implements a periodic SIRS model using IF2 in C++.

```

46
47 1 /* Author: Dexter Barrows

```

```

1  2      Github: dbarrows.github.io
2  3
3  4      */
4  5
5  6 #include <stdio.h>
6  7 #include <math.h>
7  8 #include <sys/time.h>
8  9 #include <time.h>
9 10 #include <stdlib.h>
10 11 #include <string>
11 12 #include <cmath>
12 13 #include <cstdlib>
13 14 #include <fstream>
14 15
15 16 // #include "rand.h"
16 17 // #include "timer.h"
17 18
18 19 #define Treal      100      // time to simulate over
19 20 #define R0true     3.0      // infectiousness
20 21 #define rtrue      0.1      // recovery rate
21 22 #define retrue     0.05     // resusceptibility rate
22 23 #define Nreal      500.0    // population size
23 24 #define etatrue    0.5      // real drift attraction strength
24 25 #define berrtrue   0.5      // real beta drift noise
25 26 #define merr       5.0      // expected measurement error
26 27 #define I0         5.0      // Initial infected individuals
27 28
28 29 #define PSC        0.5      // scale factor for more sensitive
29 30      parameters
30 31
31 32 #include <Rcpp.h>
32 33 using namespace Rcpp;
33 34
34 35 struct State {
35 36     double S;
36 37     double I;
37 38     double R;
38 39 };
39 40
40 41 struct Particle {
41 42     double R0;
42 43     double r;
43 44     double re;
44 45     double sigma;
45 46     double eta;
46 47     double berr;
47 48     double B;
48 49     double S;
49 50     double I;
50 51     double R;
51 52     double Sinit;

```

```

1  52      double Iinit;
2  53      double Rinit;
3  54  };
4  55
5  56  struct ParticleInfo {
6  57      double R0mean;      double R0sd;
7  58      double rmean;      double rsd;
8  59      double remean;      double resd;
9  60      double sigmamean;   double sigmasd;
10 61      double etamean;    double etasd;
11 62      double berrmean;    double berrsd;
12 63      double Sinitmean;   double Sinitsd;
13 64      double Iinitmean;   double Iinitsd;
14 65      double Rinitmean;   double Rinitsd;
15 66  };
16 67
17 68
18 69  int timeval_subtract (double *result, struct timeval *x, struct
19      timeval *y);
20 70  int check_double(double x,double y);
21 71  void exp_euler_SIRS(double h, double t0, double tn, int N, Particle *
22      particle);
23 72  void copyParticle(Particle * dst, Particle * src);
24 73  void perturbParticles(Particle * particles, int N, int NP, int
25      passnum, double coolrate);
26 74  void particleDiagnostics(ParticleInfo * partInfo, Particle *
27      particles, int NP);
28 75  void getStateMeans(State * state, Particle* particles, int NP);
29 76  NumericMatrix if2(NumericVector * data, int T, int N);
30 77  double randu();
31 78  double randn();
32 79
33 80  // [[Rcpp::export]]
34 81  Rcpp::List if2_sirs(NumericVector data, int T, int N, int NP, int
35      nPasses, double coolrate) {
36 82
37 83      int npar = 9;
38 84
39 85      NumericMatrix paramdata(NP, npar);
40 86      NumericMatrix means(nPasses, npar);
41 87      NumericMatrix sds(nPasses, npar);
42 88      NumericMatrix statemeans(T, 3);
43 89      NumericMatrix statedata(NP, 4);
44 90
45 91      srand(time(NULL));          // Seed PRNG with system time
46 92
47 93      double w[NP];              // particle weights
48 94
49 95      Particle particles[NP];     // particle estimates for current
50      step
51 96      Particle particles_old[NP]; // intermediate particle states for

```

```

1      resampling
2  97
3  98      printf("Initializing particle states\n");
4  99
5  100     // initialize particle parameter states (seeding)
6  101     for (int n = 0; n < NP; n++) {
7  102
8  103         double R0can, rcan, recan, sigmacan, Iinitcan, etacan,
9
10 104             berrcan;
11 105
12 106         do {
13 107             R0can = R0true + R0true*randn();
14 108         } while (R0can < 0);
15 109         particles[n].R0 = R0can;
16 110
17 111         do {
18 112             rcan = rtrue + rtrue*randn();
19 113         } while (rcan < 0);
20 114         particles[n].r = rcan;
21 115
22 116         do {
23 117             recan = retrue + retrue*randn();
24 118         } while (recan < 0);
25 119         particles[n].re = recan;
26 120
27 121         particles[n].B = (double) R0can * rcan / N;
28 122
29 123         do {
30 124             sigmacan = merr + merr*randn();
31 125         } while (sigmacan < 0);
32 126         particles[n].sigma = sigmacan;
33 127
34 128         do {
35 129             etacan = etatrue + PSC*etatrue*randn();
36 130         } while (etacan < 0 || etacan > 1);
37 131         particles[n].eta = etacan;
38 132
39 133         do {
40 134             berrcan = berrtrue + PSC*berrtrue*randn();
41 135         } while (berrcan < 0);
42 136         particles[n].berr = berrcan;
43 137
44 138         do {
45 139             Iinitcan = I0 + I0*randn();
46 140         } while (Iinitcan < 0 || N < Iinitcan);
47 141         particles[n].Sinit = N - Iinitcan;
48 142         particles[n].Iinit = Iinitcan;
49 143         particles[n].Rinit = 0.0;
50 144     }
51 145

```



```

1 146 // START PASSES THROUGH DATA
2 147
3 148 printf("Starting filter\n");
4 149 printf("-----\n");
5 150 printf("Pass\n");
6 151
7 152
8 153 for (int pass = 0; pass < nPasses; pass++) {
9 154
10 155     printf("...%d / %d\n", pass, nPasses);
11 156
12 157     // reset particle system evolution states
13 158     for (int n = 0; n < NP; n++) {
14 159
15 160         particles[n].S = particles[n].Sinit;
16 161         particles[n].I = particles[n].Iinit;
17 162         particles[n].R = particles[n].Rinit;
18 163         particles[n].B = (double) particles[n].R0 * particles[n].
19 164             r / N;
20 165
21 166     }
22 167
23 168     if (pass == (nPasses-1)) {
24 169         State sMeans;
25 170         getStateMeans(&sMeans, particles, NP);
26 171         statemeans(0,0) = sMeans.S;
27 172         statemeans(0,1) = sMeans.I;
28 173         statemeans(0,2) = sMeans.R;
29 174     }
30 175
31 176     for (int t = 1; t < T; t++) {
32 177
33 178         // generate individual predictions and weight
34 179         for (int n = 0; n < NP; n++) {
35 180
36 181             exp_euler_SIRS(1.0/7.0, (double) t-1, (double) t, N,
37 182                 &particles[n]);
38 183
39 184             double merr_par = particles[n].sigma;
40 185             double y_diff = data[t] - particles[n].I;
41 186
42 187             w[n] = 1.0/(merr_par*sqrt(2.0*M_PI)) * exp( - y_diff*
43 188                 y_diff / (2.0*merr_par*merr_par) );
44 189
45 190         }
46 191
47 192         // cumulative sum
48 193         for (int n = 1; n < NP; n++) {
49 194             w[n] += w[n-1];
50 195
51 196         }

```

```

1 194 // save particle states to resample from
2 195 for (int n = 0; n < NP; n++){
3 196     copyParticle(&particles_old[n], &particles[n]);
4 197 }
5 198
6 199 // resampling
7 200 for (int n = 0; n < NP; n++) {
8 201
9 202     double w_r = randu() * w[NP-1];
10 203     int i = 0;
11 204     while (w_r > w[i]) {
12 205         i++;
13 206     }
14 207
15 208     // i is now the index to copy state from
16 209     copyParticle(&particles[n], &particles_old[i]);
17 210
18 211 }
19 212
20 213 // between-iteration perturbations, not after last time
21 214 step
22 215 if (t < (T-1))
23 216     perturbParticles(particles, N, NP, pass, coolrate);
24 217
25 218 if (pass == (nPasses-1)) {
26 219     State sMeans;
27 220     getStateMeans(&sMeans, particles, NP);
28 221     statemeans(t,0) = sMeans.S;
29 222     statemeans(t,1) = sMeans.I;
30 223     statemeans(t,2) = sMeans.R;
31 224 }
32 225
33 226 }
34 227
35 228 ParticleInfo pInfo;
36 229 particleDiagnostics(&pInfo, particles, NP);
37 230
38 231 means(pass, 0) = pInfo.R0mean;
39 232 means(pass, 1) = pInfo.rmean;
40 233 means(pass, 2) = pInfo.remean;
41 234 means(pass, 3) = pInfo.sigamean;
42 235 means(pass, 4) = pInfo.etamean;
43 236 means(pass, 5) = pInfo.berrmean;
44 237 means(pass, 6) = pInfo.Sinitmean;
45 238 means(pass, 7) = pInfo.Iinitmean;
46 239 means(pass, 8) = pInfo.Rinitmean;
47 240
48 241 sds(pass, 0) = pInfo.R0sd;
49 242 sds(pass, 1) = pInfo.rsd;
50 243 sds(pass, 2) = pInfo.resd;
51 244 sds(pass, 3) = pInfo.sigmasd;

```

```

1 244         sds(pass, 4) = pInfo.etasd;
2 245         sds(pass, 5) = pInfo.berrsd;
3 246         sds(pass, 6) = pInfo.Sinit;
4 247         sds(pass, 7) = pInfo.Iinit;
5 248         sds(pass, 8) = pInfo.Rinit;
6 249
7 250         // between-pass perturbations, not after last pass
8 251         if (pass < (nPasses + 1))
9 252             perturbParticles(particles, N, NP, pass, coolrate);
10 253
11 254     }
12 255
13 256     ParticleInfo pInfo;
14 257     particleDiagnostics(&pInfo, particles, NP);
15 258
16 259     printf("Parameter results (mean | sd)\n");
17 260     printf("-----\n");
18 261     printf("R0          %f %f\n", pInfo.R0mean, pInfo.R0sd);
19 262     printf("r          %f %f\n", pInfo.rmean, pInfo.rsd);
20 263     printf("re          %f %f\n", pInfo.remean, pInfo.resd);
21 264     printf("sigma       %f %f\n", pInfo.sigmean, pInfo.sigsd);
22 265     printf("eta          %f %f\n", pInfo.eta, pInfo.etasd);
23 266     printf("berr        %f %f\n", pInfo.berrmean, pInfo.berrsd);
24 267     printf("S_init      %f %f\n", pInfo.Sinitmean, pInfo.Sinit);
25 268     printf("I_init      %f %f\n", pInfo.Iinitmean, pInfo.Iinit);
26 269     printf("R_init      %f %f\n", pInfo.Rinitmean, pInfo.Rinit);
27 270
28 271     printf("\n");
29 272
30 273
31 274
32 275     // Get particle results to pass back to R
33 276
34 277     for (int n = 0; n < NP; n++) {
35 278
36 279         paramdata(n, 0) = particles[n].R0;
37 280         paramdata(n, 1) = particles[n].r;
38 281         paramdata(n, 2) = particles[n].re;
39 282         paramdata(n, 3) = particles[n].sigma;
40 283         paramdata(n, 4) = particles[n].eta;
41 284         paramdata(n, 5) = particles[n].berr;
42 285         paramdata(n, 6) = particles[n].Sinit;
43 286         paramdata(n, 7) = particles[n].Iinit;
44 287         paramdata(n, 8) = particles[n].Rinit;
45 288
46 289     }
47 290
48 291     for (int n = 0; n < NP; n++) {
49 292
50 293         statedata(n, 0) = particles[n].S;
51 294         statedata(n, 1) = particles[n].I;

```

```

1 295         statedata(n, 2) = particles[n].R;
2 296         statedata(n, 3) = particles[n].B;
3 297
4 298     }
5 299
6 300
7 301
8 302     return Rcpp::List::create( Rcpp::Named("paramdata") = paramdata,
9 303                               Rcpp::Named("means") = means,
10 304                               Rcpp::Named("statemeans") =
11                               statemeans,
12 305                               Rcpp::Named("statedata") = statedata,
13 306                               Rcpp::Named("sds") = sds);
14 307
15 308 }
16 309
17 310
18 311 /* Use the Explicit Euler integration scheme to integrate SIR model
19 forward in time
20 double h      - time step size
21 double t0     - start time
22 double tn     - stop time
23 double * y    - current system state; a three-component vector
24                 representing [S I R], susceptible-infected-recovered
25 316
26 317 */
27 318 void exp_euler_SIRS(double h, double t0, double tn, int N, Particle *
28 particle) {
29 319
30 320     int num_steps = floor( (tn-t0) / h );
31 321
32 322     double S = particle->S;
33 323     double I = particle->I;
34 324     double R = particle->R;
35 325
36 326     double R0   = particle->R0;
37 327     double r    = particle->r;
38 328     double re   = particle->re;
39 329     double B0   = R0 * r / N;
40 330     double eta  = particle->eta;
41 331     double berr = particle->berr;
42 332
43 333     double B = particle->B;
44 334
45 335     for(int i = 0; i < num_steps; i++) {
46 336
47 337         //double Bfac = 0.5 - 0.95*cos( (2.0*M_PI/365)*(t0*num_steps+
48         i )/2.0;
49 338         double Bfac = exp(2*cos((2*M_PI/365)*(t0*num_steps+i)) - 2);
50 339         B = exp( log(B) + eta*(log(B0) - log(B)) + berr*randn() );
51 340

```

```

1 341      double BSI = Bfac*B*S*I;
2 342      double rI  = r*I;
3 343      double reR = re*R;
4 344
5 345      // get derivatives
6 346      double dS = - BSI + reR;
7 347      double dI = BSI - rI;
8 348      double dR = rI - reR;
9 349
10 350      // step forward by h
11 351      S += h*dS;
12 352      I += h*dI;
13 353      R += h*dR;
14 354
15 355  }
16 356
17 357  particle->S = S;
18 358  particle->I = I;
19 359  particle->R = R;
20 360  particle->B = B;
21 361
22 362 }
23 363
24 364
25 365 /* Particle pertubation function to be run between iterations and
26 366 passes
27 366
28 367 */
29 368 void perturbParticles(Particle * particles, int N, int NP, int
30 369 passnum, double coolrate) {
31 369
32 370     //double coolcoef = exp( - (double) passnum / coolrate );
33 371     double coolcoef = pow(coolrate, passnum);
34 372
35 373
36 374     double spreadR0      = coolcoef * R0true / 10.0;
37 375     double spreadr       = coolcoef * rtrue / 10.0;
38 376     double spreadre      = coolcoef * retrue / 10.0;
39 377     double spreadsigma   = coolcoef * merr / 10.0;
40 378     double spreadIinit   = coolcoef * I0 / 10.0;
41 379     double spreadeta     = coolcoef * etatrue / 10.0;
42 380     double spreadberr    = coolcoef * berrtrue / 10.0;
43 381
44 382
45 383     double R0can, rcan, recan, sigmacan, Iinitcan, etacan, berrcan;
46 384
47 385     for (int n = 0; n < NP; n++) {
48 386
49 387         do {
50 388             R0can = particles[n].R0 + spreadR0*randn();
51 389         } while (R0can < 0);

```

```

1 390     particles[n].R0 = R0can;
2 391
3 392     do {
4 393         rcan = particles[n].r + spreadr*randn();
5 394     } while (rcan < 0);
6 395     particles[n].r = rcan;
7 396
8 397     do {
9 398         recan = particles[n].re + spreadre*randn();
10 399     } while (recan < 0);
11 400     particles[n].re = recan;
12 401
13 402     do {
14 403         sigmacan = particles[n].sigma + spreadsigma*randn();
15 404     } while (sigmacan < 0);
16 405     particles[n].sigma = sigmacan;
17 406
18 407     do {
19 408         etacan = particles[n].eta + PSC*spreadeta*randn();
20 409     } while (etacan < 0 || etacan > 1);
21 410     particles[n].eta = etacan;
22 411
23 412     do {
24 413         berrcan = particles[n].berr + PSC*spreadberr*randn();
25 414     } while (berrcan < 0);
26 415     particles[n].berr = berrcan;
27 416
28 417     do {
29 418         Iinitcan = particles[n].Iinit + spreadIinit*randn();
30 419     } while (Iinitcan < 0 || Iinitcan > 500);
31 420     particles[n].Iinit = Iinitcan;
32 421     particles[n].Sinit = N - Iinitcan;
33 422
34 423 }
35 424
36 425 }
37 426
38 427
39 428 /* Convenience function for particle resampling process
40 429
41 430 */
42 431 void copyParticle(Particle * dst, Particle * src) {
43 432
44 433     dst->R0      = src->R0;
45 434     dst->r        = src->r;
46 435     dst->re       = src->re;
47 436     dst->sigma    = src->sigma;
48 437     dst->eta      = src->eta;
49 438     dst->berr     = src->berr;
50 439     dst->B        = src->B;
51 440     dst->S        = src->S;

```

```

1 441     dst->I      = src->I;
2 442     dst->R      = src->R;
3 443     dst->Sinit  = src->Sinit;
4 444     dst->Iinit  = src->Iinit;
5 445     dst->Rinit  = src->Rinit;
6 446
7 447 }
8 448
9 449 void particleDiagnostics(ParticleInfo * partInfo, Particle *
10 450    particles, int NP) {
11 451
12 451     double   R0mean      = 0.0,
13 452             rmean       = 0.0,
14 453             remean       = 0.0,
15 454             sigmamean    = 0.0,
16 455             etamean      = 0.0,
17 456             berrmean     = 0.0,
18 457             Sinitmean    = 0.0,
19 458             Iinitmean    = 0.0,
20 459             Rinitmean    = 0.0;
21 460
22 461     // means
23 462
24 463     for (int n = 0; n < NP; n++) {
25 464
26 465         R0mean      += particles[n].R0;
27 466         rmean       += particles[n].r;
28 467         remean      += particles[n].re;
29 468         etamean     += particles[n].eta,
30 469         berrmean    += particles[n].berr,
31 470         sigmamean   += particles[n].sigma;
32 471         Sinitmean   += particles[n].Sinit;
33 472         Iinitmean   += particles[n].Iinit;
34 473         Rinitmean   += particles[n].Rinit;
35 474
36 475     }
37 476
38 477     R0mean      /= NP;
39 478     rmean       /= NP;
40 479     remean      /= NP;
41 480     sigmamean   /= NP;
42 481     etamean     /= NP;
43 482     berrmean    /= NP;
44 483     Sinitmean   /= NP;
45 484     Iinitmean   /= NP;
46 485     Rinitmean   /= NP;
47 486
48 487     // standard deviations
49 488
50 489     double   R0sd      = 0.0,
51 490             rsd        = 0.0,

```

```

1 491         resd      = 0.0,
2 492         sigmasd   = 0.0,
3 493         etasd     = 0.0,
4 494         berrsd    = 0.0,
5 495         Sinitsd   = 0.0,
6 496         Iinitd    = 0.0,
7 497         Rinitd    = 0.0;
8 498
9 499     for (int n = 0; n < NP; n++) {
10 500
11 501         R0sd      += ( particles[n].R0 - R0mean ) * ( particles[n].R0 -
12                    R0mean );
13 502         rsd       += ( particles[n].r - rmean ) * ( particles[n].r -
14                    rmean );
15 503         resd      += ( particles[n].re - rmean ) * ( particles[n].re -
16                    rmean );
17 504         sigmasd   += ( particles[n].sigma - sigmamean ) * ( particles[n]
18                    ].sigma - sigmamean );
19 505         etasd     += ( particles[n].eta - etamean ) * ( particles[n].
20                    eta - etamean );
21 506         berrsd    += ( particles[n].berr - berrmean ) * ( particles[n].
22                    berr - berrmean );
23 507         Sinitsd   += ( particles[n].Sinit - Sinitmean ) * ( particles[n]
24                    ].Sinit - Sinitmean );
25 508         Iinitd    += ( particles[n].Iinit - Iinitmean ) * ( particles[n]
26                    ].Iinit - Iinitmean );
27 509         Rinitd    += ( particles[n].Rinit - Rinitmean ) * ( particles[n]
28                    ].Rinit - Rinitmean );
29 510
30 511     }
31 512
32 513     R0sd          /= NP;
33 514     rsd           /= NP;
34 515     resd          /= NP;
35 516     sigmasd       /= NP;
36 517     etasd         /= NP;
37 518     berrsd        /= NP;
38 519     Sinitsd       /= NP;
39 520     Iinitd        /= NP;
40 521     Rinitd        /= NP;
41 522
42 523     partInfo->R0mean    = R0mean;
43 524     partInfo->R0sd      = R0sd;
44 525     partInfo->rmean     = rmean;
45 526     partInfo->rsd       = rsd;
46 527     partInfo->remean    = remean;
47 528     partInfo->resd      = resd;
48 529     partInfo->sigmamean = sigmamean;
49 530     partInfo->sigmasd   = sigmasd;
50 531     partInfo->etamean   = etamean;
51 532     partInfo->etasd     = etasd;

```



```

1 533     partInfo->berrmean  = berrmean;
2 534     partInfo->berrsd   = berrsd;
3 535     partInfo->Sinitmean = Sinitmean;
4 536     partInfo->Sinitd    = Sinitd;
5 537     partInfo->Iinitmean = Iinitmean;
6 538     partInfo->Iinitd    = Iinitd;
7 539     partInfo->Rinitmean = Rinitmean;
8 540     partInfo->Rinitd    = Rinitd;
9 541
10 542 }
11 543
12 544 double randu() {
13 545
14 546     return (double) rand() / (double) RAND_MAX;
15 547
16 548 }
17 549
18 550 void getStateMeans(State * state, Particle* particles, int NP) {
19 551
20 552     double Smean = 0, Imean = 0, Rmean = 0;
21 553
22 554     for (int n = 0; n < NP; n++) {
23 555         Smean += particles[n].S;
24 556         Imean += particles[n].I;
25 557         Rmean += particles[n].R;
26 558     }
27 559
28 560     state->S = (double) Smean / NP;
29 561     state->I = (double) Imean / NP;
30 562     state->R = (double) Rmean / NP;
31 563
32 564 }
33 565
34 566
35 567 /* Return a normally distributed random number with mean 0 and
36 568    standard deviation 1
37 568    Uses the polar form of the Box-Muller transformation
38 569    From http://www.design.caltech.edu/erik/Misc/Gaussian.html
39 570    */
40 571 double randn() {
41 572
42 573     double x1, x2, w, y1;
43 574
44 575     do {
45 576         x1 = 2.0 * randu() - 1.0;
46 577         x2 = 2.0 * randu() - 1.0;
47 578         w = x1 * x1 + x2 * x2;
48 579     } while ( w >= 1.0 );
49 580
50 581     w = sqrt( (-2.0 * log( w ) ) / w );
51 582     y1 = x1 * w;

```

```
1 583 |  
2 584 |    return y1;  
3 585 |  
4 586 |}
```

1 Appendix F

2 Spatial Epidemics

3 F.1 Spatial SIR R Function Code

4 R code to simulate the outlined Spatial SIR function.

```
5
6 1 ## ymat:  Contains the initial conditions where:
7 2 #        - rows are locations
8 3 #        - columns are S, I, R
9 4 ## pars:  Contains the parameters: global values for R0, r, N, eta,
10        berr
11 5 ## T:     The stop time. Since 0 is included, there should be T+1
12        time steps in the simulation
13 6 ## neinum: Number of neighbors for each location, in order
14 7 ## neibmat: Contains lists of neighbors for each location
15 8 #        - rows are parent locations (nodes)
16 9 #        - columns are locations each parent is attached to (edges)
17 10 StocSSIR ← function(ymat, pars, T, steps, neinum, neibmat) {
18 11
19 12     ## number of locations
20 13     nloc ← dim(ymat)[1]
21 14
22 15     ## storage
23 16     ## dims are locations, (S,I,R,B), times
24 17     # output array
25 18     out ← array(NA, c(nloc, 4, T+1), dimnames = list(NULL, c("S", "I",
26 19     "R", "B"), NULL))
27 19     # temp storage
28 20     BSI ← numeric(nloc)
29 21     rI ← numeric(nloc)
30 22
31 23     ## extract parameters
32 24     R0 ← pars[['R0']]
33 25     r ← pars[['r']]
34 26     N ← pars[['N']]
```

```

1 27 eta ← pars[['eta']]
2 28 berr ← pars[['berr']]
3 29 phi ← pars[['phi']]
4 30
5 31 B0 ← rep(R0*r/N, nloc)
6 32
7 33 ## state vectors
8 34 S ← ymat[, 'S']
9 35 I ← ymat[, 'I']
10 36 R ← ymat[, 'R']
11 37 B ← B0
12 38
13 39 ## assign starting to output matrix
14 40 out[, , 1] ← cbind(ymat, B0)
15 41
16 42 h ← 1 / steps
17 43
18 44 for ( i in 1:(T*steps) ) {
19 45
20 46     B ← exp( log(B) + eta*(log(B0) - log(B)) + rnorm(nloc, 0,
21 47         berr) )
22 47
23 48     for (loc in 1:nloc) {
24 49         n ← neinum[loc]
25 50         sphi ← 1 - phi*(n/(n+1))
26 51         ophi ← phi/(n+1)
27 52         nBIsu ← B[neibmat[loc, 1:n]] %*% I[neibmat[loc, 1:n]]
28 53         BSI[loc] ← S[loc]*( sphi*B[loc]*I[loc] + ophi*nBIsu )
29 54     }
30 55
31 56     #if(i == 1)
32 57     # print(BSI)
33 58
34 59     rI ← r*I
35 60
36 61     dS ← -BSI
37 62     dI ← BSI - rI
38 63     dR ← rI
39 64
40 65     S ← S + h*dS
41 66     I ← I + h*dI
42 67     R ← R + h*dR
43 68
44 69     if (i %% steps == 0) {
45 70         out[, , i/steps+1] ← cbind(S, I, R, B)
46 71     }
47 72
48 73 }
49 74
50 75 #out[, , 2] ← cbind(S, I, R, B)
51 76

```

```

1 77 | return(out)
2 78 |
3 79 | }
4 80 |
5 81 | ### Suggested parameters
6 82 | #
7 83 | # T          ← 60
8 84 | # i_infec ← 5
9 85 | # steps     ← 7
10 86 | # N         ← 500
11 87 | # sigma    ← 10
12 88 | #
13 89 | # pars ← c(R0 = 3.0,      # new infected people per infected person
14 90 |           r = 0.1,      # recovery rate
15 91 |           N = 500,      # population size
16 92 |           eta = 0.5,    # geometric random walk
17 93 |           berr = 0.5)   # Beta geometric walk noise

```

19 F.2 RStan Spatial SIR Code

20 This code implements a Spatial SIR model in Rstan.

```

21 1 | data {
22 2 |
23 3 |
24 4 |   int      <lower=1> T;      // total integration steps
25 5 |   int      <lower=1> nloc;   // number of locations
26 6 |   real     y[nloc, T];     // observed number of cases
27 7 |   int      <lower=1> N;     // population size
28 8 |   real     h;              // step size
29 9 |   int      <lower=0> neinum[nloc]; // number of neighbors
30 10 |   each location has
31 11 |   int      neibmat[nloc, nloc]; // neighbor list for
32 12 |   each location
33 13 | }
34 14 |
35 15 | parameters {
36 16 |
37 17 |   real <lower=0, upper=10> R0; // R0
38 18 |   real <lower=0, upper=10> r;  // recovery rate
39 19 |   real <lower=0, upper=20> sigma; // observation error
40 20 |   real <lower=0, upper=30> Iinit[nloc]; // initial
41 21 |   infected for each location
42 22 |   real <lower=0, upper=1> eta; // geometric walk
43 23 |   attraction strength
44 24 |   real <lower=0, upper=1> berr; // beta walk noise
45 25 |   real <lower=-1.5, upper=1.5> Bnoise[nloc, T]; // Beta vector
46 26 |

```

```

1 22      real <lower=0, upper=1>          phi;    // interconnectivity
2      strength
3 23
4 24 }
5 25
6 26 model {
7 27
8 28     real S[nloc, T];
9 29     real I[nloc, T];
10 30     real R[nloc, T];
11 31     real B[nloc, T];
12 32     real B0;
13 33
14 34     real BSI[nloc, T];
15 35     real rI[nloc, T];
16 36     int n;
17 37     real sphi;
18 38     real ophi;
19 39     real nBIsun;
20 40
21 41     B0 ← R0 * r / N;
22 42
23 43     for (loc in 1:nloc) {
24 44         S[loc, 1] ← N - Iinit[loc];
25 45         I[loc, 1] ← Iinit[loc];
26 46         R[loc, 1] ← 0.0;
27 47         B[loc, 1] ← B0;
28 48     }
29 49
30 50     for (t in 2:T) {
31 51         for (loc in 1:nloc) {
32 52
33 53             Bnoise[loc, t] ~ normal(0,berr);
34 54             B[loc, t] ← exp( log(B[loc, t-1]) + eta * ( log(B0) - log
35             (B[loc, t-1]) ) + Bnoise[loc, t] );
36 55
37 56             n ← neinum[loc];
38 57             sphi ← 1.0 - phi*( n/(n+1.0) );
39 58             ophi ← phi/(n+1.0);
40 59
41 60             nBIsun ← 0.0;
42 61             for (j in 1:n)
43 62                 nBIsun ← nBIsun + B[neibmat[loc, j], t-1] * I[neibmat
44                 [loc, j], t-1];
45 63
46 64             BSI[loc, t] ← S[loc, t-1]*( sphi*B[loc, t-1]*I[loc, t-1]
47             + ophi*nBIsun );
48 65             rI[loc, t] ← r*I[loc, t-1];
49 66
50 67             S[loc, t] ← S[loc, t-1] + h*( - BSI[loc, t] );
51 68             I[loc, t] ← I[loc, t-1] + h*( BSI[loc, t] - rI[loc, t] );

```

```

1 69      R[loc, t] ← R[loc, t-1] + h*( rI[loc, t] );
2 70
3 71      if (y[loc, t] > 0) {
4 72          y[loc, t] ~ normal( I[loc, t], sigma );
5 73      }
6 74
7 75  }
8 76  }
9 77
10 78  R0      ~ lognormal(1,1);
11 79  r      ~ lognormal(1,1);
12 80  sigma   ~ lognormal(1,1);
13 81  for (loc in 1:nloc) {
14 82      Iinit[loc] ~ normal(y[loc, 1], sigma);
15 83  }
16 84
17 85 }

```

19 F.3 IF2 Spatial SIR Code

20 This code implements a Spatial SIR model using IF2 in C++.

```

21
22 1  /* Author: Dexter Barrows
23 2   Github: dbarrows.github.io
24 3
25 4   */
26 5
27 6 #include <stdio.h>
28 7 #include <math.h>
29 8 #include <sys/time.h>
30 9 #include <time.h>
31 10 #include <stdlib.h>
32 11 #include <string>
33 12 #include <cmath>
34 13 #include <cstdlib>
35 14 #include <fstream>
36 15
37 16 // #include "rand.h"
38 17 // #include "timer.h"
39 18
40 19 #define Treal      100          // time to simulate over
41 20 #define R0true     3.0          // infectiousness
42 21 #define rtrue      0.1          // recovery rate
43 22 #define Nreal      500.0        // population size
44 23 #define etatrue     0.5          // real drift attraction strength
45 24 #define berrtrue    0.5          // real beta drift noise
46 25 #define phitrue     0.5          // real connectivity strength
47 26 #define merr        10.0        // expected measurement error

```

```

1 27 #define I0          5.0          // Initial infected individuals
2 28
3 29 #define PSC          0.5          // perturbation scale factor for more
4      sensitive parameters
5 30
6 31 #include <Rcpp.h>
7 32 using namespace Rcpp;
8 33
9 34 struct Particle {
10 35     double R0;
11 36     double r;
12 37     double sigma;
13 38     double eta;
14 39     double berr;
15 40     double phi;
16 41     double * S;
17 42     double * I;
18 43     double * R;
19 44     double * B;
20 45     double * Iinit;
21 46 };
22 47
23 48
24 49 int timeval_subtract (double *result, struct timeval *x, struct
25     timeval *y);
26 50 int check_double(double x,double y);
27 51 void initializeParticles(Particle ** particles, int NP, int nloc, int
28     N);
29 52 void exp_euler_SSIR(double h, double t0, double tn, int N, Particle *
30     particle,
31     NumericVector neinum, NumericMatrix neibmat, int
32     nloc) ;
33 54 void copyParticle(Particle * dst, Particle * src, int nloc);
34 55 void perturbParticles(Particle * particles, int N, int NP, int nloc,
35     int passnum, double coolrate);
36 56 double randu();
37 57 double randn();
38 58
39 59 // [[Rcpp::export]]
40 60 Rcpp::List if2_spa(NumericMatrix data, int T, int N, int NP, int
41     nPasses, double coolrate, NumericVector neinum, NumericMatrix
42     neibmat, int nloc) {
43 61
44 62     NumericMatrix paramdata(NP, 6);          // for R0, r, sigma, eta,
45     berr, phi
46 63     NumericMatrix initInfec(nloc, NP); // for Iinit
47 64     NumericMatrix infecmeans(nloc, T); // mean infection counts for
48     each location
49 65     NumericMatrix finalstate(nloc, 4); // SIRB means for each
50     location
51 66

```



```

1 67      srand(time(NULL));          // Seed PRNG with system time
2 68
3 69      double w[NP];               // particle weights
4 70
5 71      // initialize particles
6 72      printf("Initializing particle states\n");
7 73      Particle * particles = NULL; // particle estimates for
8      current step
9 74      Particle * particles_old = NULL; // intermediate particle
10      states for resampling
11 75      initializeParticles(&particles, NP, nloc, N);
12 76      initializeParticles(&particles_old, NP, nloc, N);
13 77
14 78      /*
15 79      // copy particle test
16 80      copyParticle(&particles[0], &particles_old[0], nloc);
17 81
18 82      // perturb particle test
19 83      perturbParticles(particles, N, NP, nloc, 1, coolrate);
20 84
21 85      // evolution test
22 86      // reset particle system evolution states
23 87      for (int n = 0; n < NP; n++) {
24 88          for (int loc = 0; loc < nloc; loc++) {
25 89              particles[n].S[loc] = N - particles[n].Iinit[loc];
26 90              particles[n].I[loc] = particles[n].Iinit[loc];
27 91              particles[n].R[loc] = 0.0;
28 92              particles[n].B[loc] = (double) particles[n].R0 *
29              particles[n].r / N;
30 93          }
31 94      }
32 95      printf("Before S:%f | I:%f | R:%f\n", particles[0].S[0],
33      particles[0].I[0], particles[0].R[0]);
34 96      exp_euler_SSIR(1.0/7.0, 0.0, 1.0, N, &particles[0], neinum,
35      neibmat, nloc);
36 97      printf("After S:%f | I:%f | R:%f\n", particles[0].S[0], particles
37      [0].I[0], particles[0].R[0]);
38 98      */
39 99
40 100     // START PASSES THROUGH DATA
41 101
42 102     printf("Starting filter\n");
43 103     printf("-----\n");
44 104     printf("Pass\n");
45 105
46 106
47 107     for (int pass = 0; pass < nPasses; pass++) {
48 108
49 109         printf("...%d / %d\n", pass, nPasses);
50 110
51 111         // reset particle system evolution states

```

```

1 112     for (int n = 0; n < NP; n++) {
2 113         for (int loc = 0; loc < nloc; loc++) {
3 114             particles[n].S[loc] = N - particles[n].Iinit[loc];
4 115             particles[n].I[loc] = particles[n].Iinit[loc];
5 116             particles[n].R[loc] = 0.0;
6 117             particles[n].B[loc] = (double) particles[n].R0 *
7             particles[n].r / N;
8 118         }
9 119     }
10 120
11 121     if (pass == (nPasses-1)) {
12 122         double means[nloc];
13 123         for (int loc = 0; loc < nloc; loc++) {
14 124             means[loc] = 0.0;
15 125             for (int n = 0; n < NP; n++) {
16 126                 means[loc] += particles[n].I[loc] / NP;
17 127             }
18 128             infecmeans(loc, 0) = means[loc];
19 129         }
20 130     }
21 131
22 132     for (int t = 1; t < T; t++) {
23 133
24 134         // generate individual predictions and weight
25 135         for (int n = 0; n < NP; n++) {
26 136
27 137             exp_euler_SSIR(1.0/7.0, 0.0, 1.0, N, &particles[n],
28             neinum, neibmat, nloc);
29 138
30 139             double merr_par = particles[n].sigma;
31 140
32 141             w[n] = 1.0;
33 142             for (int loc = 0; loc < nloc; loc++) {
34 143                 double y_diff = data(loc, t) - particles[n].I[
35                 loc];
36 144                 w[n] *= 1.0/(merr_par*sqrt(2.0*M_PI)) * exp( -
37                 y_diff*y_diff / (2.0*merr_par*merr_par) );
38 145             }
39 146
40 147         }
41 148
42 149         // cumulative sum
43 150         for (int n = 1; n < NP; n++) {
44 151             w[n] += w[n-1];
45 152         }
46 153
47 154         // save particle states to resample from
48 155         for (int n = 0; n < NP; n++){
49 156             copyParticle(&particles_old[n], &particles[n], nloc);
50 157         }
51 158

```

```

1 159      // resampling
2 160      for (int n = 0; n < NP; n++) {
3 161
4 162          double w_r = randu() * w[NP-1];
5 163          int i = 0;
6 164          while (w_r > w[i]) {
7 165              i++;
8 166          }
9 167
10 168          // i is now the index to copy state from
11 169          copyParticle(&particles[n], &particles_old[i], nloc);
12 170
13 171      }
14 172
15 173      // between-iteration perturbations, not after last time
16      step
17 174      if (t < (T-1))
18 175          perturbParticles(particles, N, NP, nloc, pass,
19      coolrate);
20 176
21 177      if (pass == (nPasses-1)) {
22 178          double means[nloc];
23 179          for (int loc = 0; loc < nloc; loc++) {
24 180              means[loc] = 0.0;
25 181              for (int n = 0; n < NP; n++) {
26 182                  means[loc] += particles[n].I[loc] / NP;
27 183              }
28 184              infecmeans(loc, t) = means[loc];
29 185          }
30 186      }
31 187
32 188  }
33 189
34 190      // between-pass perturbations, not after last pass
35 191      if (pass < (nPasses + 1))
36 192          perturbParticles(particles, N, NP, nloc, pass, coolrate);
37 193
38 194  }
39 195
40 196      // pack parameter data (minus initial conditions)
41 197      for (int n = 0; n < NP; n++) {
42 198          paramdata(n, 0) = particles[n].R0;
43 199          paramdata(n, 1) = particles[n].r;
44 200          paramdata(n, 2) = particles[n].sigma;
45 201          paramdata(n, 3) = particles[n].eta;
46 202          paramdata(n, 4) = particles[n].berr;
47 203          paramdata(n, 5) = particles[n].phi;
48 204      }
49 205
50 206      // Pack initial condition data
51 207      for (int n = 0; n < NP; n++) {

```

```

1 208     for (int loc = 0; loc < nloc; loc++) {
2 209         initInfec(loc, n) = particles[n].Iinit[loc];
3 210     }
4 211 }
5 212
6 213 // Pack final state means data
7 214 double Smeans[nloc], Imeans[nloc], Rmeans[nloc], Bmeans[nloc];
8 215 for (int loc = 0; loc < nloc; loc++) {
9 216     Smeans[loc] = 0.0;
10 217     Imeans[loc] = 0.0;
11 218     Rmeans[loc] = 0.0;
12 219     Bmeans[loc] = 0.0;
13 220     for (int n = 0; n < NP; n++) {
14 221         Smeans[loc] += particles[n].S[loc] / NP;
15 222         Imeans[loc] += particles[n].I[loc] / NP;
16 223         Rmeans[loc] += particles[n].R[loc] / NP;
17 224         Bmeans[loc] += particles[n].B[loc] / NP;
18 225     }
19 226     finalstate(loc, 0) = Smeans[loc];
20 227     finalstate(loc, 1) = Imeans[loc];
21 228     finalstate(loc, 2) = Rmeans[loc];
22 229     finalstate(loc, 3) = Bmeans[loc];
23 230 }
24 231
25 232
26 233 return Rcpp::List::create( Rcpp::Named("paramdata") = paramdata,
27 234                           Rcpp::Named("initInfec") = initInfec,
28 235                           Rcpp::Named("infecmeans") =
29                               infecmeans,
30 236                           Rcpp::Named("finalstate") =
31                               finalstate);
32 237
33 238
34 239 }
35 240 }
36 241
37 242
38 243 /* Use the Explicit Euler integration scheme to integrate SIR model
39     forward in time
40 244     double h      - time step size
41 245     double t0     - start time
42 246     double tn     - stop time
43 247     double * y    - current system state; a three-component vector
44                     representing [S I R], susceptible-infected-recovered
45 248
46 249 */
47 250 void exp_euler_SSIR(double h, double t0, double tn, int N, Particle *
48     particle,
49 251                     NumericVector neinum, NumericMatrix neibmat, int
50                     nloc) {
51 252

```

```

1 253     int num_steps = floor( (tn-t0) / h );
2 254
3 255     double * S = particle->S;
4 256     double * I = particle->I;
5 257     double * R = particle->R;
6 258     double * B = particle->B;
7 259
8 260     // create last state vectors
9 261     double S_last[nloc];
10 262     double I_last[nloc];
11 263     double R_last[nloc];
12 264     double B_last[nloc];
13 265
14 266     double R0    = particle->R0;
15 267     double r     = particle->r;
16 268     double B0    = R0 * r / N;
17 269     double eta   = particle->eta;
18 270     double berr  = particle->berr;
19 271     double phi   = particle->phi;
20 272
21 273     //printf("sphi \t\t| ophi \t\t| BSI \t\t| rI \t\t| dS \t\t| dI \t\t|
22 274     \t| dR \t\t| S \t\t| I \t\t| R |\n");
23 275
24 276     for(int t = 0; t < num_steps; t++) {
25 277
26 278         for (int loc = 0; loc < nloc; loc++) {
27 279             S_last[loc] = S[loc];
28 280             I_last[loc] = I[loc];
29 281             R_last[loc] = R[loc];
30 282             B_last[loc] = B[loc];
31 283         }
32 284
33 285         for (int loc = 0; loc < nloc; loc++) {
34 286
35 287             B[loc] = exp( log(B_last[loc]) + eta*(log(B0) - log(
36 288                 B_last[loc])) + berr*randn() );
37 289
38 290             int n = neinum[loc];
39 291             double sphi = 1.0 - phi*( (double) n/(n+1.0) );
40 292             double ophi = phi/(n+1.0);
41 293
42 294             double nBIsum = 0.0;
43 295             for (int j = 0; j < n; j++)
44 296                 nBIsum += B_last[(int) neibmat(loc, j) - 1] * I_last
45 297                     [(int) neibmat(loc, j) - 1];
46 298
47 299             double BSI = S_last[loc]*( sphi*B_last[loc]*I_last[loc] +
48 300                 ophi*nBIsum );
49 301             double rI  = r*I_last[loc];
50 302
51 303             // get derivatives

```

```

1 300         double dS = - BSI;
2 301         double dI = BSI - rI;
3 302         double dR = rI;
4 303
5 304         // step forward by h
6 305         S[loc] += h*dS;
7 306         I[loc] += h*dI;
8 307         R[loc] += h*dR;
9 308
10 309         //if (loc == 1)
11 310         // printf("%f\t%f\t%f\t%f\t%f\t%f\t%f\t%f\t%f\t"
12         //         "%f\t\n", sphi, ophi, BSI, rI, dS, dI, dR, S[1], I
13         //         [1], R[1]);
14 311
15 312     }
16 313
17 314 }
18 315
19 316 /*particle->S = S;
20 317 particle->I = I;
21 318 particle->R = R;
22 319 particle->B = B;*/
23 320
24 321 }
25 322
26 323 /*  Initializes particles
27 324 */
28 325 void initializeParticles(Particle ** particles, int NP, int nloc, int
29 326 N) {
30 327
31 328     // allocate space for doubles
32 329     *particles = (Particle*) malloc (NP*sizeof(Particle));
33 330
34 331     // allocate space for arrays inside particles
35 332     for (int n = 0; n < NP; n++) {
36 333         (*particles)[n].S = (double*) malloc(nloc*sizeof(double));
37 334         (*particles)[n].I = (double*) malloc(nloc*sizeof(double));
38 335         (*particles)[n].R = (double*) malloc(nloc*sizeof(double));
39 336         (*particles)[n].B = (double*) malloc(nloc*sizeof(double));
40 337         (*particles)[n].Iinit = (double*) malloc(nloc*sizeof(double))
41 338         ;
42 339     }
43 340
44 341     // initialize all all parameters
45 342     for (int n = 0; n < NP; n++) {
46 343
47 344         double R0can, rcan, sigmacan, Iinitcan, etacan, berrcan,
48 345         phican;
49 346
50 347         do {
51 348             R0can = R0true + R0true*randn();

```

```

1 346     } while (R0can < 0);
2 347     (*particles)[n].R0 = R0can;
3 348
4 349     do {
5 350         rcan = rtrue + rtrue*randn();
6 351     } while (rcan < 0);
7 352     (*particles)[n].r = rcan;
8 353
9 354     for (int loc = 0; loc < nloc; loc++)
10 355         (*particles)[n].B[loc] = (double) R0can * rcan / N;
11 356
12 357     do {
13 358         sigma = merr + merr*randn();
14 359     } while (sigma < 0);
15 360     (*particles)[n].sigma = sigma;
16 361
17 362     do {
18 363         etacan = etatrue + PSC*etatrue*randn();
19 364     } while (etacan < 0 || etacan > 1);
20 365     (*particles)[n].eta = etacan;
21 366
22 367     do {
23 368         berrcan = berrtrue + PSC*berrtrue*randn();
24 369     } while (berrcan < 0);
25 370     (*particles)[n].berr = berrcan;
26 371
27 372     do {
28 373         phican = phitrue + PSC*phitrue*randn();
29 374     } while (phican <= 0 || phican >= 1);
30 375     (*particles)[n].phi = phican;
31 376
32 377     for (int loc = 0; loc < nloc; loc++) {
33 378         do {
34 379             Iinitcan = I0 + I0*randn();
35 380             } while (Iinitcan < 0 || N < Iinitcan);
36 381             (*particles)[n].Iinit[loc] = Iinitcan;
37 382         }
38 383
39 384     }
40 385
41 386 }
42 387
43 388 /* Particle pertubation function to be run between iterations and
44     passes
45 389
46 390 */
47 391 void perturbParticles(Particle * particles, int N, int NP, int nloc,
48     int passnum, double coolrate) {
49 392
50 393     //double coolcoef = exp( - (double) passnum / coolrate );
51 394     double coolcoef = pow(coolrate, passnum);

```

```

1 395
2 396 double spreadR0      = coolcoef * R0true / 10.0;
3 397 double spreadr       = coolcoef * rtrue / 10.0;
4 398 double spreadsigma   = coolcoef * merr / 10.0;
5 399 double spreadIinit   = coolcoef * I0 / 10.0;
6 400 double spreadeta     = coolcoef * etatrue / 10.0;
7 401 double spreadberr    = coolcoef * berrtrue / 10.0;
8 402 double spreadphi     = coolcoef * phitrue / 10.0;
9 403
10 404 double R0can, rcan, sigmacan, Iinitcan, etacan, berrcan, phican;
11 405
12 406 for (int n = 0; n < NP; n++) {
13 407
14 408     do {
15 409         R0can = particles[n].R0 + spreadR0*randn();
16 410     } while (R0can < 0);
17 411     particles[n].R0 = R0can;
18 412
19 413     do {
20 414         rcan = particles[n].r + spreadr*randn();
21 415     } while (rcan < 0);
22 416     particles[n].r = rcan;
23 417
24 418     do {
25 419         sigmacan = particles[n].sigma + spreadsigma*randn();
26 420     } while (sigmacan < 0);
27 421     particles[n].sigma = sigmacan;
28 422
29 423     do {
30 424         etacan = particles[n].eta + PSC*spreadeta*randn();
31 425     } while (etacan < 0 || etacan > 1);
32 426     particles[n].eta = etacan;
33 427
34 428     do {
35 429         berrcan = particles[n].berr + PSC*spreadberr*randn();
36 430     } while (berrcan < 0);
37 431     particles[n].berr = berrcan;
38 432
39 433     do {
40 434         phican = particles[n].phi + PSC*spreadphi*randn();
41 435     } while (phican <= 0 || phican >= 1);
42 436     particles[n].phi = phican;
43 437
44 438     for (int loc = 0; loc < nloc; loc++) {
45 439         do {
46 440             Iinitcan = particles[n].Iinit[loc] + spreadIinit*
47 441                 randn();
48 441         } while (Iinitcan < 0 || Iinitcan > 500);
49 442         particles[n].Iinit[loc] = Iinitcan;
50 443     }
51 444 }

```



```

1 445 |
2 446 | }
3 447 |
4 448 | /* Convenience function for particle resampling process
5 449 | */
6 450 | void copyParticle(Particle * dst, Particle * src, int nloc) {
7 451 |
8 452 |     dst->R0      = src->R0;
9 453 |     dst->r        = src->r;
10 454 |     dst->sigma    = src->sigma;
11 455 |     dst->eta      = src->eta;
12 456 |     dst->berr     = src->berr;
13 457 |     dst->phi      = src->phi;
14 458 |
15 459 |     for (int n = 0; n < nloc; n++) {
16 460 |         dst->S[n]      = src->S[n];
17 461 |         dst->I[n]      = src->I[n];
18 462 |         dst->R[n]      = src->R[n];
19 463 |         dst->B[n]      = src->B[n];
20 464 |         dst->Iinit[n]  = src->Iinit[n];
21 465 |     }
22 466 |
23 467 | }
24 468 |
25 469 |
26 470 |
27 471 | double randu() {
28 472 |
29 473 |     return (double) rand() / (double) RAND_MAX;
30 474 |
31 475 | }
32 476 |
33 477 | /*
34 478 | void getStateMeans(State * state, Particle* particles, int NP) {
35 479 |
36 480 |     double Smean = 0, Imean = 0, Rmean = 0;
37 481 |
38 482 |     for (int n = 0; n < NP; n++) {
39 483 |         Smean += particles[n].S;
40 484 |         Imean += particles[n].I;
41 485 |         Rmean += particles[n].R;
42 486 |     }
43 487 |
44 488 |     state->S = (double) Smean / NP;
45 489 |     state->I = (double) Imean / NP;
46 490 |     state->R = (double) Rmean / NP;
47 491 |
48 492 | }
49 493 | */
50 494 |
51 495 | /* Return a normally distributed random number with mean 0 and

```

```

1      standard deviation 1
2 496  Uses the polar form of the Box-Muller transformation
3 497  From http://www.design.caltech.edu/erik/Misc/Gaussian.html
4 498  */
5 499 double randn() {
6 500
7 501     double x1, x2, w, y1;
8 502
9 503     do {
10 504         x1 = 2.0 * randu() - 1.0;
11 505         x2 = 2.0 * randu() - 1.0;
12 506         w = x1 * x1 + x2 * x2;
13 507     } while ( w >= 1.0 );
14 508
15 509     w = sqrt( (-2.0 * log( w ) ) / w );
16 510     y1 = x1 * w;
17 511
18 512     return y1;
19 513
20 514 }

```

22 F.4 CUDA IF2 Spatial Fitting Code

23 Below is the nascent CUDA code that will be expanded upon in future work. At
 24 present it only implements the core IF2 fitting algorithm and does not implement
 25 parametric bootstrapping nor produce forecasts.

```

26
27 1 /* Author: Dexter Barrows
28 2   Github: dbarrows.github.io
29 3   */
30 4
31 5 /* Runs a particle filter on synthetic noisy data and attempts to
32 6    reconstruct underlying true state at each time step. Note that
33 7    this program uses gnuplot to plot the data, so an x11
34 8    environment must be present. Also the multiplier of 1024 in the
35 9    definition of NP below should be set to a multiple of the number
36 10   of multiprocessors of your GPU for optimal results.
37 11
38 12   Also, the accompanying "pf.plg" file contains the instructions
39 13   gnuplot will use. It must be present in the same directory as
40 14   the executable generated by compiling this file.
41 15
42 16   Compile with:
43 17
44 18   nvcc -arch=sm_20 -O2 pf_cuda.cu timer.cpp rand.cpp -o pf_cuda.x
45 19
46 20   */
47 21

```

```

1 22 #include <cuda.h>
2 23 #include <iostream>
3 24 #include <fstream>
4 25 #include <curand.h>
5 26 #include <curand_kernel.h>
6 27 #include <string>
7 28 #include <sstream>
8 29 #include <cmath>
9 30
10 31 #include "timer.h"
11 32 #include "rand.h"
12 33 #include "readdata.h"
13 34
14 35 #define NP          (2*2500)    // number of particles
15 36 #define N           500.0      // population size
16 37 #define R0true      3.0        // infectiousness
17 38 #define rtrue       0.1        // recovery rate
18 39 #define etatrue     0.5        // real drift attraction strength
19 40 #define berrtrue    0.5        // real beta drift noise
20 41 #define phitrue     0.5        // real connectivity strength
21 42 #define merr        10.0       // expected measurement error
22 43 #define I0          5.0        // Initial infected individuals
23 44 #define PSC         0.5        // sensitive parameter perturbation
24      scaling
25 45 #define NLOC        10
26 46
27 47 #define PI          3.141592654f
28 48
29 49 // Wrapper for CUDA calls, from CUDA API
30 50 // Modified to also print the error code and string
31 51 # define CUDA_CALL(x) do { if ((x) != cudaSuccess ) {
32      \
33 52      std::cout << " Error at " << __FILE__ << ":" << __LINE__ << std::
34      endl;      \
35 53      std::cout << " Error was " << x << " " << cudaGetErrorString(x)
36      << std::endl;      \
37 54      return EXIT_FAILURE ;}} while (0)
38      \
39 55
40 56 typedef struct {
41 57     float R0;
42 58     float r;
43 59     float sigma;
44 60     float eta;
45 61     float berr;
46 62     float phi;
47 63     float S[NLOC];
48 64     float I[NLOC];
49 65     float R[NLOC];
50 66     float B[NLOC];
51 67     float Iinit[NLOC];

```

```

1 68     curandState randState; // PRNG state
2 69 } Particle;
3 70
4 71 __host__ std::string getHRmemsize (size_t memsize);
5 72 __host__ std::string getHRtime (float runtime);
6 73
7 74 __device__ void exp_euler_SSIR(float h, float t0, float tn, Particle
8     * particle, int * neinum, int * neibmat, int nloc);
9 75 __device__ void copyParticle(Particle * dst, Particle * src, int nloc
10    );
11 76
12 77
13 78 /* Initialize all PRNG states, get starting state vector using
14     initial distribution
15     */
16 80 __global__ void initializeParticles (Particle * particles, int nloc)
17    {
18 81
19 82     int id = blockIdx.x*blockDim.x + threadIdx.x; // global thread
20        ID
21 83
22 84     if (id < NP) {
23 85
24 86         // initialize PRNG state
25 87         curandState state;
26 88         curand_init(id, 0, 0, &state);
27 89
28 90         float R0can, rcan, sigmacan, Iinitcan, etacan, berrcan,
29            phican;
30 91
31 92         do {
32 93             R0can = R0true + R0true*curand_normal(&state);
33 94         } while (R0can < 0);
34 95         particles[id].R0 = R0can;
35 96
36 97         do {
37 98             rcan = rtrue + rtrue*curand_normal(&state);
38 99         } while (rcan < 0);
39 100         particles[id].r = rcan;
40 101
41 102         for (int loc = 0; loc < nloc; loc++)
42 103             particles[id].B[loc] = (float) R0can * rcan / N;
43 104
44 105         do {
45 106             sigmacan = merr + merr*curand_normal(&state);
46 107         } while (sigmacan < 0);
47 108         particles[id].sigma = sigmacan;
48 109
49 110         do {
50 111             etacan = etatrue + PSC*etatrue*curand_normal(&state);
51 112         } while (etacan < 0 || etacan > 1);

```

```

1 113     particles[id].eta = etacan;
2 114
3 115     do {
4 116         berrcan = berrtrue + PSC*berrtrue*curand_normal(&state);
5 117     } while (berrcan < 0);
6 118     particles[id].berr = berrcan;
7 119
8 120     do {
9 121         phican = phitrue + PSC*phitrue*curand_normal(&state);
10 122     } while (phican <= 0 || phican >= 1);
11 123     particles[id].phi = phican;
12 124
13 125     for (int loc = 0; loc < nloc; loc++) {
14 126         do {
15 127             Iinitcan = I0 + I0*curand_normal(&state);
16 128         } while (Iinitcan < 0 || N < Iinitcan);
17 129         particles[id].Iinit[loc] = Iinitcan;
18 130     }
19 131
20 132     particles[id].randState = state;
21 133
22 134 }
23 135
24 136 }
25 137
26 138 __global__ void resetStates (Particle * particles, int nloc) {
27 139
28 140     int id = blockIdx.x*blockDim.x + threadIdx.x; // global thread
29 141     ID
30 141
31 142     if (id < NP) {
32 143
33 144         for (int loc = 0; loc < nloc; loc++) {
34 145             particles[id].S[loc] = N - particles[id].Iinit[loc];
35 146             particles[id].I[loc] = particles[id].Iinit[loc];
36 147             particles[id].R[loc] = 0.0;
37 148         }
38 149
39 150     }
40 151
41 152 }
42 153
43 154 __global__ void clobberParams (Particle * particles, int nloc) {
44 155
45 156     int id = blockIdx.x*blockDim.x + threadIdx.x; // global thread
46 157     ID
47 157
48 158     if (id < NP) {
49 159
50 160         particles[id].R0 = R0true;
51 161         particles[id].r = rtrue;

```

```

1 162     particles[id].sigma = merr;
2 163     particles[id].eta = etatrue;
3 164     particles[id].berr = berrtrue;
4 165     particles[id].phi = phitrue;
5 166
6 167     for (int loc = 0; loc < nloc; loc++) {
7 168         particles[id].Iinit[loc] = I0;
8 169     }
9 170
10 171 }
11 172
12 173 }
13 174
14 175
15 176 /* Project particles forward, perturb, and save weight based on data
16 177    int t - time step number (1,...,T)
17 178    */
18 179 __global__ void project (Particle * particles, int * neinum, int *
19 180    neibmat, int nloc) {
20 181
21 181     int id = blockIdx.x*blockDim.x + threadIdx.x;    // global id
22 182
23 183     if (id < NP) {
24 184         // project forward
25 185         exp_euler_SSIR(1.0/7.0, 0.0, 1.0, &particles[id], neinum,
26 186         neibmat, nloc);
27 186     }
28 187
29 188 }
30 189
31 190 __global__ void weight(float * data, Particle * particles, double * w
32 191    , int t, int T, int nloc) {
33 191
34 192     int id = blockIdx.x*blockDim.x + threadIdx.x;    // global id
35 193
36 194     if (id < NP) {
37 195
38 196         float merr_par = particles[id].sigma;
39 197
40 198         // Get weight and save
41 199         double w_local = 1.0;
42 200         for (int loc = 0; loc < nloc; loc++) {
43 201             float y_diff = data[loc*T + t] - particles[id].I[loc];
44 202             w_local *= 1.0/(merr_par*sqrt(2.0*PI)) * exp( - y_diff*
45 203             y_diff / (2.0*merr_par*merr_par) );
46 203         }
47 204
48 205         w[id] = w_local;
49 206
50 207     }
51 208

```

```

1 209 }
2 210
3 211 __global__ void stashParticles (Particle * particles, Particle *
4 212   particles_old, int nloc) {
5 213
6 214     int id = blockIdx.x*blockDim.x + threadIdx.x;    // global id
7 215
8 216     if (id < NP) {
9 217         // COPY PARTICLE
10 218         copyParticle(&particles_old[id], &particles[id], nloc);
11 219     }
12 220 }
13 221
14 222
15 223 /* The 0th thread will perform cumulative sum on the weights.
16 224    There may be a faster way to do this, will investigate.
17 225    */
18 226 __global__ void cumsumWeights (double * w) {
19 227
20 228     int id = blockIdx.x*blockDim.x + threadIdx.x;    // global thread
21 229     ID
22 230
23 231     // compute cumulative weights
24 232     if (id == 0) {
25 233         for (int i = 1; i < NP; i++)
26 234             w[i] += w[i-1];
27 235     }
28 236 }
29 237
30 238
31 239 /* Resample from all particle states within cell
32 240    */
33 241 __global__ void resample (Particle * particles, Particle *
34 242   particles_old, double * w, int nloc) {
35 243
36 244     int id = blockIdx.x*blockDim.x + threadIdx.x;
37 245
38 246     if (id < NP) {
39 247
40 248         // resampling proportional to weights
41 249         double w_r = curand_uniform(&particles[id].randState) * w[NP
42 250         -1];
43 251         int i = 0;
44 252         while (w_r > w[i]) {
45 253             i++;
46 254         }
47 255
48 256         // i is now the index of the particle to copy from
49 257         copyParticle(&particles[id], &particles_old[i], nloc);

```

```

1 256
2 257     }
3 258
4 259 }
5 260
6 261 // launch this with probably just nloc threads... block structure/
7 262 // size probably not important
8 262 __global__ void reduceStates (Particle * particles, float *
9 263 countmeans, int t, int T, int nloc) {
10 263
11 264     int id = blockIdx.x*blockDim.x + threadIdx.x;
12 265
13 266     if (id < nloc) {
14 267
15 268         int loc = id;
16 269
17 270         double countmean_local = 0.0;
18 271         for (int n = 0; n < NP; n++) {
19 272             countmean_local += particles[n].I[loc] / NP;
20 273         }
21 274
22 275         countmeans[loc*T + t] = (float) countmean_local;
23 276
24 277     }
25 278
26 279 }
27 280
28 281 __global__ void perturbParticles(Particle * particles, int nloc, int
29 282 passnum, double coolrate) {
30 282
31 283     //double coolcoef = exp( - (double) passnum / coolrate );
32 284     double coolcoef = pow(coolrate, passnum);
33 285
34 286     double spreadR0 = coolcoef * R0true / 10.0;
35 287     double spreadr = coolcoef * rtrue / 10.0;
36 288     double spreadsigma = coolcoef * merr / 10.0;
37 289     double spreadIinit = coolcoef * I0 / 10.0;
38 290     double spreadeta = coolcoef * etatrue / 10.0;
39 291     double spreadberr = coolcoef * berrtrue / 10.0;
40 292     double spreadphi = coolcoef * phitrue / 10.0;
41 293
42 294     double R0can, rcan, sigmacan, Iinitcan, etacan, berrcan, phican;
43 295
44 296     int id = blockIdx.x*blockDim.x + threadIdx.x;
45 297
46 298     if (id < NP) {
47 299
48 300         do {
49 301             R0can = particles[id].R0 + spreadR0*curand_normal(&
50 302                 particles[id].randState);
51 302             } while (R0can < 0);

```



```

1 303     particles[id].R0 = R0can;
2 304
3 305     do {
4 306         rcan = particles[id].r + spreadr*curand_normal(&particles
5         [id].randState);
6 307     } while (rcan < 0);
7 308     particles[id].r = rcan;
8 309
9 310     do {
10 311         sigmacan = particles[id].sigma + spreadsigma*
11         curand_normal(&particles[id].randState);
12 312     } while (sigmacan < 0);
13 313     particles[id].sigma = sigmacan;
14 314
15 315     do {
16 316         etacan = particles[id].eta + PSC*spreadeta*curand_normal
17         (&particles[id].randState);
18 317     } while (etacan < 0 || etacan > 1);
19 318     particles[id].eta = etacan;
20 319
21 320     do {
22 321         berrcan = particles[id].berr + PSC*spreadberr*
23         curand_normal(&particles[id].randState);
24 322     } while (berrcan < 0);
25 323     particles[id].berr = berrcan;
26 324
27 325     do {
28 326         phican = particles[id].phi + PSC*spreadphi*curand_normal
29         (&particles[id].randState);
30 327     } while (phican <= 0 || phican >= 1);
31 328     particles[id].phi = phican;
32 329
33 330     for (int loc = 0; loc < nloc; loc++) {
34 331         do {
35 332             Iinitcan = particles[id].Iinit[loc] + spreadIinit*
36             curand_normal(&particles[id].randState);
37 333         } while (Iinitcan < 0 || Iinitcan > 500);
38 334         particles[id].Iinit[loc] = Iinitcan;
39 335     }
40 336
41 337 }
42 338
43 339 }
44 340
45 341
46 342 int main (int argc, char *argv[]) {
47 343
48 344
49 345     int T, nloc;
50 346
51 347     double restime;

```

```

1 348     struct timeval tdr0, tdr1, tdrMaster;
2 349
3 350     // Parse arguments *****
4 351
5 352     if (argc < 4) {
6 353         std::cout << "Not enough arguments" << std::endl;
7 354         return 0;
8 355     }
9 356
10 357     std::string arg1(argv[1]); // infection counts
11 358     std::string arg2(argv[2]); // neighbour counts
12 359     std::string arg3(argv[3]); // neighbour indices
13 360     std::string arg4(argv[4]); // outfile: params + runtime
14 361
15 362     std::cout << "Arguments:" << std::endl;
16 363     std::cout << "Infection data: " << arg1 << std::endl;
17 364     std::cout << "Neighbour counts: " << arg2 << std::endl;
18 365     std::cout << "Neighbour indices: " << arg3 << std::endl;
19 366     std::cout << "Outfile " << arg4 << std::endl;
20 367
21 368     // *****
22 369
23 370
24 371     // Read count data *****
25 372
26 373     std::cout << "Getting count data" << std::endl;
27 374     float * data = getDataFloat(arg1, &T, &nloc);
28 375     size_t datasize = nloc*T*sizeof(float);
29 376
30 377     // *****
31 378
32 379     // Read neinum matrix data *****
33 380
34 381     std::cout << "Getting neighbour count data" << std::endl;
35 382     int * neinum = getDataInt(arg2, NULL, NULL);
36 383     size_t neinumsize = nloc * sizeof(int);
37 384
38 385     // *****
39 386
40 387     // Read neibmat matrix data *****
41 388
42 389     std::cout << "Getting neighbour count data" << std::endl;
43 390     int * neibmat = getDataInt(arg3, NULL, NULL);
44 391     size_t neibmatsize = nloc * nloc * sizeof(int);
45 392
46 393     // *****
47 394
48 395     //
49
50
51 396     *****

```

```

1 397 // start timing
2 398 gettimeofday (&tdr0, NULL);
3 399
4 400 // CUDA data *****
5 401
6 402 std::cout << "Allocating device storage" << std::endl;
7 403
8 404 float * d_data; // device copy of data
9 405 Particle * particles; // particles
10 406 Particle * particles_old; // intermediate particle states
11 407 double * w; // weights
12 408 int * d_neinum; // device copy of adjacency
13 // matrix
14 409 int * d_neibmat; // device copy of neighbour
15 // counts matrix
16 410 float * countmeans; // host copy of reduced infection
17 // count means from last pass
18 411 float * d_countmeans; // device copy of reduced
19 // infection count means from last pass
20 412
21 413 CUDA_CALL( cudaMalloc( (void**) &d_data , datasize )
22 );
23 414 CUDA_CALL( cudaMalloc( (void**) &particles , NP*sizeof(
24 Particle)) );
25 415 CUDA_CALL( cudaMalloc( (void**) &particles_old , NP*sizeof(
26 Particle)) );
27 416 CUDA_CALL( cudaMalloc( (void**) &w , NP*sizeof(
28 double)) );
29 417 CUDA_CALL( cudaMalloc( (void**) &d_neinum , neinumsize)
30 );
31 418 CUDA_CALL( cudaMalloc( (void**) &d_neibmat , neibmatsize)
32 );
33 419 CUDA_CALL( cudaMalloc( (void**) &d_countmeans , nloc*T*sizeof(
34 float)) );
35 420
36 421
37 422 gettimeofday (&tdr1, NULL);
38 423 timeval_subtract (&restime, &tdr1, &tdr0);
39 424
40 425 std::cout << "\t" << getHRtime(restime) << std::endl;
41 426
42 427 size_t avail, total;
43 428 cudaMemGetInfo( &avail, &total );
44 429 size_t used = total - avail;
45 430
46 431 std::cout << "\t[" << getHRmemsize(used) << "]" used of [" <<
47 getHRmemsize(total) << "]" <<std::endl;
48 432
49 433 std::cout << "Copying data to device" << std::endl;
50 434
51 435 gettimeofday (&tdr0, NULL);

```

```

1 436
2 437   CUDA_CALL( cudaMemcpy(d_data      , data      , datasize      ,
3      cudaMemcpyHostToDevice) );
4 438   CUDA_CALL( cudaMemcpy(d_neinum   , neinum    , neinumsize   ,
5      cudaMemcpyHostToDevice) );
6 439   CUDA_CALL( cudaMemcpy(d_neibmat   , neibmat    , neibmatsize   ,
7      cudaMemcpyHostToDevice) );
8 440
9 441   gettimeofday (&tdr1, NULL);
10 442   timeval_subtract (&restime, &tdr1, &tdr0);
11 443
12 444   std::cout << "\t" << getHRtime(restime) << std::endl;
13 445
14 446   // *****
15 447
16 448
17 449
18 450   // Initialize particles *****
19 451
20 452   std::cout << "Initializing particles" << std::endl;
21 453
22 454   //gettimeofday (&tdr0, NULL);
23 455
24 456   int nThreads      = 32;
25 457   int nBlocks       = ceil( (float) NP / nThreads);
26 458
27 459   initializeParticles <<< nBlocks, nThreads >>> (particles, nloc);
28 460   CUDA_CALL( cudaGetLastError() );
29 461   CUDA_CALL( cudaDeviceSynchronize() );
30 462
31 463   initializeParticles <<< nBlocks, nThreads >>> (particles_old,
32      nloc);
33 464   CUDA_CALL( cudaGetLastError() );
34 465   CUDA_CALL( cudaDeviceSynchronize() );
35 466
36 467   //gettimeofday (&tdr1, NULL);
37 468   //timeval_subtract (&restime, &tdr1, &tdr0);
38 469   //std::cout << "\t" << getHRtime(restime) << std::endl;
39 470
40 471   cudaMemGetInfo( &avail, &total );
41 472   used = total - avail;
42 473   std::cout << "\t[" << getHRmemsize(used) << "]" used of [" <<
43      getHRmemsize(total) << "]" <<std::endl;
44 474
45 475   // *****
46 476
47 477   // Starting filtering *****
48 478
49 479   for (int pass = 0; pass < 50; pass++) {
50 480
51 481       //std::cout << "pass = " << pass << std::endl;

```

```

1 482
2 483 // ** TEMP **
3 484 //clobberParams <<< nBlocks, nThreads >>> (particles, nloc);
4 485 // ** TEMP **
5 486
6 487 nThreads    = 32;
7 488 nBlocks     = ceil( (float) NP / nThreads);
8 489
9 490 resetStates <<< nBlocks, nThreads >>> (particles, nloc);
10 491 CUDA_CALL( cudaGetLastError() );
11 492 CUDA_CALL( cudaDeviceSynchronize() );
12 493
13 494 //std::cout << "Filtering over [1," << Tlim << "]"<< std::
14 495 endl;
15 496
16 496 //gettimeofday (&tdr0, NULL);
17 497
18 498 nThreads = 1;
19 499 nBlocks  = 10;
20 500
21 501 if (pass == 49) {
22 502     reduceStates <<< nBlocks, nThreads >>> (particles,
23 503         d_countmeans, 0, T, nloc);
24 504     CUDA_CALL( cudaGetLastError() );
25 505     CUDA_CALL( cudaDeviceSynchronize() );
26 506 }
27 507
28 507 //gettimeofday (&tdr1, NULL);
29 508 //timeval_subtract (&restime, &tdr1, &tdr0);
30 509 //std::cout << "\tReduction " << getHRtime(restime) << std::
31 510 endl;
32 510
33 511 int Tlim = T;
34 512
35 513 for (int t = 1; t < Tlim; t++) {
36 514
37 515     // Projection
38 516     *****
39 517
40 517 nThreads    = 32;
41 518 nBlocks     = ceil( (float) NP / nThreads);
42 519
43 520 //if (t == 1)
44 521 //    gettimeofday (&tdr0, NULL);
45 522
46 523 project <<< nBlocks, nThreads >>> (particles, d_neinum,
47 524     d_neibmat, nloc);
48 525 CUDA_CALL( cudaGetLastError() );
49 526 CUDA_CALL( cudaDeviceSynchronize() );
50 527
51 527 //if (t == 1 && pass == 0) {

```

```

1 528 // gettimeofday (&tdr1, NULL);
2 529 // timeval_subtract (&restime, &tdr1, &tdr0);
3 530 // std::cout << "\tProjection " << getHRtime(restime) <<
4 // std::endl;
5 531 //}
6 532
7 533 // Weighting
8 //*****
9 534
10 535 nThreads    = 32;
11 536 nBlocks     = ceil( (float) NP / nThreads);
12 537
13 538 weight <<< nBlocks, nThreads >>>(d_data, particles, w, t,
14 // T, nloc);
15 539 CUDA_CALL( cudaGetLastError() );
16 540 CUDA_CALL( cudaDeviceSynchronize() );
17 541
18 542 // Cumulative sum
19 //*****
20 543
21 544 nThreads    = 1;
22 545 nBlocks     = 1;
23 546
24 547 //if (t == 1)
25 548 // gettimeofday (&tdr0, NULL);
26 549
27 550 cumsumWeights <<< nBlocks, nThreads >>> (w);
28 551 CUDA_CALL( cudaGetLastError() );
29 552 CUDA_CALL( cudaDeviceSynchronize() );
30 553
31 554 //if (t == 1) {
32 555 // gettimeofday (&tdr1, NULL);
33 556 // timeval_subtract (&restime, &tdr1, &tdr0);
34 557 // std::cout << "\tCumulative sum " << getHRtime(
35 // restime) << std::endl;
36 558 //}
37 559
38 560 // Save particles for resampling from
39 //*****
40 561
41 562 nThreads    = 32;
42 563 nBlocks     = ceil( (float) NP / nThreads);
43 564
44 565 stashParticles <<< nBlocks, nThreads >>> (particles,
45 // particles_old, nloc);
46 566 CUDA_CALL( cudaGetLastError() );
47 567 CUDA_CALL( cudaDeviceSynchronize() );
48 568
49 569
50 570 // Resampling
51 //*****

```

```

1 571
2 572     nThreads    = 32;
3 573     nBlocks    = ceil( (float) NP/ nThreads);
4 574
5 575     //if (t == 1)
6 576     //  gettimeofday (&tdr0, NULL);
7 577
8 578     resample <<< nBlocks, nThreads >>> (particles,
9     particles_old, w, nloc);
10 579     CUDA_CALL( cudaGetLastError() );
11 580     CUDA_CALL( cudaDeviceSynchronize() );
12 581
13 582     //if (t == 1) {
14 583     //  gettimeofday (&tdr1, NULL);
15 584     //  timeval_subtract (&restime, &tdr1, &tdr0);
16 585     //  std::cout << "\tResampling " << getHRtime(restime) <<
17     std::endl;
18 586     //}
19 587
20 588     // Reduction
21     *****
22 589
23 590     //if (t == (Tlim-1)) {
24 591
25 592     if (pass == 49) {
26 593
27 594         //if (t == 1)
28 595         //  gettimeofday (&tdr0, NULL);
29 596
30 597         nThreads = 1;
31 598         nBlocks  = 10;
32 599
33 600         reduceStates <<< nBlocks, nThreads >>> (particles,
34         d_countmeans, t, T, nloc);
35 601         CUDA_CALL( cudaGetLastError() );
36 602         CUDA_CALL( cudaDeviceSynchronize() );
37 603
38 604         //if (t == 1) {
39 605         //  gettimeofday (&tdr1, NULL);
40 606         //  timeval_subtract (&restime, &tdr1, &tdr0);
41 607         //  std::cout << "Reduction          " << getHRtime(
42         restime) << std::endl;
43 608         //}
44 609
45 610     }
46 611
47 612     // Perturb particles
48     *****
49 613
50 614     nThreads    = 32;
51 615     nBlocks    = ceil( (float) NP/ nThreads);

```

```

1 616
2 617         perturbParticles <<< nBlocks, nThreads >>> (particles,
3         nloc, pass, 0.975);
4 618         CUDA_CALL( cudaGetLastError() );
5 619         CUDA_CALL( cudaDeviceSynchronize() );
6 620
7 621
8 622         } // end time
9 623
10 624     } // end pass
11 625
12 626     std::cout.precision(10);
13 627
14 628     countmeans = (float*) malloc (nloc*T*sizeof(float));
15 629     cudaMemcpy(countmeans, d_countmeans, nloc*T*sizeof(float),
16         cudaMemcpyDeviceToHost);
17 630
18 631     // stop master timer and print
19 632
20 633     gettimeofday (&tdrMaster, NULL);
21 634     timeval_subtract(&restime, &tdrMaster, &tdr0);
22 635     std::cout << "Time: " << getHRtime(restime) << std::endl;
23 636     std::cout << "Rawtime: " << restime << std::endl;
24 637
25 638     // Write results out
26 639
27 640     std::string filename = arg4;
28 641
29 642     std::cout << "Writing results to file '" << filename << "' ..."
30         << std::endl;
31 643
32 644     std::ofstream outfile;
33 645     outfile.open(filename.c_str());
34 646
35 647     for(int loc = 0; loc < nloc; loc++) {
36 648         for (int t = 0; t < T; t++) {
37 649             outfile << countmeans[loc*T + t] << " ";
38 650         }
39 651         outfile << std::endl;
40 652     }
41 653
42 654     outfile.close();
43 655
44 656     //gettimeofday (&tdr1, NULL);
45 657     //timeval_subtract (&restime, &tdr1, &tdrMaster);
46 658     //std::cout << "Total PF time (excluding setup) " << getHRtime(
47         restime) << std::endl;
48 659
49 660     cudaFree(d_data);
50 661     cudaFree(particles);
51 662     cudaFree(particles_old);

```



```

1 663     cudaFree(w);
2 664     cudaFree(d_neinum);
3 665     cudaFree(d_neibmat);
4 666     cudaFree(d_countmeans);
5 667
6 668     exit (EXIT_SUCCESS);
7 669
8 670 }
9 671
10 672
11 673 /* Use the Explicit Euler integration scheme to integrate SIR model
12     forward in time
13 674     float h      - time step size
14 675     float t0     - start time
15 676     float tn     - stop time
16 677     float * y    - current system state; a three-component vector
17                     representing [S I R], susceptible-infected-recovered
18 678     */
19 679 __device__ void exp_euler_SSIR(float h, float t0, float tn, Particle
20     * particle, int * neinum, int * neibmat, int nloc) {
21 680
22 681     int num_steps = floor( (tn-t0) / h );
23 682
24 683     float * S = particle->S;
25 684     float * I = particle->I;
26 685     float * R = particle->R;
27 686     float * B = particle->B;
28 687
29 688     // create last state vectors
30 689     float * S_last = (float*) malloc (nloc*sizeof(float));
31 690     float * I_last = (float*) malloc (nloc*sizeof(float));
32 691     float * R_last = (float*) malloc (nloc*sizeof(float));
33 692     float * B_last = (float*) malloc (nloc*sizeof(float));
34 693
35 694     float R0      = particle->R0;
36 695     float r       = particle->r;
37 696     float B0      = R0 * r / N;
38 697     float eta     = particle->eta;
39 698     float berr    = particle->berr;
40 699     float phi     = particle->phi;
41 700
42 701     for(int t = 0; t < num_steps; t++) {
43 702
44 703         for (int loc = 0; loc < nloc; loc++) {
45 704             S_last[loc] = S[loc];
46 705             I_last[loc] = I[loc];
47 706             R_last[loc] = R[loc];
48 707             B_last[loc] = B[loc];
49 708         }
50 709
51 710         for (int loc = 0; loc < nloc; loc++) {

```

```

1 711
2 712     B[loc] = exp( log(B_last[loc]) + eta*(log(B0) - log(
3       B_last[loc])) + berr*curand_normal(&(particle->
4       randState)) );
5 713
6 714     int n = neinum[loc];
7 715     float sphi = 1.0 - phi*( (float) n/(n+1.0) );
8 716     float ophi = phi/(n+1.0);
9 717
10 718     float nBIsun = 0.0;
11 719     for (int j = 0; j < n; j++)
12 720         nBIsun += B_last[neibmat[nloc*loc + j]-1] * I_last[
13           neibmat[nloc*loc + j]-1];
14 721
15 722     float BSI = S_last[loc]*( sphi*B_last[loc]*I_last[loc] +
16       ophi*nBIsun );
17 723     float rI = r*I_last[loc];
18 724
19 725     // get derivatives
20 726     float dS = - BSI;
21 727     float dI = BSI - rI;
22 728     float dR = rI;
23 729
24 730     // step forward by h
25 731     S[loc] += h*dS;
26 732     I[loc] += h*dI;
27 733     R[loc] += h*dR;
28 734
29 735     }
30 736
31 737 }
32 738
33 739 free(S_last);
34 740 free(I_last);
35 741 free(R_last);
36 742 free(B_last);
37 743
38 744 }
39 745
40 746 /* Convenience function for particle resampling process
41 747 */
42 748 __device__ void copyParticle(Particle * dst, Particle * src, int nloc
43 ) {
44 749
45 750     dst->R0 = src->R0;
46 751     dst->r = src->r;
47 752     dst->sigma = src->sigma;
48 753     dst->eta = src->eta;
49 754     dst->berr = src->berr;
50 755     dst->phi = src->phi;
51 756

```

```

1 757     for (int n = 0; n < nloc; n++) {
2 758         dst->S[n]      = src->S[n];
3 759         dst->I[n]      = src->I[n];
4 760         dst->R[n]      = src->R[n];
5 761         dst->B[n]      = src->B[n];
6 762         dst->Iinit[n]  = src->Iinit[n];
7 763     }
8 764
9 765 }
10 766
11 767 /* Convert memory size in bytes to human-readable format
12 768 */
13 769 std::string getHRmemsize (size_t memsize) {
14 770
15 771     std::stringstream ss;
16 772     std::string valstring;
17 773
18 774     int kb = 1024;
19 775     int mb = kb*1024;
20 776     int gb = mb*1024;
21 777
22 778     if (memsize <= kb)
23 779         ss << memsize << " B";
24 780     else if (memsize > kb && memsize <= mb)
25 781         ss << (float) memsize/ kb << " KB";
26 782     else if (memsize > mb && memsize <= gb)
27 783         ss << (float) memsize/ mb << " MB";
28 784     else
29 785         ss << (float) memsize/ gb << " GB";
30 786
31 787     valstring = ss.str();
32 788
33 789     return valstring;
34 790
35 791 }
36 792
37 793
38 794 /* Convert time in seconds to human readable format
39 795 */
40 796 std::string getHRtime (float runtime) {
41 797
42 798     std::stringstream ss;
43 799     std::string valstring;
44 800
45 801     int mt = 60;
46 802     int ht = mt*60;
47 803     int dt = ht*24;
48 804
49 805     if (runtime <= mt)
50 806         ss << runtime << " s";
51 807     else if (runtime > mt && runtime <= ht)

```

```

1 808         ss << runtime/mt << " m";
2 809     else if (runtime > ht && runtime <= dt)
3 810         ss << runtime/dt << " h";
4 811     else
5 812         ss << runtime/ht << " d";
6 813
7 814     valstring = ss.str();
8 815
9 816     return valstring;
10 817
11 818 }

```

13 Figure [F.1] shows the running times for parameter fitting as compared to IF2 and
14 HMC MC.




Figure F.1: Running times for fitting the spatial SIR model to data.

15 The means from the data in Figure [F.1] are about 61.5 seconds for cuIF2, 574 seconds
16 for IF2, and 38,800 seconds for HMC MC. For cuIF2 This is a speedup of over 9.33x
17 against IF2 and over 617x against HMC MC.