

Particle Filters

Dexter Barrows
November 11, 2015

1 Intro

Particle filters are similar to MCMC-based methods in that they attempt to draw samples from an approximation of the posterior distribution of model parameters θ given observed data D . Instead of constructing a Markov chain and approximating its stationary distribution, a cohort of “particles” are used to move through the data in an on-line (sequential) fashion with the cohort being culled of poorly-performing particles at each iteration via importance sampling. If the culled particles are not replenished, this will be a Sequential Importance Sampling (SIS) particle filter. If the culled particles are replenished from surviving particles, in a sense setting up a process not dissimilar from Darwinian selection, then this will be a Sequential Importance Resampling (SIR) particle filter.

2 Formulation

Particle filters, also called Sequential Monte-Carlo (SMC) or bootstrap filters, feature similar core functionality as the venerable Kalman Filter. As the algorithm moves through the data (sequence of observations), a prediction-update cycle is used to simulate the evolution of the model M with different particular parameter selections, track how closely these predictions approximate the new observed value, and update the current cohort appropriately.

Two separate functions are used to simulate the evolution and observation processes. The “true” state evolution is specified by

$$X_{t+1} \sim f_1(X_t, \theta), \tag{1}$$

And the observation process by

$$Y_t \sim f_2(X_t, \theta). \tag{2}$$

Note that components of θ can contribute to both functions, but a typical formulation is to have some components contribute to $f_1(\cdot, \theta)$ and others to $f_2(\cdot, \theta)$.

The prediction part of the cycle utilises $f_1(\cdot, \theta)$ to update each particle's current state estimate to the next time step, while $f_2(\cdot, \theta)$ is used to evaluate a weighting w for each particle which will be used to determine how closely that particle is estimating the true underlying state of the system. Note that $f_2(\cdot, \theta)$ could be thought of as a probability of observing a piece of data y_t given the particle's current state estimate and parameter set, $P(y_t|X_t, \theta)$. Then, the new cohort of particles is drawn from the old cohort proportional to the weights. This process is repeated until the set of observations D is exhausted.

3 Algorithm

Now we can formalize the particle filter.

We will denote each particle $p^{(j)}$ as the j^{th} particle consisting of a state estimate at time t , $X_t^{(j)}$, a parameter set $\theta^{(j)}$, and a weight $w^{(j)}$. Note that the state estimates will evolve with the system as the cohort traverses the data.

The algorithm for a Sequential Importance Resampling particle is shown in Algorithm 1.

Algorithm 1: SIR particle filter

```

/* Select a starting point */
Input : Observations  $D = y_1, y_2, \dots, y_T$ , initial particle distribution  $P_0$  of size  $J$ 

/* Setup */
1 Initialize particle cohort by sampling  $(p^{(1)}, p^{(2)}, \dots, p^{(J)})$  from  $P_0$ 
2 for  $t = 1 : T$  do
    /* Evolve */
    3 for  $j = 1 : J$  do
    4    $X_t^{(j)} \leftarrow f_1(X_{t-1}^{(j)}, \theta^{(j)})$ 

    /* Weight */
    5 for  $j = 1 : J$  do
    6    $w^{(j)} \leftarrow P(y_t | X_t^{(j)}, \theta^{(j)}) = f_2(X_t^{(j)}, \theta^{(j)})$ 

    /* Normalize */
    7 for  $j = 1 : J$  do
    8    $w^{(j)} \leftarrow w^{(j)} / \sum_1^J w^{(j)}$ 

    /* Resample */
    9  $p^{(1:J)} \leftarrow \text{sample}(p^{(1:J)}, \text{prob} = w, \text{replace} = \text{true})$ 

/* Samples from approximated posterior distribution */
Output: Cohort of posterior samples  $(\theta^{(1)}, \theta^{(2)}, \dots, \theta^{(J)})$ 

```

4 Particle Collapse

Not uncommonly, a situation may arise in which a single particle gets assigned a normalized weight very close to 1 and all the other particles have weights very close to 0. When this occurs, the next generation of the cohort will overwhelmingly consist of descendants of the heavily-weighted particle, termed particle collapse or degeneracy.

Since the basic SIR particle filter does not perturb either the particle system states or system parameter values, the cohort will quickly consist solely of identical particles, effectively halting further exploration of the parameter space as new data is introduced.

A similar situation occurs when a small number of particles (but not necessarily a single particle) split almost all of the normalized weight between them, then jointly dominate the resampling process for the remainder of the iterations. This again halts the exploration of the parameter space with new data.

In either case, the hallmark feature used to detect collapse is the same - at some point the cohort will consist of particles with very similar or identical parameter sets which will consequently result in their assigned weights being extremely close.

Mathematically, we are interested in the number of effective particles, N_{eff} , which represents the number of particles that are acceptably dissimilar. This is estimated by evaluating

$$N_{eff} = \frac{1}{\sum_1^J (w^{(j)})^2}. \quad (3)$$

This can be used to diagnose not only when collapse has occurred, but can also indicate when it is near.

Diagram!!!!!!

5 Iterated Filtering and Data Cloning

A particle filter hinges on the idea that as it progresses through the data set D , its estimate of the the posterior carried in the cohort of particles approaches maximum likelihood. However, this convergence may not converge fast enough that the estimate it produces is of quality before the data runs out. One way around this problem is to “clone” the data and make multiple passes through it as if it were a continuation of the original time series. Note that the system state contained in each particle will have to be reset with each pass.

Rigorous proofs have been developed (references to Ionides et. al. work) that show that by treating the parameters as stochastic processes instead of fixed values, the multiple passes through the data will indeed force convergence of the process mean toward maximum likelihood, and the process variance toward 0.

6 IF2

The successor to Iterated Filtering 1 (reference), Iterated Filtering 2 is simpler, faster, and demonstrated better convergence toward maximum likelihood (reference). The core concept involves a two-pronged approach. First, Data cloning is used to allow more time for the parameter stochastic process means to converge to maximum likelihood, and frequent cooled perturbation of the particle parameters allow better exploration of the parameter space while still allowing convergence to good point estimates.

It is worth noting that IF2 is not designed to estimate the full posterior distribution, but in practice can be used to do so within reason. Further, IF2 thwarts the problem of particle collapse by keeping at least some perturbation in the system at all times. It is important to note that while true particle collapse will not occur, there is still risk of a pseudo-collapse in which all particles will be extremely close to one another so as to be virtually indistinguishable. However this will only occur with the use of overly-aggressive cooling strategies or by specifying an excessive number of passes through the data.

An important new quantity is the particle perturbation density denoted $h(\theta|\sigma)$. Typically this is multi-normal with σ being a vector of variances proportional to the expected values of θ . In practice the proportionality can be derived from current means or specified ahead of time. Further, these intensities must decrease over time. This can be done via exponential cooling, a decreasing step function, a combination of these, or some similar way.

The algorithm for IF2 can be seen in Algorithm 2.

Algorithm 2: IF2

```
/* Select a starting point */
Input : Observations  $D = y_1, y_2, \dots, y_T$ , initial particle distribution  $P_0$  of size  $J$ ,
        decreasing sequence of perturbation intensity vectors  $\sigma_1, \sigma_2, \dots, \sigma_M$ 

/* Setup */
1 Initialize particle cohort by sampling  $(p^{(1)}, p^{(2)}, \dots, p^{(J)})$  from  $P_0$ 

/* Particle seeding distribution */
2  $\Theta \leftarrow P_0$ 
3 for  $m = 1 : M$  do
    /* Pass perturbation */
    4 for  $j = 1 : J$  do
    5      $p^{(j)} \sim h(\Theta^{(j)}, \sigma_m)$ 
    6 for  $t = 1 : T$  do
    7     for  $j = 1 : J$  do
    8         /* Iteration perturbation */
    9          $p^{(j)} \sim h(p^{(j)}, \sigma_m)$ 
    10        /* Evolve */
    11         $X_t^{(j)} \leftarrow f_1(X_{t-1}^{(j)}, \theta^{(j)})$ 
    12        /* Weight */
    13         $w^{(j)} \leftarrow P(y_t | X_t^{(j)}, \theta^{(j)}) = f_2(X_t^{(j)}, \theta^{(j)})$ 
    14        /* Normalize */
    15        for  $j = 1 : J$  do
    16             $w^{(j)} \leftarrow w^{(j)} / \sum_1^J w^{(j)}$ 
    17        /* Resample */
    18         $p^{(1:J)} \leftarrow \text{sample}(p^{(1:J)}, \text{prob} = w, \text{replace} = \text{true})$ 
    19        /* Collect particles for next pass */
    20        for  $j = 1 : J$  do
    21             $\Theta^{(j)} \leftarrow p^{(j)}$ 

/* Samples from approximated posterior distribution */
Output: Cohort of posterior samples  $(\theta^{(1)}, \theta^{(2)}, \dots, \theta^{(J)})$ 
```

7 Fitting an SIR Model to Synthetic Epidemic Data with IF2

Here we will examine a test case in which IF2 will be used to fit a Susceptible-Infected-Removed (SIR) epidemic model to mock infectious count data.

The synthetic data was produced by taking the solution to a basic SIR ODE model, sampling it at regular intervals, and perturbing those values by adding in observation noise. The SIR model used was

$$\begin{aligned}\frac{dS}{dt} &= -\beta IS \\ \frac{dI}{dt} &= \beta IS - rI \\ \frac{dR}{dt} &= rI\end{aligned}\tag{4}$$

where S is the number of individuals susceptible to infection, I is the number of infectious individuals, R is the number of recovered individuals, $\beta = R_0 r / N$ is the force of infection, R_0 is the number of secondary cases per infected individual, r is the recovery rate, and N is the population size.

The solution to this system was obtained using the `ode()` function from the `deSolve` package. The required derivative array function in the format required by `ode()` was specified as

```

1      SIR <- function(Time, State, Pars) {
2
3          with(as.list(c(State, Pars)), {
4
5              B    <- R0*r/N      # calculate Beta
6              BSI  <- B*S*I      # save product
7              rI   <- r*I        # save product
8
9              dS   = -BSI        # change in Susceptible people
10             dI   = BSI - rI    # change in Infected people
11             dR   = rI          # change in Removed (recovered people)
12
13             return(list(c(dS, dI, dR)))
14
15         })
16
17     }
```

The true parameter values were set to $R_0 = 3.0$, $r = 0.1$, $N = 500$ by

```

1      pars <- c(R0 = 3.0, # new infected people per infected person
2              r   = 0.1, # recovery rate
3              N   = 500) # population size
```

The initial conditions were set to 5 infectious individuals, 495 people susceptible to infection, and no one had yet recovered from infection and been removed. These were set using

```
1 true_init_cond <- c(S = N - i_infec,
2                     I = i_infec,
3                     R = 0)
```

The `ode()` function is called as

```
1 odeout <- ode(y = true_init_cond, times = 0:(T-1), func = SIR, parms =
  true_pars)
```

where `odeout` is a $T \times 4$ matrix where the rows correspond to solutions at the given times (the first row is the initial condition), and the columns correspond to the solution times and S-I-R counts at those times.

The observation error was taken to be $\varepsilon_{obs} \sim \mathcal{N}(0, \sigma)$, where individual values were drawn for each synthetic data point.

These “true” values were perturbed to mimic observation error by

```
1 set.seed(1001) # set RNG seed for reproducibility
2 sigma <- 10    # observation error standard deviation
3 infec_counts_raw <- odeout[,3] + rnorm(101, 0, sigma)
4 infec_counts <- ifelse(infec_counts_raw < 0, 0, infec_counts)
```

where the last two lines simply set negative observations (impossible) to 0.

Plotting the data using the `ggplot2` package by

```
1 plotdata <- data.frame(times=1:T, true=trueTraj, data=infec_counts)
2
3 g <- ggplot(plotdata, aes(times)) +
4   geom_line(aes(y = true, colour = "True")) +
5   geom_point(aes(y = data, color = "Data")) +
6   labs(x = "Time", y = "Infection count", color = "") +
7   scale_color_brewer(palette="Paired") +
8   theme(panel.background = element_rect(fill = "#F0F0F0"))
```

we obtain Figure 1.




Figure 1: True SIR ODE solution infected counts, and with added observation noise

The IF2 algorithm was implemented in C++ for speed, and integrated into the R workflow using the `Rcpp` package. The C++ code is compiled using

```
1 sourceCpp(paste(getwd(), "if2.cpp", sep="/"))
```

Then run and packed into a data frame using

```
1 paramdata <- data.frame(if2(infec_counts[1:Tlim], Tlim, N))
2 colnames(paramdata) <- c("R0", "r", "sigma", "Sinit", "Iinit", "Rinit")
```

The final kernel estimates for four of the key parameters are shown in Figure 2.

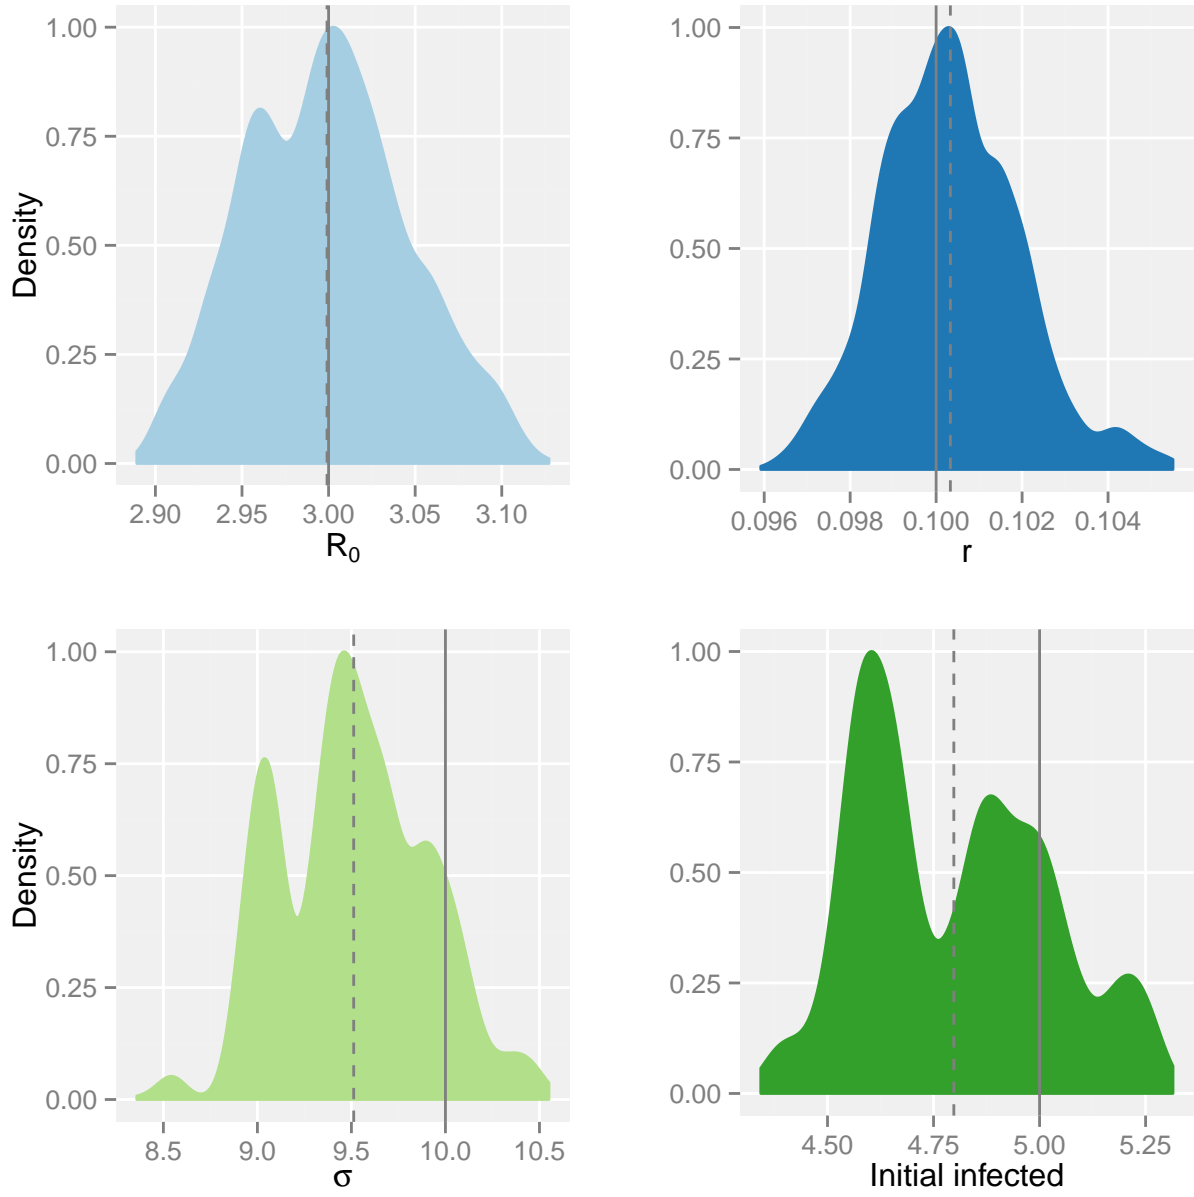


Figure 2: Kernel estimates for four essential system parameters. True values are indicated by solid vertical lines, sample means by dashed lines.

Bootstrapping from the final particle sample gives us awesome stuff too as we can see in Figure 3.




Figure 3: IF2 estimate of the system state at each time point using bootstrapping. Ribbon shows 2.5% and 97.5% quantiles.

Appendices

A Full R code

This code will run all the indicated analysis and produce all plots.

```
1 library(deSolve)
2 library(ggplot2)
3 library(reshape2)
4 library(gridExtra)
5 library(Rcpp)
6 library(RColorBrewer)
7
8 SIR ← function(Time, State, Pars) {
9
10     with(as.list(c(State, Pars)), {
11
12         B ← R0*r/N
13         BSI ← B*S*I
14         rI ← r*I
15
16         dS = -BSI
17         dI = BSI - rI
18         dR = rI
19
20         return(list(c(dS, dI, dR)))
21     })
22 }
23
24 }
25
26 T ← 100
27 N ← 500
28 sigma ← 10
29 i_infec ← 5
30 Tlim ← T
31
32 ## Generate true trajecory and synthetic data
33 ##
34
35 true_init_cond ← c(S = N - i_infec,
36                   I = i_infec,
37                   R = 0)
38
39 true_pars ← c(R0 = 3.0,
40              r = 0.1,
41              N = 500.0)
42
43 odeout ← ode(true_init_cond, 0:(T-1), SIR, true_pars)
44 trueTraj ← odeout[,3]
```

```

45
46 set.seed(1000)
47
48 infec_counts_raw ← odeout[,3] + rnorm(T, 0, sigma)
49 infec_counts     ← ifelse(infec_counts_raw < 0, 0, infec_counts_raw)
50
51 plotdata ← data.frame(times=1:T,true=trueTraj,data=infec_counts)
52
53 g ← ggplot(plotdata, aes(times)) +
54     geom_line(aes(y = true, colour = "True")) +
55     geom_point(aes(y = data, color = "Data")) +
56     labs(x = "Time", y = "Infection count", color = "") +
57     scale_color_brewer(palette="Paired") +
58     theme(panel.background = element_rect(fill = "#F0F0F0"))
59
60 print(g)
61 ggsave(g, filename="dataplot.pdf", height=4, width=6.5)
62
63 ## Rcpp stuff
64 ##
65
66 sourceCpp(paste(getwd(),"if2.cpp",sep="/"))
67
68 paramdata ← data.frame(if2(infec_counts[1:Tlim], Tlim, N))
69 colnames(paramdata) ← c("R0", "r", "sigma", "Sinit", "Iinit", "Rinit")
70
71 ## sample from parameter distributions
72 ##
73
74 nTraj ← 100
75 datlen ← dim(paramdata)[1]
76 inds ← sample.int(datlen,nTraj,replace = TRUE)
77 params ← paramdata[inds,]
78
79 bootstrapdata ← matrix(NA, nrow = nTraj, ncol = T)
80
81 for (i in 1:nTraj) {
82
83     init_cond ← c(S = params$Sinit[i],
84                  I = params$Iinit[i],
85                  R = params$Rinit[i])
86     pars ← c(R0 = params$R0[i],
87             r = params$r[i],
88             N = 500.0)
89
90     odeout ← ode(init_cond, 0:(T-1), SIR, pars)
91
92     bootstrapdata[i,] ← odeout[,3]
93
94 }
95
96
97 meanTraj ← colMeans(bootstrapdata)
98 quantTraj ← apply(bootstrapdata, 2, quantile, probs = c(0.025,0.975))

```

```

99
100 datapart <- c(infec_counts[1:Tlim], rep(NA, T-Tlim))
101
102 plotdata <- data.frame(times=1:T, true=trueTraj, est=meanTraj, quants=t(
  quantTraj), datapart=datapart)
103
104 g <- ggplot(plotdata, aes(times)) +
105   geom_ribbon(aes(ymin = quants.2.5., ymax=quants.97.5.), alpha=0.1)
106   +
107   geom_line(aes(y = true, colour = "True")) +
108   geom_line(aes(y = est, colour = "IF2")) +
109   geom_point(aes(y = datapart, color = "Data")) +
110   labs(x = "Time", y = "Infection count", color = "") +
111   scale_color_brewer(palette="Paired") +
112   theme(panel.background = element_rect(fill = "#F0F0F0"))
113 print(g)
114 ggsave(g, filename="if2plot.pdf", height=4, width=6.5)
115
116 f <- function(pal) brewer.pal(brewer.pal.info[pal, "maxcolors"], pal)
117 kcolours <- f("Paired")
118
119
120 trueval.R0 <- 3.0
121 trueval.r <- 0.1
122 trueval.sigma <- 10.0
123 trueval.Iinit <- 5
124
125 meanval.R0 <- mean(paramdata$R0)
126 meanval.r <- mean(paramdata$r)
127 meanval.sigma <- mean(paramdata$sigma)
128 meanval.Iinit <- mean(paramdata$Iinit)
129
130 linecolour <- "grey50"
131 lineweight <- 0.5
132
133 kerdataR0 <- data.frame(R0points = paramdata$R0)
134 R0kernel <- ggplot(kerdataR0, aes(x = R0points, y = ..scaled..)) +
135   geom_density(color = kcolours[1], fill = kcolours[1]) +
136   theme(panel.background = element_rect(fill = "#F0F0F0")) +
137   scale_color_brewer(palette="Paired") +
138   labs(x = expression(R[0]), y = "Density", color = "") +
139   geom_vline(aes(xintercept=trueval.R0), linetype="solid",
140     size=lineweight, color=linecolour) +
141   geom_vline(aes(xintercept=meanval.R0), linetype="dashed",
142     size=lineweight, color=linecolour)
143
144 kerdatar <- data.frame(rpoints = paramdata$r)
145 rkern <- ggplot(kerdatar, aes(x = rpoints, y = ..scaled..)) +
146   geom_density(color = kcolours[2], fill = kcolours[2]) +
147   theme(panel.background = element_rect(fill = "#F0F0F0")) +
148   scale_color_brewer(palette="Paired") +
149   labs(x = "r", y = "", color = "") +

```

```

148         geom_vline(aes(xintercept=trueval.r), linetype="solid",
149                     size=lineweight, color=linecolour) +
150         geom_vline(aes(xintercept=meanval.r), linetype="dashed",
151                     size=lineweight, color=linecolour)
152
153     kerdatasigma <- data.frame(sigmappoints = paramdata$sigma)
154     sigmakernel <- ggplot(kerdatasigma, aes(x = sigmappoints, y = ..scaled..)) +
155         geom_density(color = kcolours[3], fill = kcolours[3]) +
156         theme(panel.background = element_rect(fill = "#F0F0F0")) +
157         scale_color_brewer(palette="Paired") +
158         labs(x = expression(sigma), y = "Density", color = "") +
159         geom_vline(aes(xintercept=trueval.sigma), linetype="solid",
160                     size=lineweight, color=linecolour) +
161         geom_vline(aes(xintercept=meanval.sigma), linetype="dashed",
162                     size=lineweight, color=linecolour)
163
164     kerdatainfec <- data.frame(infecpoints = paramdata$Iinit)
165     infeckernel <- ggplot(kerdatainfec, aes(x = infecpoints, y = ..scaled..)) +
166         geom_density(color = kcolours[4], fill = kcolours[4]) +
167         theme(panel.background = element_rect(fill = "#F0F0F0")) +
168         scale_color_brewer(palette="Paired") +
169         labs(x = "Initial infected", y = "", color = "") +
170         geom_vline(aes(xintercept=trueval.Iinit), linetype="solid",
171                     size=lineweight, color=linecolour) +
172         geom_vline(aes(xintercept=meanval.Iinit), linetype="dashed",
173                     size=lineweight, color=linecolour)
174
175     # show grid
176     grid.arrange(R0kernel, rkernel, sigmakernel, infeckernel, ncol = 2, nrow =
177                   2)
178
179     pdf("if2kernels.pdf", height = 6.5, width = 6.5)
180     grid.arrange(R0kernel, rkernel, sigmakernel, infeckernel, ncol = 2, nrow =
181                   2)
182     dev.off()

```

B Full C++ code

Stan model code to be used with the preceding R code.

```

1  /* Author: Dexter Barrows
2     Github: dbarrows.github.io
3
4     */
5
6  /* Runs a particle filter on synthetic noisy data and attempts to
7     reconstruct underlying true state at each time step. Note that
8     this program uses gnuplot to plot the data, so an x11
9     environment must be present.
10

```

```

11     Also, the accompanying "pf.plg" file contains the instructions
12     gnuplot will use. It must be present in the same directory as
13     the executable generated by compiling this file.
14
15     */
16
17 #include <stdio.h>
18 #include <math.h>
19 #include <sys/time.h>
20 #include <time.h>
21 #include <stdlib.h>
22 #include <string>
23
24 #include "rand.h"
25 #include "timer.h"
26
27 #define T      100          // time to simulate over
28 #define R0true  3.0          // infectiousness
29 #define rtrue   0.1          // recovery rate
30 #define N      500.0        // population size
31 #define merr    10.0         // expected measurement error
32
33
34 struct Particle {
35     float R0;
36     float r;
37     float sigma;
38     float S;
39     float I;
40     float R;
41     float Sinit;
42     float Iinit;
43     float Rinit;
44 };
45
46 struct ParticleInfo {
47     float R0mean;          float R0sd;
48     float rmean;           float rsd;
49     float sigmamean;       float sigmasd;
50     float Sinitmean;       float Sinitstd;
51     float Iinitmean;       float Iinitstd;
52     float Rinitmean;       float Rinitstd;
53 };
54
55
56 int timeval_subtract (double *result, struct timeval *x, struct timeval *y)
57 ;
58 int check_float(float x, float y);
59 void exp_euler_SIR(float h, float t0, float tn, Particle * particle);
60 void copyParticle(Particle * dst, Particle * src);
61 void perturbParticles(Particle * particles, int NP, int passnum, float
    coolrate);
62 bool isCollapsed(Particle * particles, int NP);

```

```

62 void particleDiagnostics(ParticleInfo * partInfo, Particle * particles, int
    NP);
63
64
65 int main(int argc, char *argv[]) {
66
67     float    i_infec      = 5;
68     int       Tlim        = T;
69     int       NP          = 3000;
70     int       nPasses     = 40;
71     float     coolrate    = 7;
72
73     srand(time(NULL)); // Seed PRNG with
        system time
74
75     Particle particle_true;
76     particle_true.R0      = R0true;
77     particle_true.r       = rtrue;
78     particle_true.sigma   = merr;
79     particle_true.S       = N - i_infec;
80     particle_true.I       = i_infec;
81     particle_true.R       = 0;
82
83
84     printf("System parameters\n");
85     printf("-----\n");
86     printf("R0:      %f\n", R0true);
87     printf("r:       %f\n", rtrue);
88     printf("merr:    %f\n", merr);
89
90     float y_true[T]; // true number of infected peeps
91     float y_noise[T]; // true number of infected peeps with observation
        noise
92     float y_est[T]; // particle mean state estimation
93
94     float y_par_noise; // particle estimates with noise
95     float w[NP]; // particle weights
96
97     Particle particles[NP]; // particle estimates for current step
98     Particle particles_old[NP]; // intermediate particle states for
        resampling
99
100    // generate our true trajectory and noisy observation data
101    y_true[0] = particle_true.I;
102    y_noise[0] = y_true[0] + merr*randn();
103    if (y_noise[0] < 0)
104        y_noise[0] = 0;
105    for (int i = 1; i < T; i++) {
106        exp_euler_SIR( 1.0/100, 0.0, 1.0, &particle_true);
107        y_true[i] = particle_true.I;
108        y_noise[i] = y_true[i] + merr*randn();
109        if (y_noise[i] < 0)
110            y_noise[i] = 0;
111    }

```



```

112
113     double restime;
114     struct timeval  tdr0, tdr1;
115
116     gettimeofday (&tdr0, NULL);
117
118
119
120     printf("Initializing particle states\n");
121
122     // initialize particle parameter states (seeding)
123     for (int n = 0; n < NP; n++) {
124
125         float R0can, rcan, sigmacan, Iinitcan;
126
127         do {
128             R0can = R0true + R0true*randn();
129         } while (R0can < 0);
130         particles[n].R0 = R0can;
131
132         do {
133             rcan = rtrue + rtrue*randn();
134         } while (rcan < 0);
135         particles[n].r = rcan;
136
137         do {
138             sigmacan = merr + merr*randn();
139         } while (sigmacan < 0);
140         particles[n].sigma = sigmacan;
141
142         do {
143             Iinitcan = i_infec + i_infec*randn();
144         } while (Iinitcan < 0 || N < Iinitcan);
145         particles[n].Sinit = N - Iinitcan;
146         particles[n].Iinit = Iinitcan;
147         particles[n].Rinit = 0.0;
148     }
149
150
151     // START PASSES THROUGH DATA
152
153
154
155     printf("Starting filter\n");
156     printf("-----\n");
157     printf("Pass");
158
159
160     for (int pass = 0; pass < nPasses; pass++) {
161
162         printf("...%d", pass);
163
164         perturbParticles(particles, NP, pass, coolrate);
165

```

```

166 // initialize particle system states
167 for (int n = 0; n < NP; n++) {
168
169     particles[n].S = particles[n].Sinit;
170     particles[n].I = particles[n].Iinit;
171     particles[n].R = particles[n].Rinit;
172
173 }
174
175 // between-pass perturbations
176
177 for (int t = 1; t < Tlim; t++) {
178
179     // between-iteration perturbations
180     perturbParticles(particles, NP, pass, coolrate);
181
182     // generate individual predictions and weight
183     for (int n = 0; n < NP; n++) {
184
185         exp_euler_SIR(1.0/10.0, 0.0, 1.0, &particles[n]);
186
187         float merr_par = particles[n].sigma;
188         float y_diff = y_noise[t] - particles[n].I;
189
190         w[n] = 1.0/(merr_par*sqrt(2.0*M_PI)) * exp( - y_diff*y_diff
191             / (2.0*merr_par*merr_par) );
192
193     }
194
195     // cumulative sum
196     for (int n = 1; n < NP; n++) {
197         w[n] += w[n-1];
198     }
199
200     // save particle states to resample from
201     for (int n = 0; n < NP; n++){
202         copyParticle(&particles_old[n], &particles[n]);
203     }
204
205     // resampling
206     for (int n = 0; n < NP; n++) {
207
208         float w_r = randu() * w[NP-1];
209         int i = 0;
210         while (w_r > w[i]) {
211             i++;
212         }
213
214         // i is now the index to copy state from
215         copyParticle(&particles[n], &particles_old[i]);
216     }
217
218 }

```

```

219
220     }
221
222     ParticleInfo pInfo;
223     particleDiagnostics(&pInfo, particles, NP);
224
225     printf("\n");
226
227     gettimeofday (&tdr1, NULL);
228     timeval_subtract (&restime, &tdr1, &tdr0);
229     printf ("Single threaded runtime %e\n", restime);
230
231
232     // Save parameter distribution results for post-processing
233
234     std::string paramfile("pfddata.dat");
235     FILE * pfout = fopen(paramfile.c_str(), "w");
236
237     //printf("Writing parameter results to file '%s'...\n", paramfile.c_str
238         ());
239
240     for (int n = 0; n < NP; n++) {
241
242         fprintf(pfout, "%f ", particles[n].R0);
243         fprintf(pfout, "%f ", particles[n].r);
244         fprintf(pfout, "%f ", particles[n].sigma);
245         fprintf(pfout, "%f ", particles[n].Sinit);
246         fprintf(pfout, "%f ", particles[n].Iinit);
247         fprintf(pfout, "%f\n", particles[n].Rinit);
248     }
249
250     fclose(pfout);
251
252     // Save results for plotting
253
254     std::string datafile("plotdata.dat");
255
256     printf("Writing plotting results to file '%s'...\n", datafile.c_str());
257
258     FILE * paramout = fopen(datafile.c_str(), "w");
259
260     for (int t = 0; t < T; t++) {
261
262         fprintf(paramout, "%d ", t);
263         fprintf(paramout, "%f ", y_true[t]);
264
265         if (t < Tlim)
266             fprintf(paramout, "%f ", y_noise[t]);
267         else
268             fprintf(paramout, "%d ", -1);
269
270         fprintf(paramout, "%f\n", y_est[t]);
271

```

```

272     }
273
274     fclose(paramout);
275
276     printf("Plotting using gnuplot...\n");
277     printf("Press ENTER close plot and continue\n");
278
279     std::string syscall("gnuplot -e \"filename='\"");
280     syscall += datafile;
281     syscall += "'\" pf.plg";
282
283     //system( syscall.c_str() );
284
285 }
286
287
288 /* Use the Explicit Euler integration scheme to integrate SIR model
    forward in time
289     float h      - time step size
290     float t0     - start time
291     float tn     - stop time
292     float * y    - current system state; a three-component vector
                     representing [S I R], susceptible-infected-recovered
293
294     */
295 void exp_euler_SIR(float h, float t0, float tn, Particle * particle) {
296
297     float t = t0;
298
299     int num_steps = floor( (tn-t0) / h );
300
301     float S = particle->S;
302     float I = particle->I;
303     float R = particle->R;
304
305     float R0    = particle->R0;
306     float r     = particle->r;
307     float B     = R0 * r / N;
308
309     for(int i = 0; i < num_steps; i++) {
310         // get derivatives
311         float dS = - B*S*I;
312         float dI = B*S*I - r*I;
313         float dR = r*I;
314         // step forward by h
315         S += h*dS;
316         I += h*dI;
317         R += h*dR;
318     }
319
320     particle->S = S;
321     particle->I = I;
322     particle->R = R;
323

```

```

324 }
325
326
327 /* Particle perturbation function to be run between iterations and passes
328
329 */
330 void perturbParticles(Particle * particles, int NP, int passnum, float
    coolrate) {
331
332     float coolcoef = exp( - (float) passnum / coolrate );
333
334     float spreadR0      = coolcoef * R0true / 3.0;
335     float spreadr       = coolcoef * rtrue  / 3.0;
336     float spreadsigma   = coolcoef * merr   / 3.0;
337     float spreadIinit   = coolcoef * 5.0    / 3.0;
338
339     float R0can, rcan, sigmacan, Iinitcan;
340
341     for (int n = 0; n < NP; n++) {
342
343         do {
344             R0can = particles[n].R0 + spreadR0*randn();
345         } while (R0can < 0);
346         particles[n].R0 = R0can;
347
348         do {
349             rcan = particles[n].r + spreadr*randn();
350         } while (rcan < 0);
351         particles[n].r = rcan;
352
353         do {
354             sigmacan = particles[n].sigma + spreadsigma*randn();
355         } while (sigmacan < 0);
356         particles[n].sigma = sigmacan;
357
358         do {
359             Iinitcan = particles[n].Iinit + spreadIinit*randn();
360         } while (Iinitcan < 0 || Iinitcan > 500);
361         particles[n].Iinit = Iinitcan;
362         particles[n].Sinit = N - Iinitcan;
363
364     }
365 }
366
367
368
369 /* Convenience function for particle resampling process
370
371 */
372 void copyParticle(Particle * dst, Particle * src) {
373
374     dst->R0      = src->R0;
375     dst->r       = src->r;
376     dst->sigma    = src->sigma;

```

```

377     dst->S      = src->S;
378     dst->I      = src->I;
379     dst->R      = src->R;
380     dst->Sinit  = src->Sinit;
381     dst->Iinit  = src->Iinit;
382     dst->Rinit  = src->Rinit;
383
384 }
385
386
387 /* Checks to see if particles are collapsed
388    This is done by checking if the standard deviations between the
389    particles' parameter
390    values are significantly close to one another. Spread threshold may
391    need to be tuned.
392    */
393 bool isCollapsed(Particle * particles, int NP) {
394     bool retVal;
395
396     float R0mean = 0, rmean = 0, sigmamean = 0, Sinitmean = 0, Iinitmean =
        0, Rinitmean = 0;
397
398     // means
399
400     for (int n = 0; n < NP; n++) {
401
402         R0mean      += particles[n].R0;
403         rmean       += particles[n].r;
404         sigmamean   += particles[n].sigma;
405         Sinitmean   += particles[n].Sinit;
406         Iinitmean   += particles[n].Iinit;
407         Rinitmean   += particles[n].Rinit;
408
409     }
410
411     R0mean      /= NP;
412     rmean       /= NP;
413     sigmamean   /= NP;
414     Sinitmean   /= NP;
415     Iinitmean   /= NP;
416     Rinitmean   /= NP;
417
418     float R0sd = 0, rsd = 0, sigmasd = 0, Sinitsd = 0, Iinitsd = 0,
        Rinitsd = 0;
419
420     for (int n = 0; n < NP; n++) {
421
422         R0sd      += ( particles[n].R0 - R0mean ) * ( particles[n].R0 -
            R0mean );
423         rsd       += ( particles[n].r - rmean ) * ( particles[n].r - rmean );
424         sigmasd   += ( particles[n].sigma - sigmamean ) * ( particles[n].
            sigma - sigmamean );

```

```

425         Sinitstd += ( particles[n].Sinit - Sinitmean ) * ( particles[n].
           Sinit - Sinitmean );
426         Iinitstd += ( particles[n].Iinit - Iinitmean ) * ( particles[n].
           Iinit - Iinitmean );
427         Rinitstd += ( particles[n].Rinit - Rinitmean ) * ( particles[n].
           Rinit - Rinitmean );
428
429     }
430
431     R0std      /= NP;
432     rsd        /= NP;
433     sigmasd    /= NP;
434     Sinitstd   /= NP;
435     Iinitstd   /= NP;
436     Rinitstd   /= NP;
437
438     if ( (R0std + rsd + sigmasd) < 1e-5)
439         retVal = true;
440     else
441         retVal = false;
442
443     return retVal;
444 }
445 }
446
447 void particleDiagnostics(ParticleInfo * partInfo, Particle * particles, int
    NP) {
448
449     float    R0mean      = 0.0,
450             rmean        = 0.0,
451             sigmamean    = 0.0,
452             Sinitmean    = 0.0,
453             Iinitmean    = 0.0,
454             Rinitmean    = 0.0;
455
456     // means
457
458     for (int n = 0; n < NP; n++) {
459
460         R0mean      += particles[n].R0;
461         rmean        += particles[n].r;
462         sigmamean    += particles[n].sigma;
463         Sinitmean    += particles[n].Sinit;
464         Iinitmean    += particles[n].Iinit;
465         Rinitmean    += particles[n].Rinit;
466
467     }
468
469     R0mean      /= NP;
470     rmean        /= NP;
471     sigmamean    /= NP;
472     Sinitmean    /= NP;
473     Iinitmean    /= NP;
474     Rinitmean    /= NP;

```

```

475
476 // standard deviations
477
478 float    R0sd    = 0.0,
479          rsd     = 0.0,
480          sigmasd = 0.0,
481          Sinitds = 0.0,
482          Iinitds = 0.0,
483          Rinitds = 0.0;
484
485 for (int n = 0; n < NP; n++) {
486
487     R0sd    += ( particles[n].R0 - R0mean ) * ( particles[n].R0 -
488               R0mean );
489     rsd     += ( particles[n].r - rmean ) * ( particles[n].r - rmean );
490     sigmasd += ( particles[n].sigma - sigmamean ) * ( particles[n].
491               sigma - sigmamean );
492     Sinitds += ( particles[n].Sinit - Sinitmean ) * ( particles[n].
493               Sinit - Sinitmean );
494     Iinitds += ( particles[n].Iinit - Iinitmean ) * ( particles[n].
495               Iinit - Iinitmean );
496     Rinitds += ( particles[n].Rinit - Rinitmean ) * ( particles[n].
497               Rinit - Rinitmean );
498
499 }
500
501 R0sd    /= NP;
502 rsd     /= NP;
503 sigmasd /= NP;
504 Sinitds /= NP;
505 Iinitds /= NP;
506 Rinitds /= NP;
507
508 partInfo->R0mean    = R0mean;
509 partInfo->R0sd      = R0sd;
510 partInfo->sigmamean = sigmamean;
511 partInfo->sigmasd   = sigmasd;
512 partInfo->rmean     = rmean;
513 partInfo->rsd       = rsd;
514 partInfo->Sinitmean = Sinitmean;
515 partInfo->Sinitds   = Sinitds;
516 partInfo->Iinitmean = Iinitmean;
517 partInfo->Iinitds   = Iinitds;
518 partInfo->Rinitmean = Rinitmean;
519 partInfo->Rinitds   = Rinitds;
520
521 }

```